

# IMPLEMENTING A PRIORITY QUEUE

## As a **sorted** Doubly-linked list:

Elements **are inserted based on their key** to ensure the priority queue is always sorted

What is the complexity of these methods?

- `size()`:  $O(1)$
- `empty()`:  $O(1)$
- `insert(e)`:  $O(n)$
- `min()`:  $O(1)$
- `removeMin()`:  $O(1)$

## As an **unsorted** Doubly-linked list:

This means that elements are **inserted at the end of the list**


What is the complexity of these methods?

- `size()`:  $O(1)$
- `empty()`:  $O(1)$
- `insert(e)`:  $O(1)$
- `min()`:  $O(n)$
- `removeMin()`:  $O(n)$

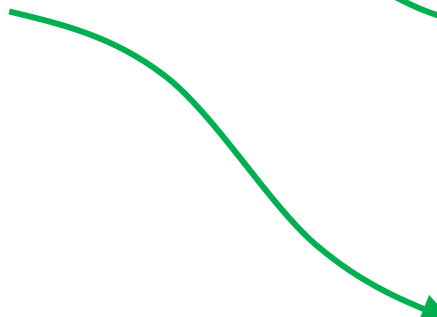
# PRIORITY QUEUE IMPLEMENTATION

- Let us look at a possible implementation of a priority queue as a **sorted linked list**:

```
template <typename E, typename C>
class ListPriorityQueue {
public:
    int size( ) const;
    bool empty( ) const;
    void insert(const E& e);
    const E& min( ) const;
    void removeMin( );
private:
    std::list<E> L;
    C isLess;
};
```



```
template <typename E, typename C>
const E& ListPriorityQueue<E,C>::min( ) const
{ return L.front( ); }
```

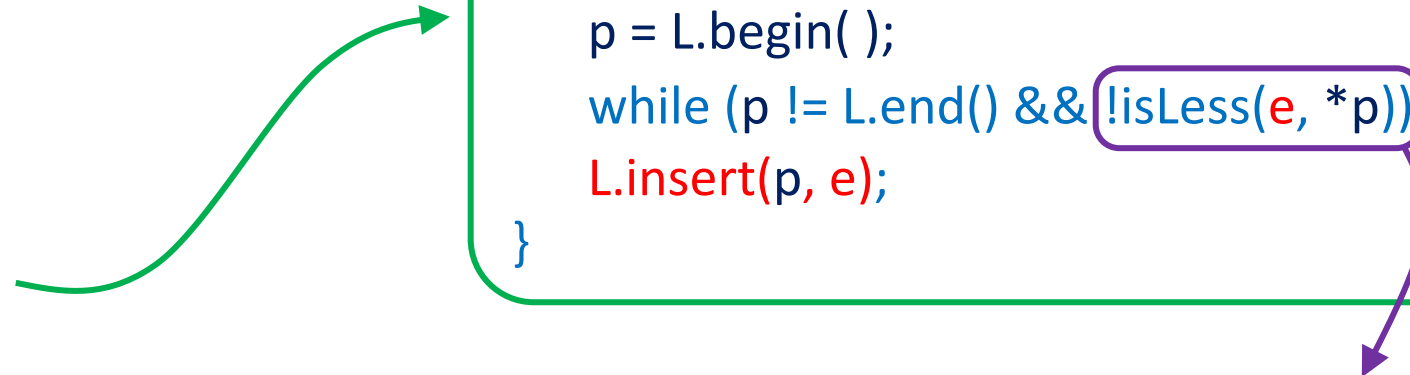


```
template <typename E, typename C>
void ListPriorityQueue<E,C>::removeMin( )
{ L.pop_front( ); }
```

# PRIORITY QUEUE IMPLEMENTATION

- Let us look at a possible implementation of a priority queue as a **sorted linked list**:

```
template <typename E, typename C>
class ListPriorityQueue {
public:
    int size( ) const;
    bool empty( ) const;
    void insert(const E& e);
    const E& min( ) const;
    void removeMin( );
private:
    std::list<E> L;
    C isLess;
};
```



```
template <typename E, typename C>
void ListPriorityQueue<E,C>::insert(const E& e) {
    typename std::list<E>::iterator p;
    p = L.begin( );
    while (p != L.end() && !isLess(e, *p)) ++p;
    L.insert(p, e);
}
```

This implies that the list is sorted in an ascending order (i.e., the first element has the smallest key, while the last element has the largest)

# SORTING WITH A PRIORITY QUEUE

- Given an **unsorted list**  $L$  of  $n$  elements, we can **sort**  $L$  using a priority queue  $Q$  as follows:
  - **Phase 1:** Put the elements of  $L$  into a priority queue  $P$  using  $n$  `insert( $e$ )` operations
  - **Phase 2:** Extract the elements from  $P$  in an **ascending order** using  $n$  combinations of `min()` and `removeMin()` operations

**Algorithm** PriorityQueueSort(  $L, P$  ):

```
while ! $L$ .empty() do
     $e \leftarrow L$ .front
     $L$ .pop_front() {remove element  $e$  from the list}
     $P$ .insert( $e$ )    {add element  $e$  to the priority queue}
while ! $P$ .empty() do
     $e \leftarrow P$ .min()
     $P$ .removeMin() {remove the smallest element  $e$  from the queue}
     $L$ .push_back( $e$ ) {append element  $e$  to the back of  $L$ }
```

Phase 1 ← { while ! $L$ .empty() do ... }

Phase 2 ← { while ! $P$ .empty() do ... }

# SORTING WITH A PRIORITY QUEUE

- Given an **unsorted list** **L** of **n** elements, we can **sort** **L** using a priority queue **Q** as follows:
  - **Phase 1:** Put the elements of **L** into a priority queue **P** using **n** **insert(e)** operations
  - **Phase 2:** Extract the elements from **P** in an **ascending order** using **n** combinations of **min()** and **removeMin()** operations
- This approach depends on how **insert(e)**, **min()** and **removeMin()** are implemented:
  - **Option 1:** If **insert(e)** inserts **e** at the **end of the queue** (which can be done instantly), then **min()** and **removeMin()** must **search the queue** to find the minimum element!
  - **Option 2:** If **insert(e)** inserts **e** **based on its key** (so that, after the insert, the queue is always sorted), then **min()** and **removeMin()** can **instantly retrieve the minimum element** (this would be the first element in the queue)!
- If we go with **Option 1**, we end up with an algorithm called “**selection sort**”, while if we go with **Option 2** we end up with an algorithm called “**insertion sort**”

# SORTING WITH A PRIORITY QUEUE

	<i>List L</i>	<i>Priority Queue P</i>
Input	(7, 4, 8, 2, 5, 3, 9)	()
Phase 1	(4, 8, 2, 5, 3, 9)	(7)
	(8, 2, 5, 3, 9)	(7, 4)
	⋮	⋮
	( )	(7, 4, 8, 2, 5, 3, 9)
Phase 2	(2)	(7, 4, 8, 5, 3, 9)
	(2, 3)	(7, 4, 8, 5, 9)
	(2, 3, 4)	(7, 8, 5, 9)
	(2, 3, 4, 5)	(7, 8, 9)
	(2, 3, 4, 5, 7)	(8, 9)
	(2, 3, 4, 5, 7, 8)	(9)
	(2, 3, 4, 5, 7, 8, 9)	( )

Example of **selection sort**

	<i>List L</i>	<i>Priority Queue P</i>
Input	(7, 4, 8, 2, 5, 3, 9)	()
Phase 1	(4, 8, 2, 5, 3, 9)	(7)
	(8, 2, 5, 3, 9)	(4, 7)
	(2, 5, 3, 9)	(4, 7, 8)
	(5, 3, 9)	(2, 4, 7, 8)
	(3, 9)	(2, 4, 5, 7, 8)
	(9)	(2, 3, 4, 5, 7, 8)
	( )	(2, 3, 4, 5, 7, 8, 9)
Phase 2	(2)	(3, 4, 5, 7, 8, 9)
	(2, 3)	(4, 5, 7, 8, 9)
	⋮	⋮
	(2, 3, 4, 5, 7, 8, 9)	( )

Example of **insertion sort**

# COMPLEXITY ANALYSIS

- Let's compare the **complexity**. Remember:
  - **Selection sort:** `insert(e)` inserts `e` at the **end of the queue**, while `removeMin()` and `min()` **search the queue** to find the minimum element.
  - **Insertion sort:** `insert(e)` inserts `e` **based on its key** (so that, after the insert, the queue is always sorted), while `min()` and `removeMin()` **retrieve the minimum element**
- With **selection-sort**:
  - The total time needed for the *first phase* →  $O(n)$
  - The total time needed for the *second phase* →  $O(n^2)$

The first `removeMin()` operation takes  $O(n)$  time, the second one takes  $O(n - 1)$  time, etc. Thus, inserting `n` elements back to `L` takes:

$$n + n - 1 + \dots + 2 + 1 = \sum_{x=1}^n x = \frac{n(n+1)}{2} \quad \text{which is } O(n^2)$$

- With **insertion sort**, the opposite is true; the *first phase* takes  $O(n^2)$  time, while the *second* takes  $O(n)$  time.

# SORTING WITH A PRIORITY QUEUE

- To summarize, given a list  $L$  of  $n$  elements, we can sort  $L$  using a priority queue  $Q$ :
  - **Phase 1:** Put the elements of  $L$  into a priority queue  $P$  using  $n$  `insert(e)` operations
  - **Phase 2:** Extract the elements from  $P$  in an **ascending order** using  $n$  combinations of `min()` and `removeMin()` operations
- Depending on how the elements are inserted in  $P$ , we have **two options**:
  - **Selection sort:** Each `insert(e)` takes  $O(1)$ , while each `removeMin()/min()` takes  $O(n)$
  - **Insertion sort:** Each `insert(e)` takes  $O(n)$ , while each `removeMin()/min()` takes  $O(1)$
- Both of these algorithms run in  $O(n^2)$  time. However, there is a **third option**:
  - **Heap sort:** `insert(e)`, `removeMin()` and `min()` take  $O(n \log n)$  time!
- Before explaining how it works, we need to discuss “**complete binary tree**” and “**heaps**”.



# PRIORITY QUEUE STL

- C++ provides a readily-available **priority queue** as part of the **queue STL** in C++. The **priority\_queue** class is **templated** with three parameters:
  1. The **base type** of the elements; **you must specify this argument**.
  2. The **STL container** in which the elements are stored; if you don't specify this argument, **an STL vector is used by default**.
  3. The **comparator object**. If you don't specify this argument, **the standard C++ less-than operator ("<") is used by default**.
- Here are example's of how to define a priority queue:

```
#include <queue>
using namespace std;
priority_queue<int> p1;
priority_queue<Point2D, vector<Point2D>, LeftRight> p2;
```

# PRIORITY QUEUE ADT VS. STL

- In the case of **priority queue**, the differences are:
  - Instead of `insert(e)` which can be implemented either by inserting `e` at the end or based on its key, we have `push(e)` which **inserts `e` based on its key**.
  - Instead of `min()` and `removeMin()` which return and remove an element with the smallest key, we have `top()` and `pop()` which return and remove an element with the **largest key**.





# COMPLETE BINARY TREES



# COMPLETE BINARY TREE


- A binary tree  $T$  with height  $h$  is **complete** if:
  - Levels  $0, 1, 2, \dots, h-1$  of  $T$  have the **maximum number of nodes** possible
  - The remaining nodes fill level  $h$  **from left to right**
- Examples:




# COMPLETE BINARY TREE

**Proposition:** A complete binary tree  $T$  with  $n$  nodes has  $h = \lfloor \log n \rfloor$


**Experimental Justification:**



$$h = \lfloor \log 12 \rfloor = \lfloor 3.58 \rfloor = 3$$






$$h = \lfloor \log 15 \rfloor = \lfloor 3.90 \rfloor = 3$$



$$h = \lfloor \log 21 \rfloor = \lfloor 4.39 \rfloor = 4$$


# COMPLETE BINARY TREE **ADT**

- A **complete binary tree ADT** supports all the functions of a **binary tree**, in addition to the following:
  - **add(e)**: Add **e** to **T** (and return a new external node **v** storing element **e**) such that the resulting tree is a complete binary tree with last node **v**.
  - **remove()**: Remove the last node of **T** and return its element.
- By using only these update operations, the binary tree is guaranteed to be complete



# VECTOR-BASED IMPLEMENTATION


- Remember how we said **we can represent a binary tree as a vector**



- This representation is especially suitable for complete binary trees, because we always **fill the last level of the tree from left to right**, which corresponds to **filling the vector from left to right**!

# VECTOR-BASED IMPLEMENTATION

- Remember how we said **we can represent a binary tree as a vector**




- This representation is especially suitable for complete binary trees, because we always **fill the last level of the tree from left to right**, which corresponds to **filling the vector from left to right**!



# VECTOR-BASED IMPLEMENTATION

- Remember how we said **we can represent a binary tree as a vector**



- What is the complexity of **add(e)** and **remove()**?

If we use an extendable array these methods would take  **$O(1)$  amortized time!**

# C++ VECTOR-BASED IMPLEMENTATION

```
template <typename E>
class VectorCompleteTree {
private:
    std::vector<E> V; // the vector in which the tree will be stored
```

protected:

```
    Position pos(int i) { return V.begin() + i; }
```

```
    int idx(const Position& p) const { return p - V.begin(); }
```

public:

```
    typedef typename std::vector<E>::iterator Position; // a position in the tree
```

```
    VectorCompleteTree() : V(1) { } // constructor
```

```
    int size() const
```

getters {

```
    Position left(const Position& p)
    Position right(const Position& p)
    Position parent(const Position& p)
    Position root()
    Position last()
```

checkers {


```
    bool hasLeft(const Position& p) const
    bool hasRight(const Position& p) const
    bool isRoot(const Position& p) const
```

modifiers {

```
    void add(const E& e)
    void remove()
    void swap(const Position& p, const Position& q)
};
```

This method maps an index  $i$  to a position


It returns the position of the 1<sup>st</sup> element +  $i$



# C++ VECTOR-BASED IMPLEMENTATION

```
template <typename E>
class VectorCompleteTree {
private:
    std::vector<E> V; // the vector in which the tree will be stored
protected:
    Position pos(int i) { return V.begin() + i; }
    int idx(const Position& p) const { return p - V.begin(); }
public:
    typedef typename std::vector<E>::iterator Position; // a position in the tree
    VectorCompleteTree() : V(1) {} // constructor
    int size() const
    {
        Position left(const Position& p)
        Position right(const Position& p)
        Position parent(const Position& p)
        Position root()
        Position last()
    }
    {
        bool hasLeft(const Position& p) const
        bool hasRight(const Position& p) const
        bool isRoot(const Position& p) const
    }
    {
        void add(const E& e)
        void remove()
        void swap(const Position& p, const Position& q)
    }
};
```

This method returns the index of the element at position  $p$ . It returns the difference between the position of this element and the position of the 1<sup>st</sup> element.




# C++ VECTOR-BASED IMPLEMENTATION

```
template <typename E>
class VectorCompleteTree {
private:
    std::vector<E> V; // the vector in which the tree will be stored
protected:
    Position pos(int i) { return V.begin() + i; }
    int idx(const Position& p) const { return p - V.begin(); }
public:
    typedef typename std::vector<E>::iterator Position; // a position in the tree
    VectorCompleteTree() : V(1) {} // constructor
    int size() const { return V.size() - 1; }
    Position left(const Position& p) { return pos(2*idx(p)); }
    Position right(const Position& p) { return pos(2*idx(p) + 1); }
    Position parent(const Position& p) { return pos(idx(p)/2); }
    Position root() { return pos(1); }
    Position last() { return pos(size()); }
    bool hasLeft(const Position& p) const { return 2*idx(p) <= size(); }
    bool hasRight(const Position& p) const { return 2*idx(p) + 1 <= size(); }
    bool isRoot(const Position& p) const { return idx(p) == 1; }
    void add(const E& e) { V.push_back(e); }
    void remove() { V.pop_back(); }
    void swap(const Position& p, const Position& q) { E e = *q; *q = *p; *p = e; }
};
```

getters

checkers

modifiers






# HEAPS



# THE HEAP DATA STRUCTURE

- A **heap** is a **complete binary tree** that stores a collection of **elements with their associated keys** at its nodes, and that **satisfies the following property**:
  - **Heap-Order Property**: For every node **v** other than the root:  
 $(\text{key associated with } v) \geq (\text{key associated with } v\text{'s parent})$



- To implement a **priority queue** using a heap, we'll need a **comparator**.