

42

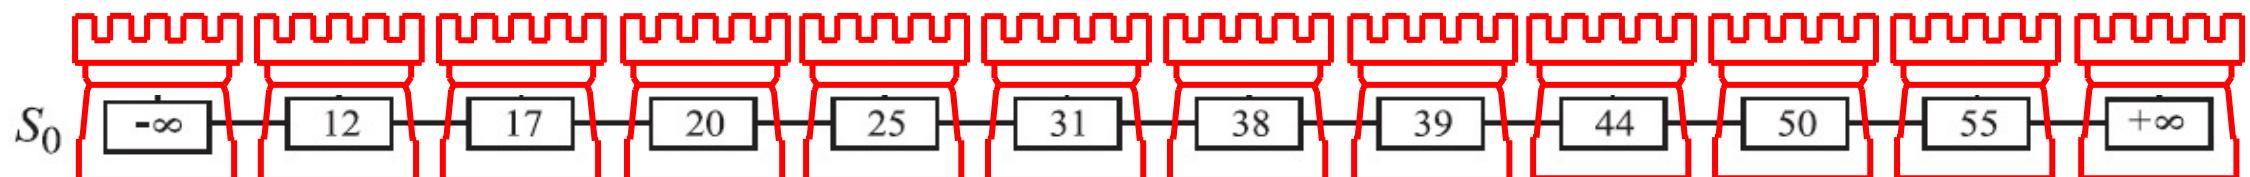
SKIP LISTS

SKIP LISTS

- As mentioned earlier, when implementing an ordered map using a **search table** :
 - **Finding** an entry whose key = k takes $O(\log n)$ time (since it uses binary search)
 - **Inserting** or **erasing** an entry takes $O(n)$ time (since we must shift other entries)
- However, if we implement an ordered map using a “**skip list**”, then:
 - **Finding** an entry whose key = k takes $O(\log n)$ time **on average**
 - **Inserting** or **erasing** an entry takes $O(\log n)$ time **on average**
- Here, the notion of **average time complexity** depends on the use of a **random-number generator** in the implementation of the insertions, to help decide where to place the new entry.
- The running time is averaged over all possible outcomes of the random numbers used when inserting entries.

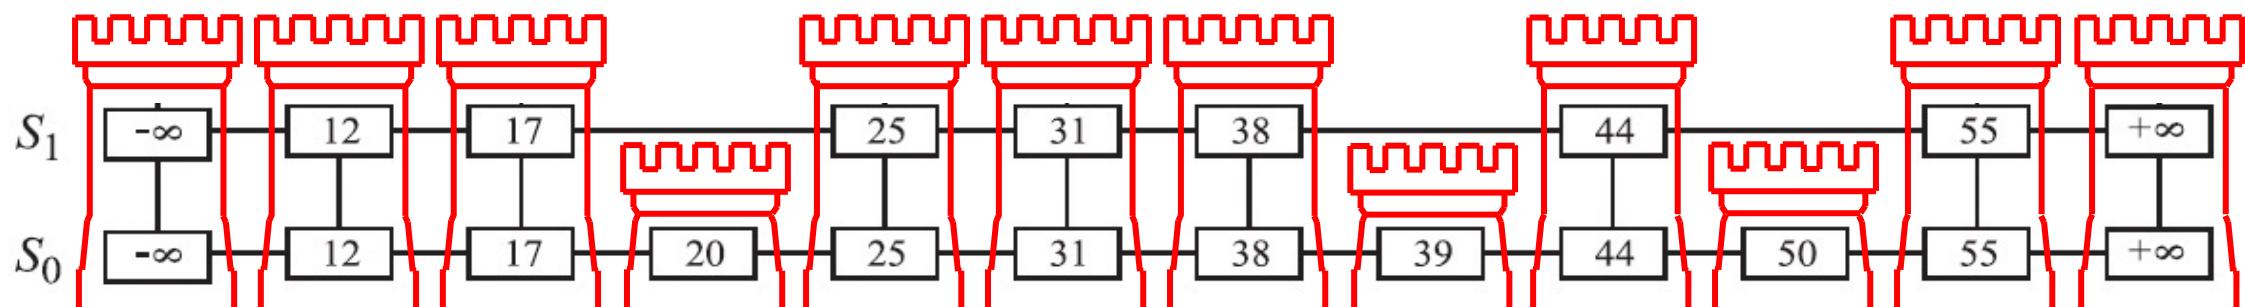
A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list



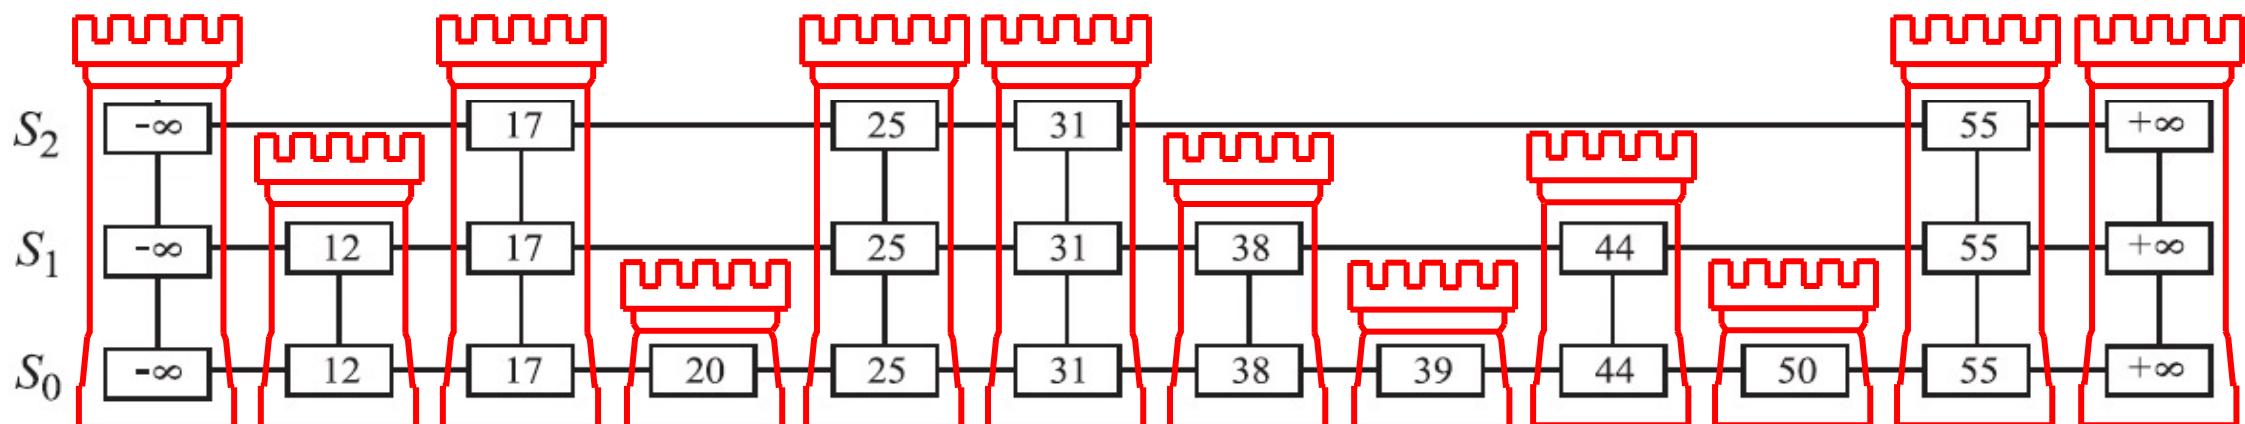
A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list
- At each level, each entry has a **50/50 chance of building an additional level!**
 - Each **tower** is implemented as a doubly-linked list



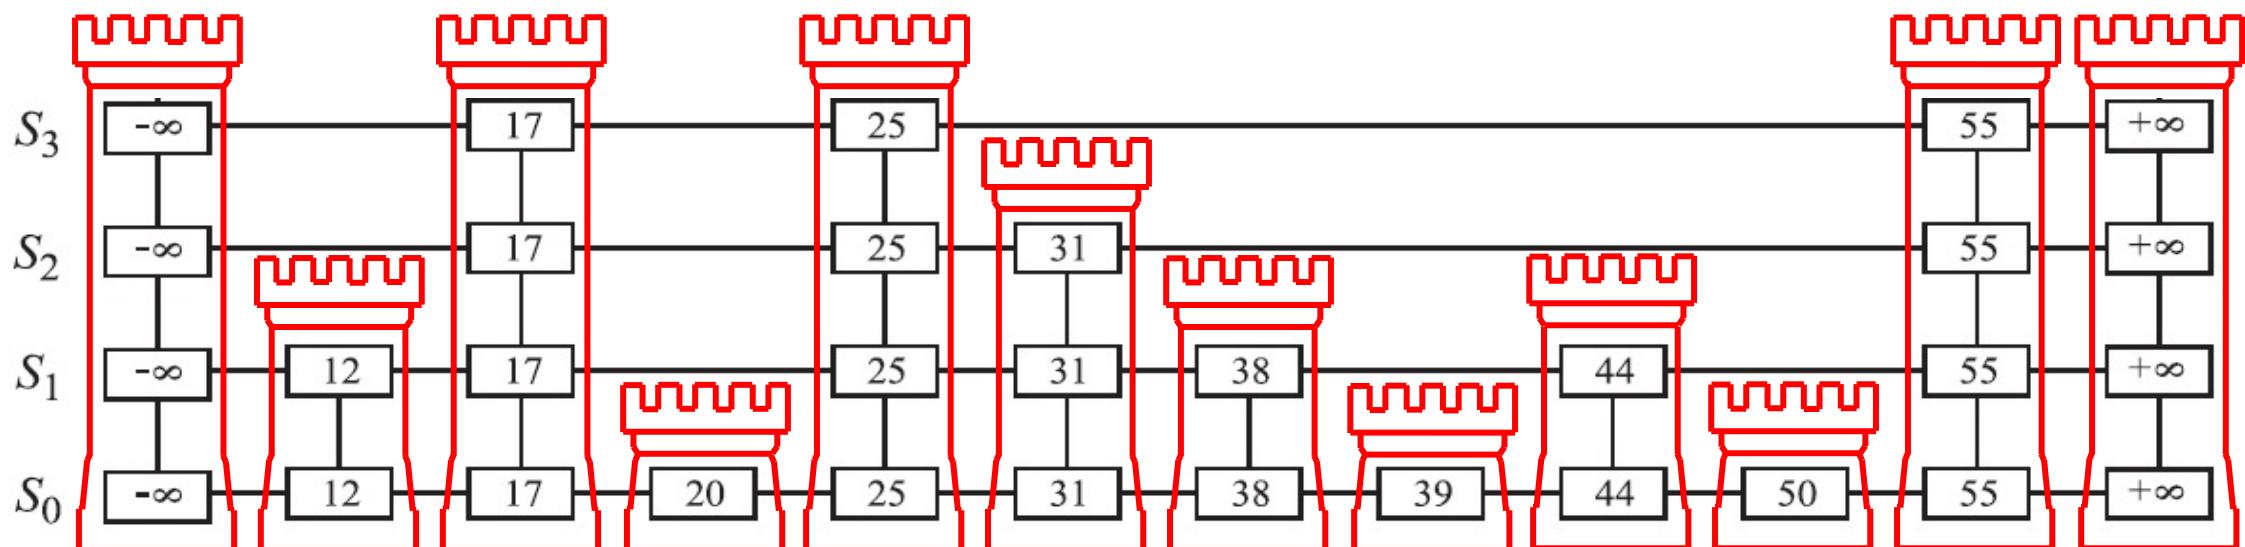
A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list
- At each level, each entry has a **50/50 chance of building an additional level!**
 - Each **tower** is implemented as a doubly-linked list



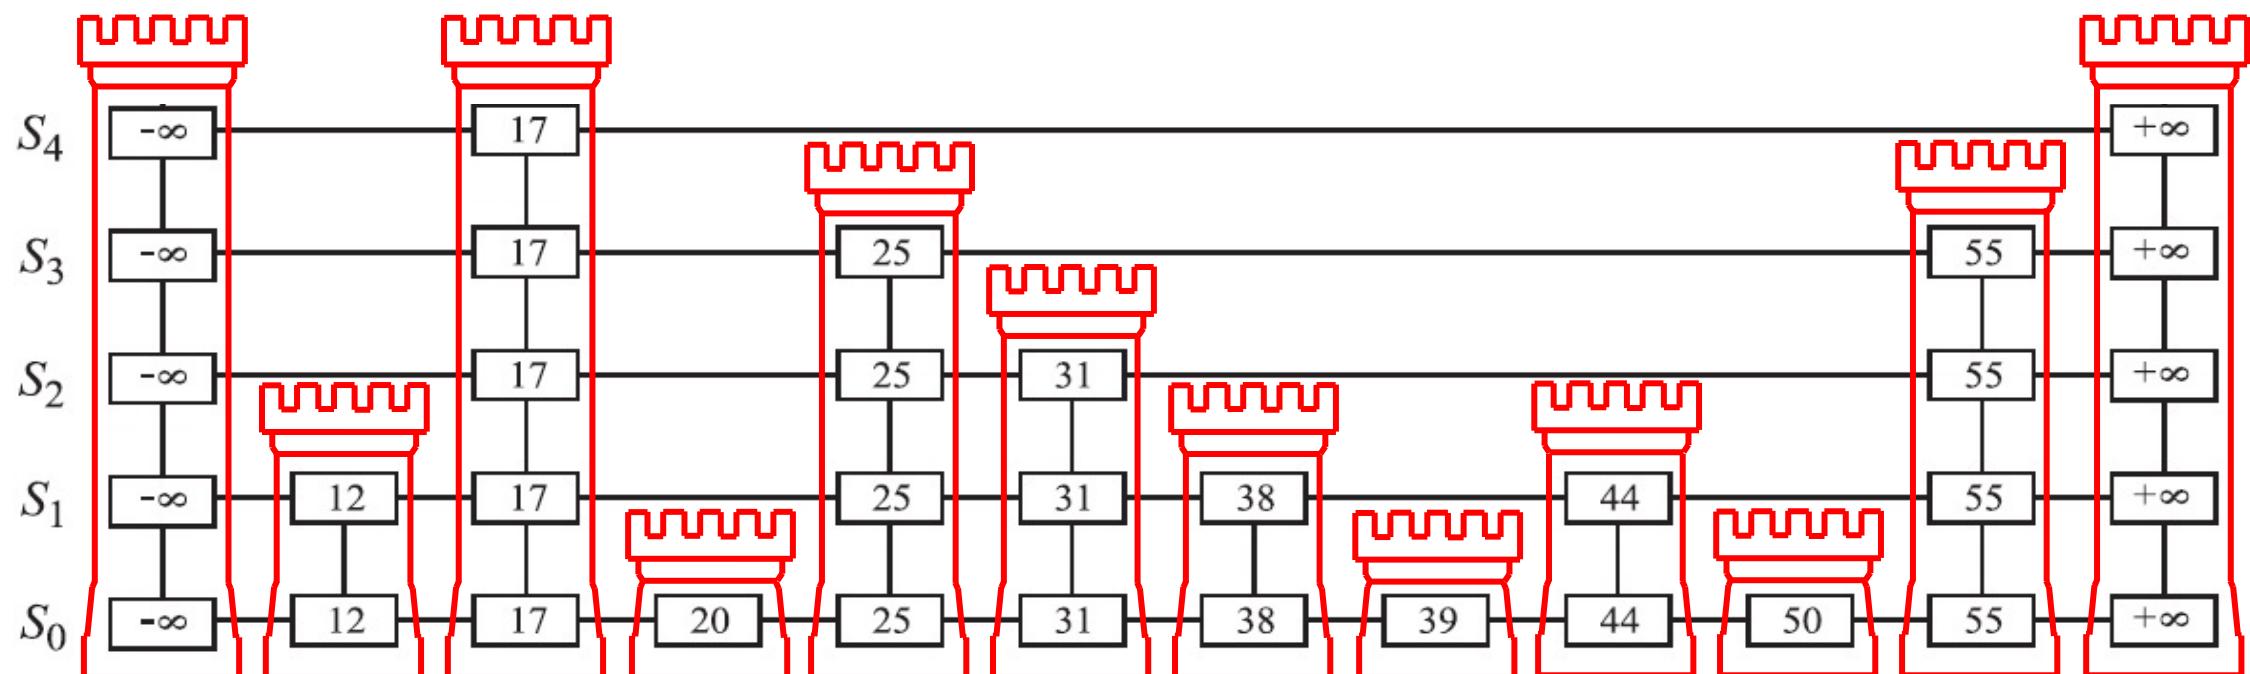
A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list
- At each level, each entry has a **50/50 chance of building an additional level!**
 - Each **tower** is implemented as a doubly-linked list



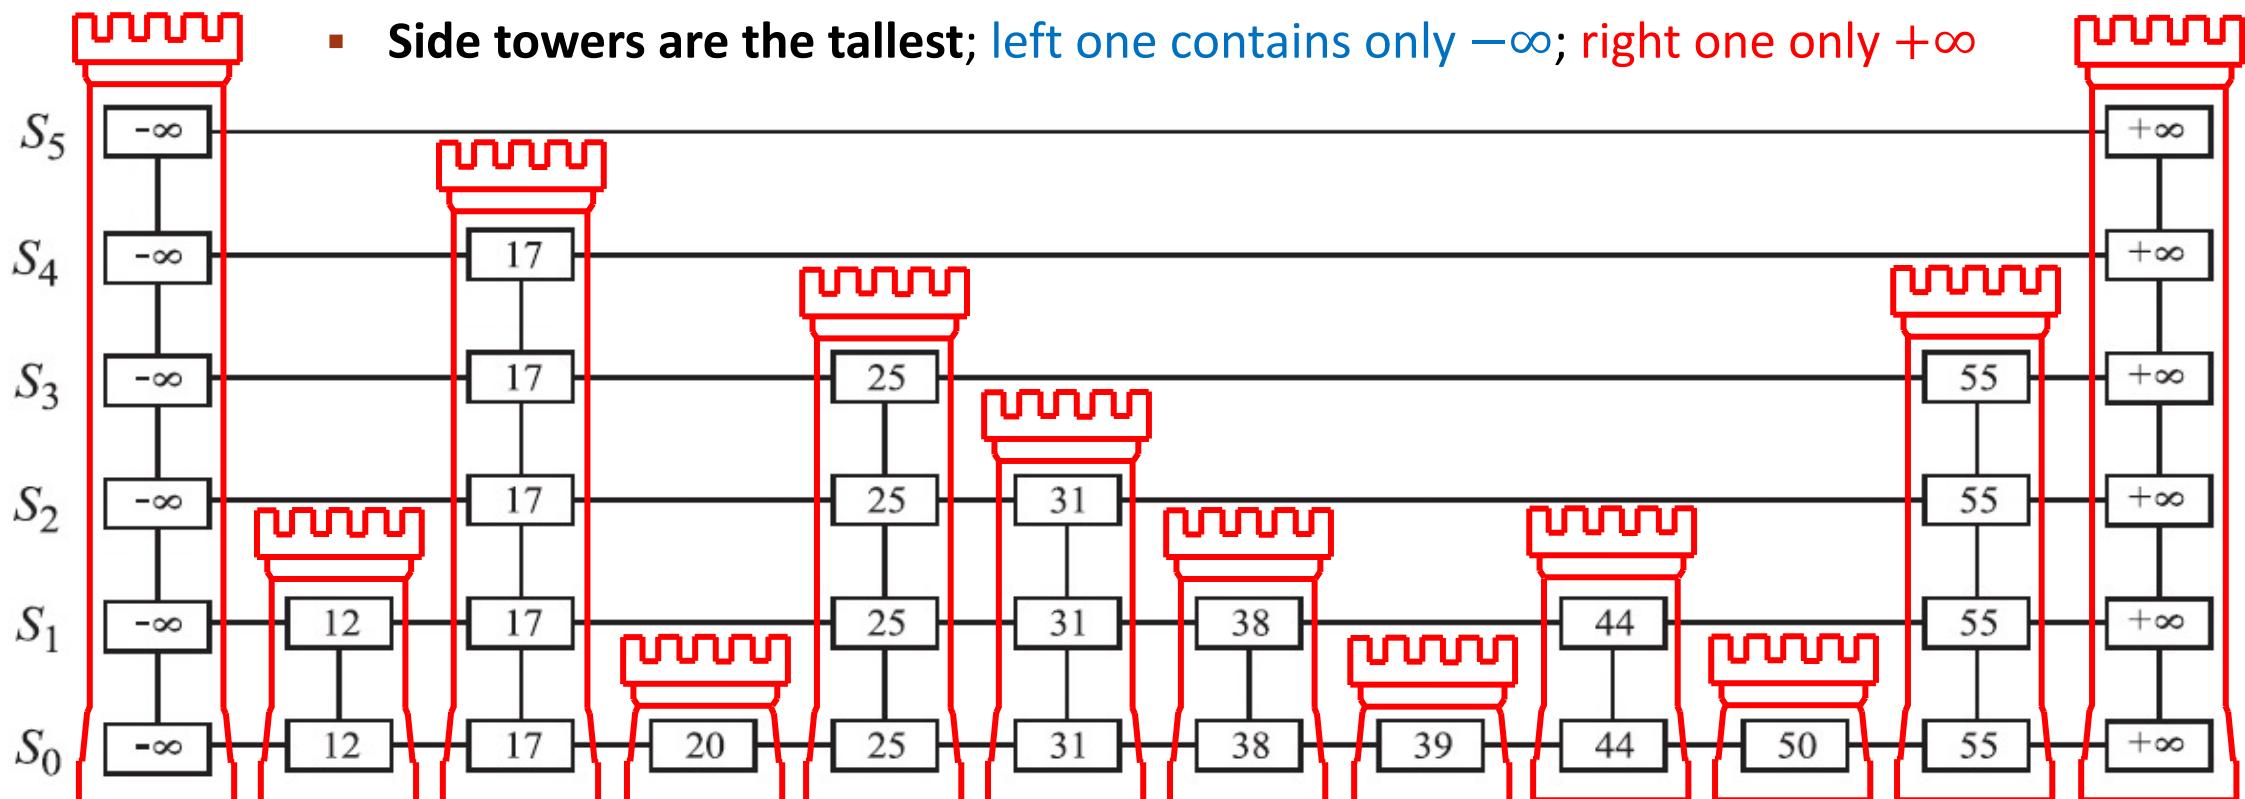
A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list
- At each level, each entry has a **50/50 chance of building an additional level!**
 - Each **tower** is implemented as a doubly-linked list



A skip list can be arranged into **levels** and **towers**:

- The **bottom level**, S_0 , contains all entries sorted (plus the $-\infty$ and $+\infty$)
 - For simplicity, **this illustration only shows the keys** of the entries
 - Each **level** is implemented as a doubly-linked list
- At each level, each entry has a **50/50 chance of building an additional level!**
 - Each **tower** is implemented as a doubly-linked list
- **Side towers are the tallest; left one contains only $-\infty$; right one only $+\infty$**



SKIP LIST – TRAVERSAL

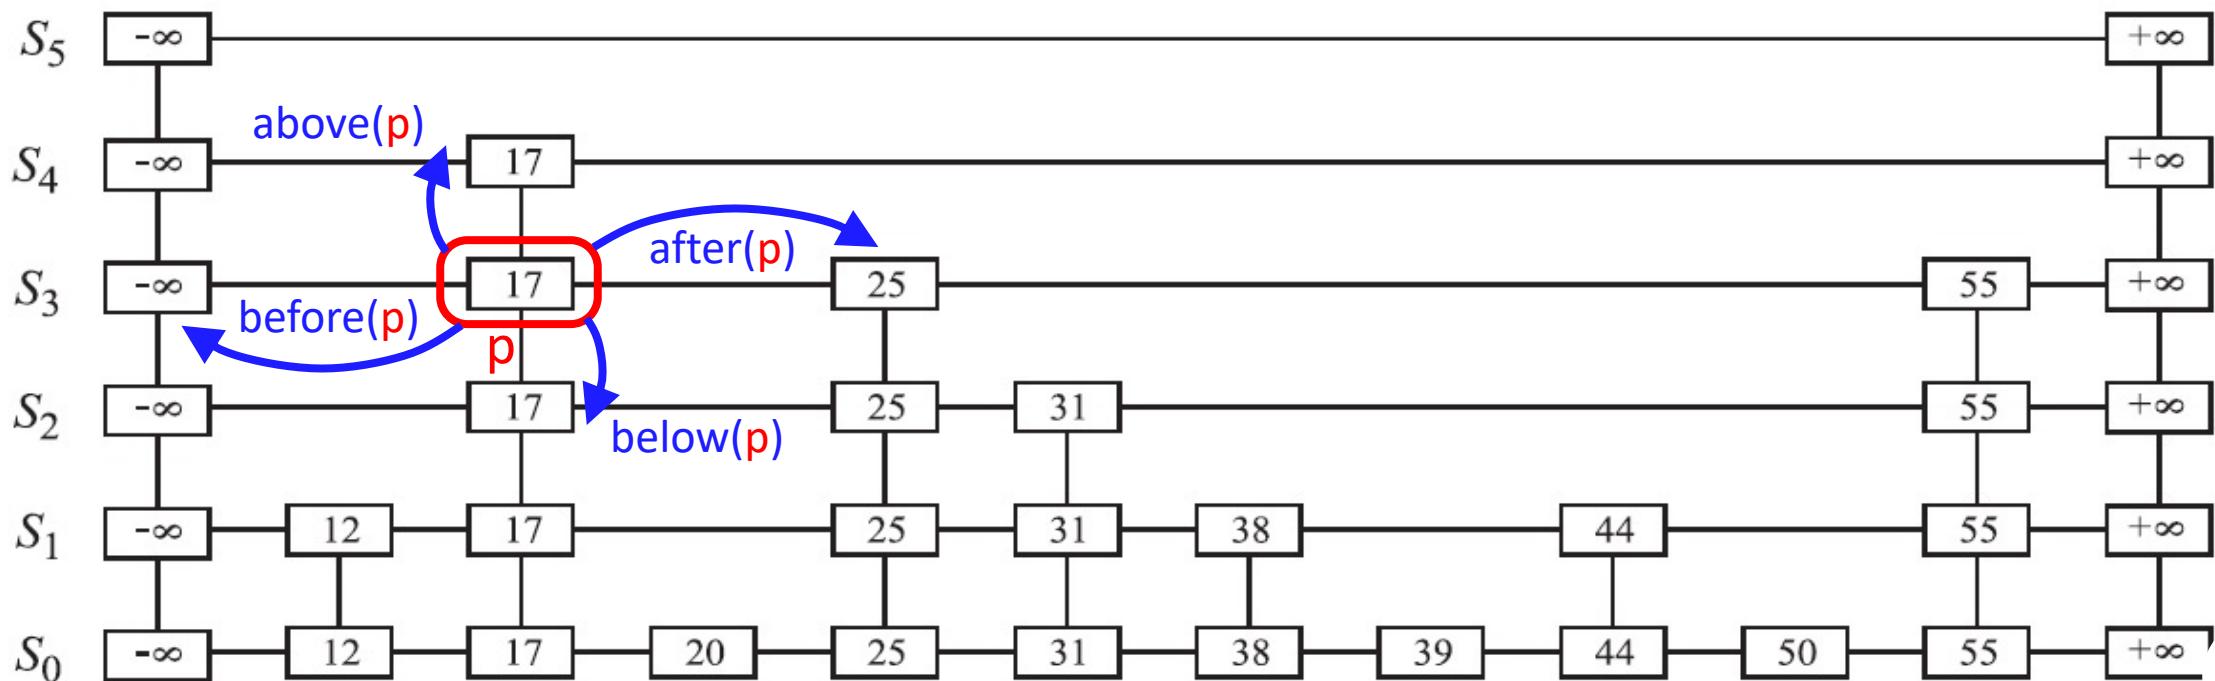
A skip list can be traversed via **positions**, p , using the following operations:

$\text{after}(p)$: Return the position following p on the same level.

$\text{before}(p)$: Return the position preceding p on the same level.

$\text{below}(p)$: Return the position below p in the same tower.

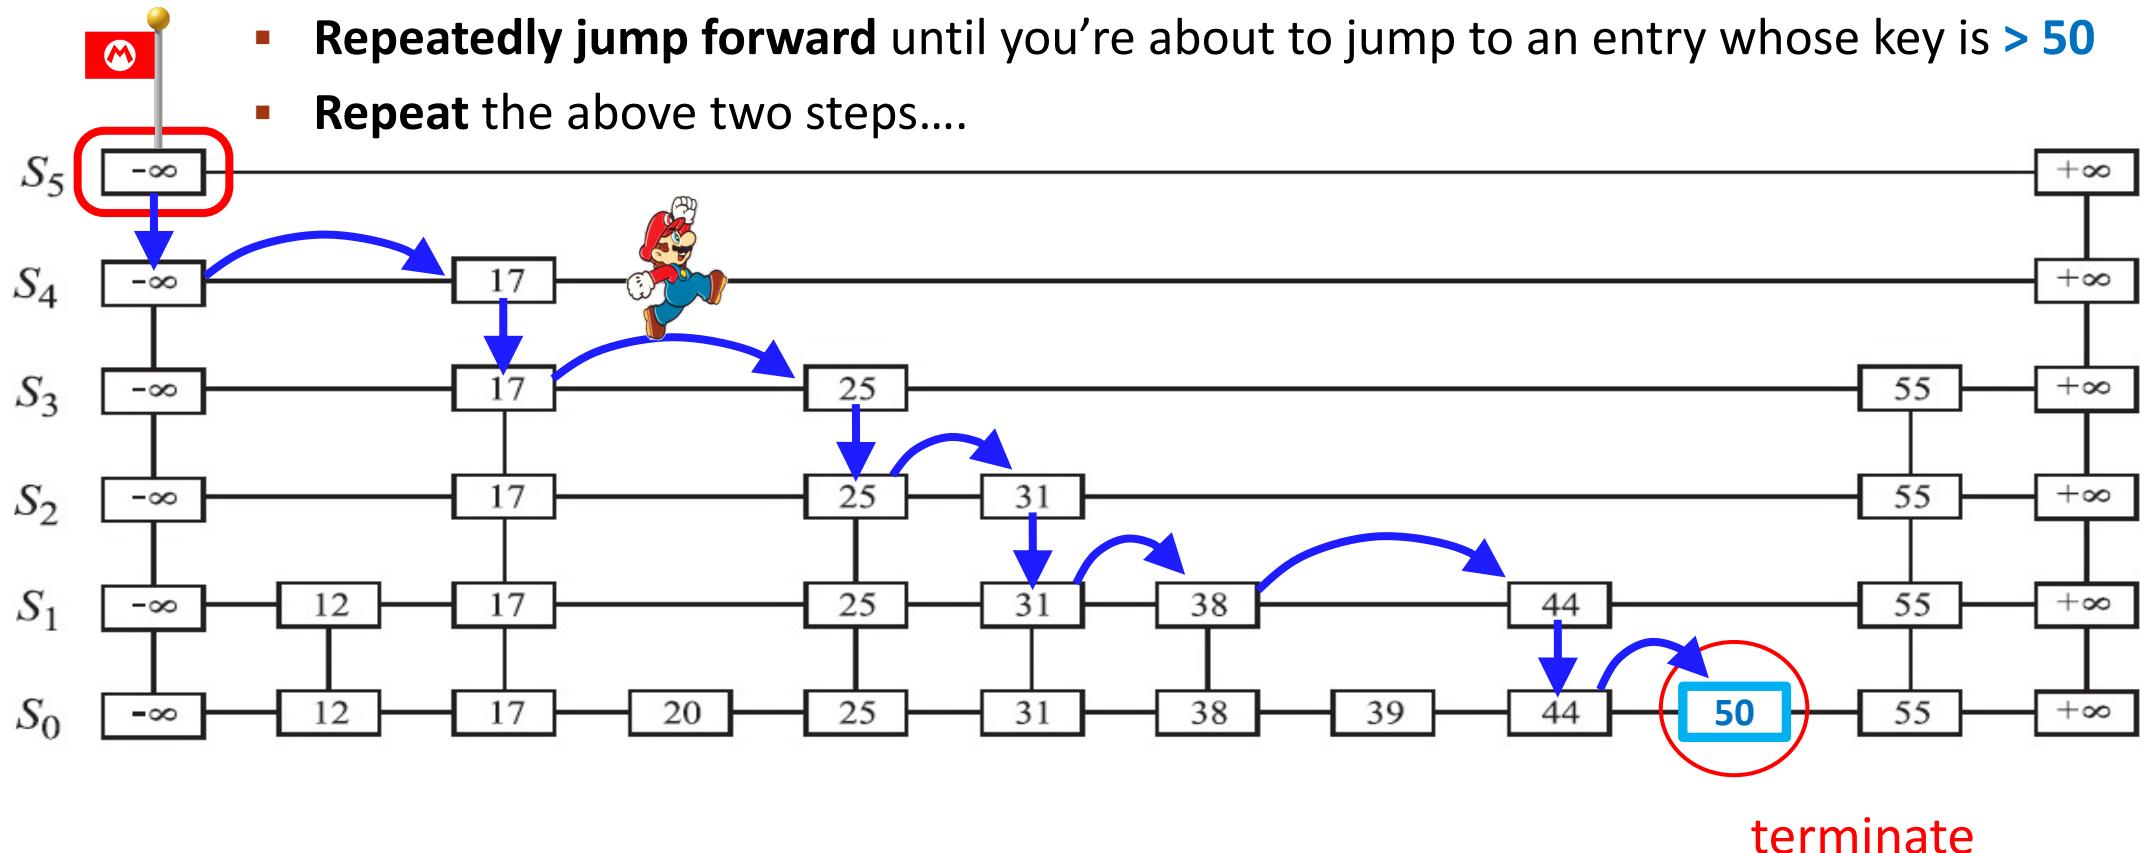
$\text{above}(p)$: Return the position above p in the same tower.



SKIP LIST – SEARCH

Here is an example of how to **search** a skip list, S , for the entry whose key = 50:

- Set position p to be at the **top of the left tower**:
- **Jump once downward** one step by calling $p = S.\text{below}(p)$; unless you've reached the bottom, in which case the search terminates.
- **Repeatedly jump forward** until you're about to jump to an entry whose key is > 50
- **Repeat** the above two steps....



SKIP LIST – SEARCH

Algorithm SkipSearch(k):

Input: A search key k

Output: Position p in the bottom list S_0 such that the entry at p has the largest key *less than or equal* to k

$p \leftarrow s$

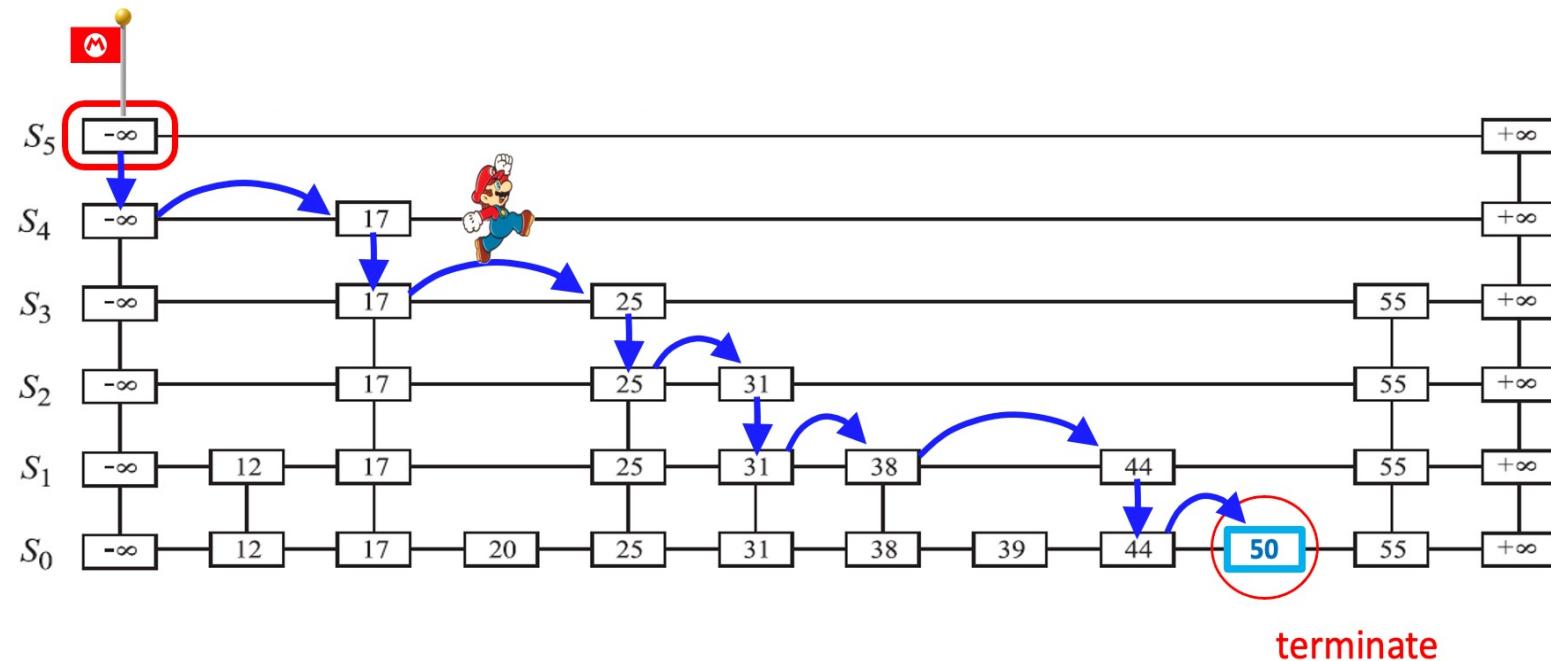
while $\text{below}(p) \neq \text{null}$ **do**

$p \leftarrow \text{below}(p)$ {drop down}

while $k \geq \text{after}(p).\text{key}()$ **do**

$p \leftarrow \text{after}(p)$ {scan forward}

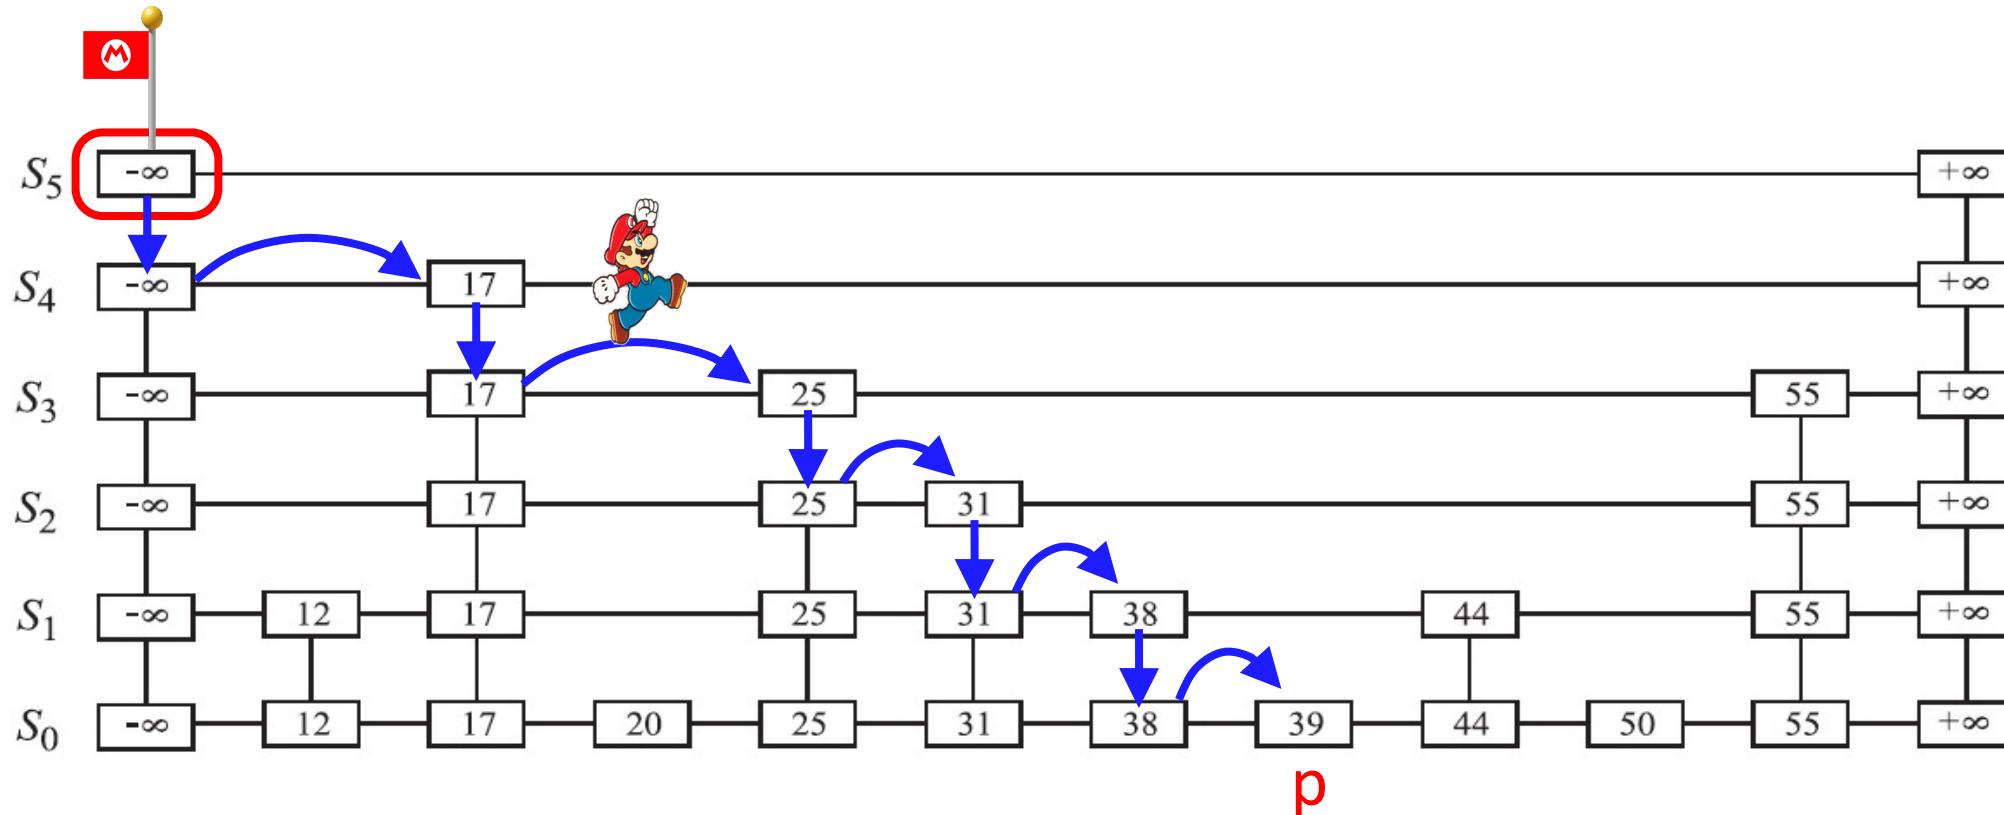
return p



SKIP LIST – INSERTION

Here is an example of how to **insert** an entry with key = **42**:

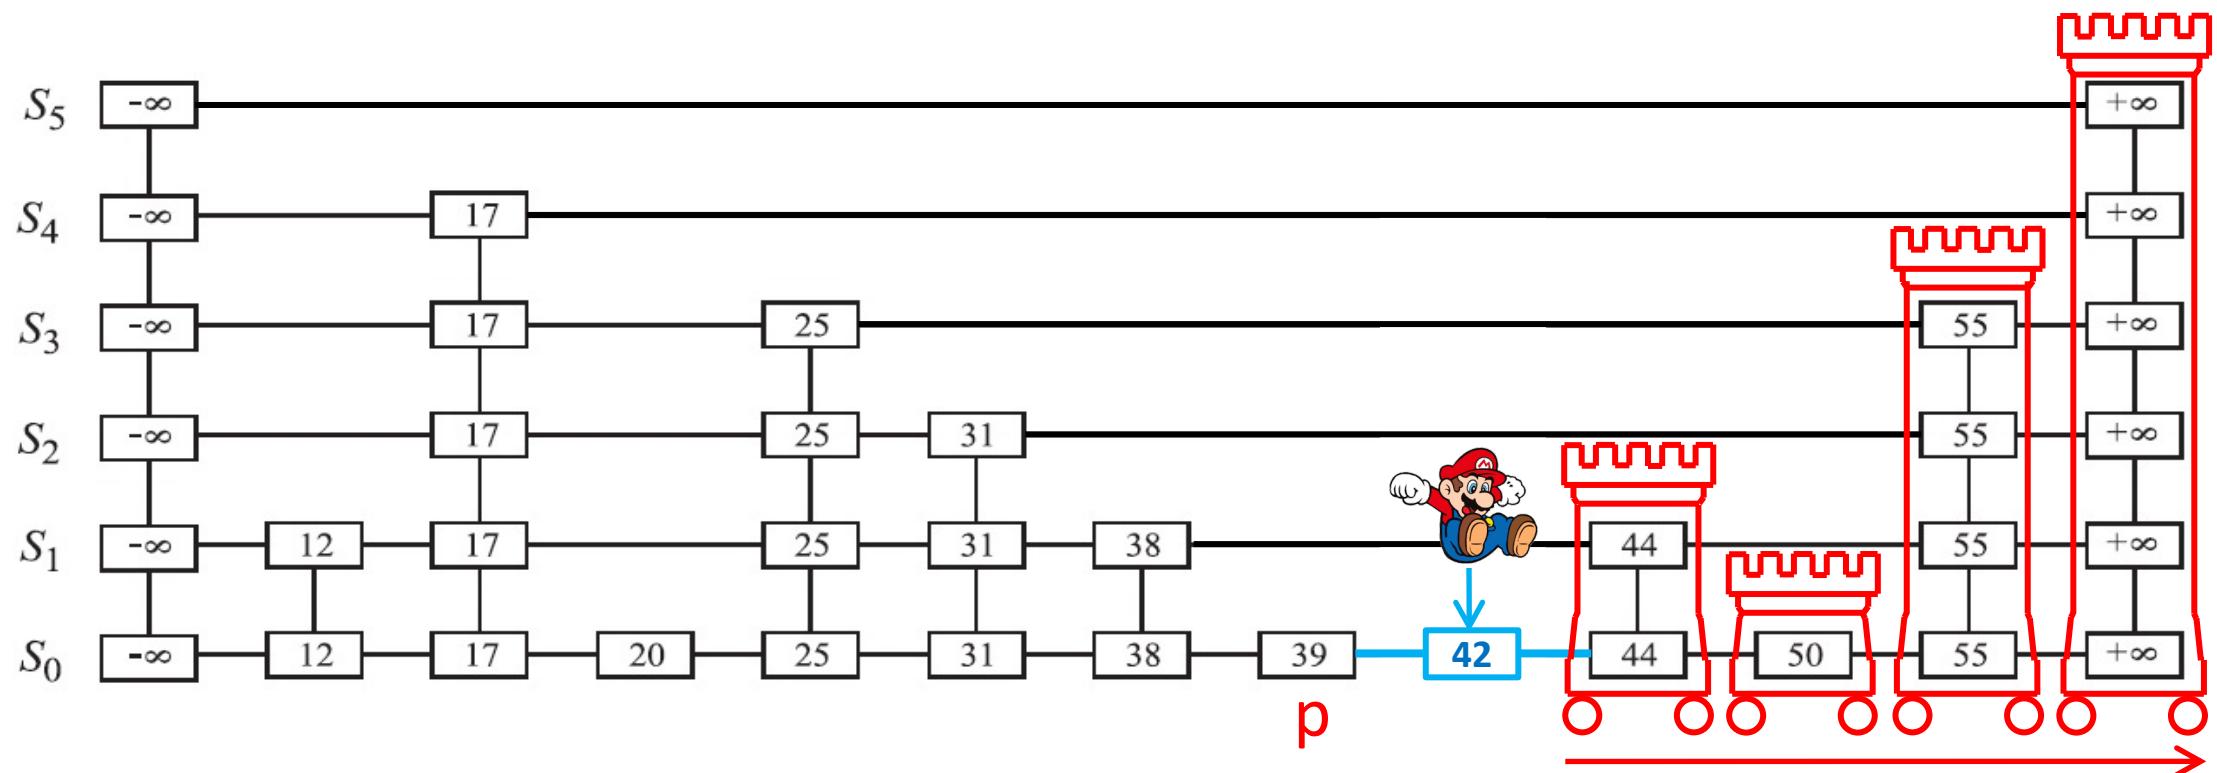
- **Search** for the largest key that is ≤ 42 , and set **p** to be its position



SKIP LIST – INSERTION

Here is an example of how to **insert** an entry with key = **42**:

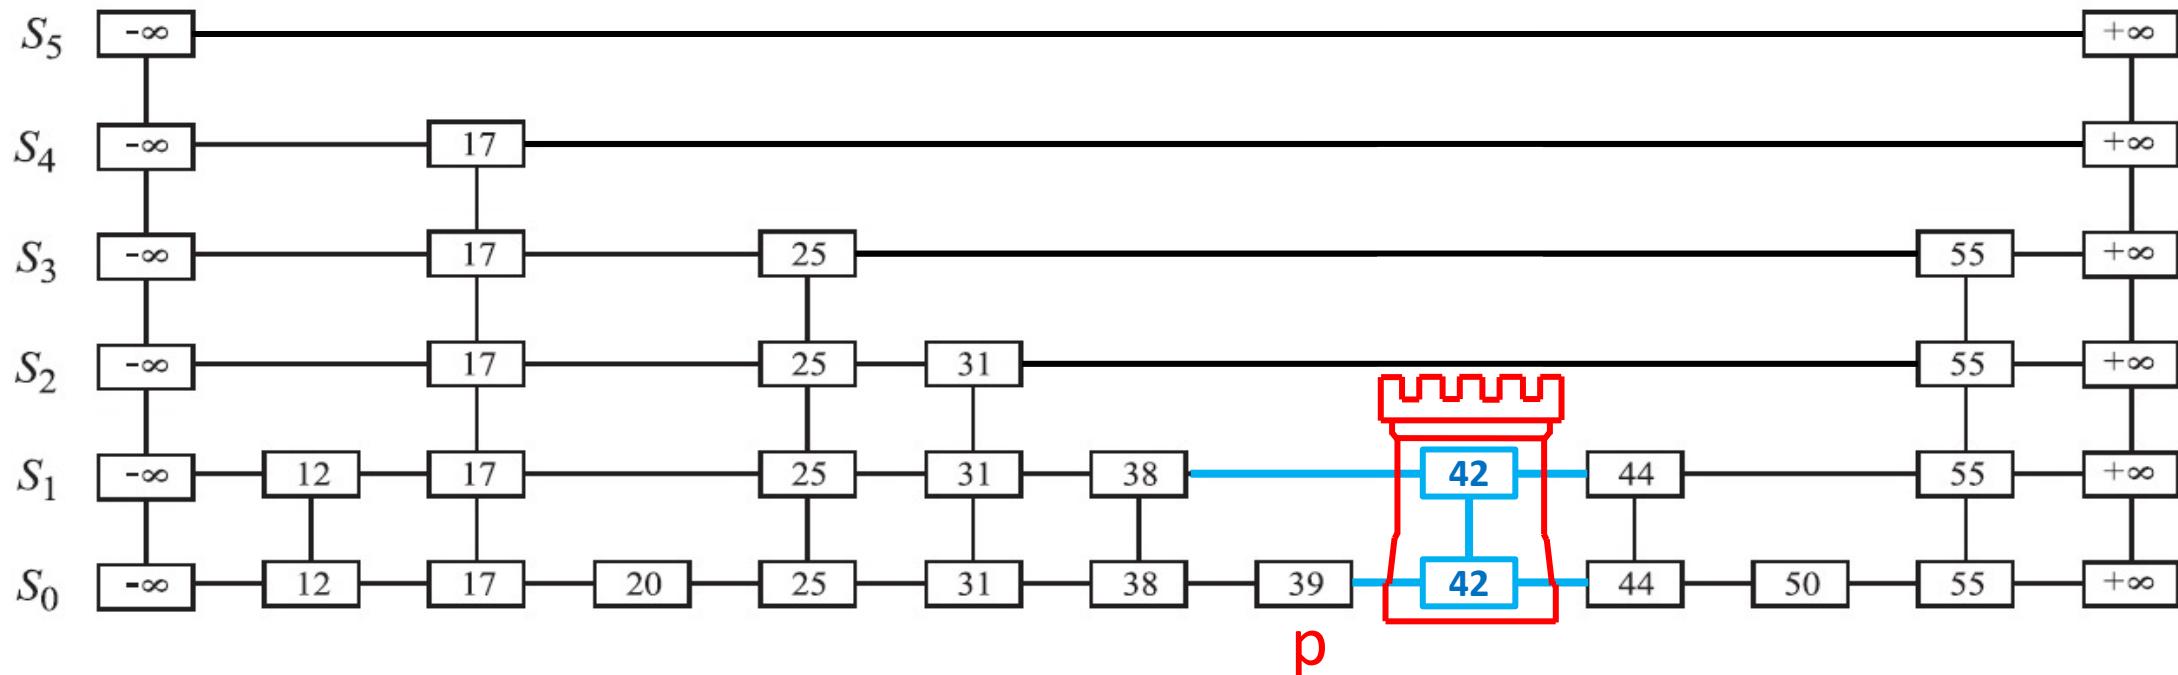
- **Search** for the largest key that is ≤ 42 , and set **p** to be its position
- **Insert** the entry right after **p** in the same level



SKIP LIST – INSERTION

Here is an example of how to **insert** an entry with key = **42**:

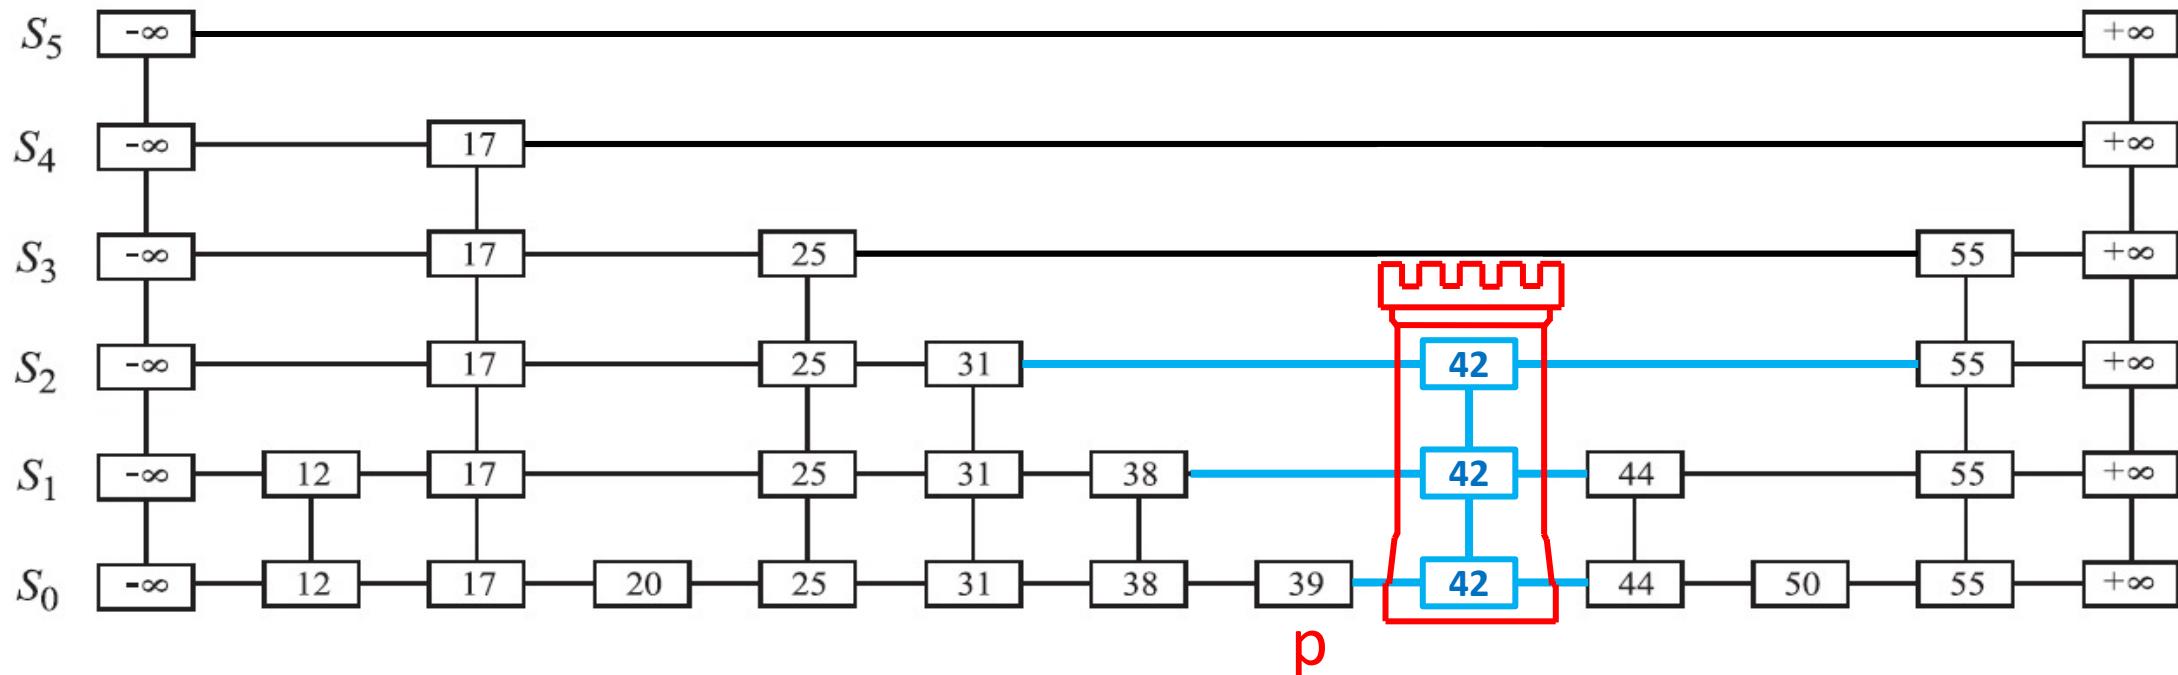
- **Search** for the largest key that is ≤ 42 , and set **p** to be its position
- **Insert** the entry right after **p** in the same level
- **Build a tower** at **p** by repeatedly flip a coin to decide whether to go one more level up!



SKIP LIST – INSERTION

Here is an example of how to **insert** an entry with key = **42**:

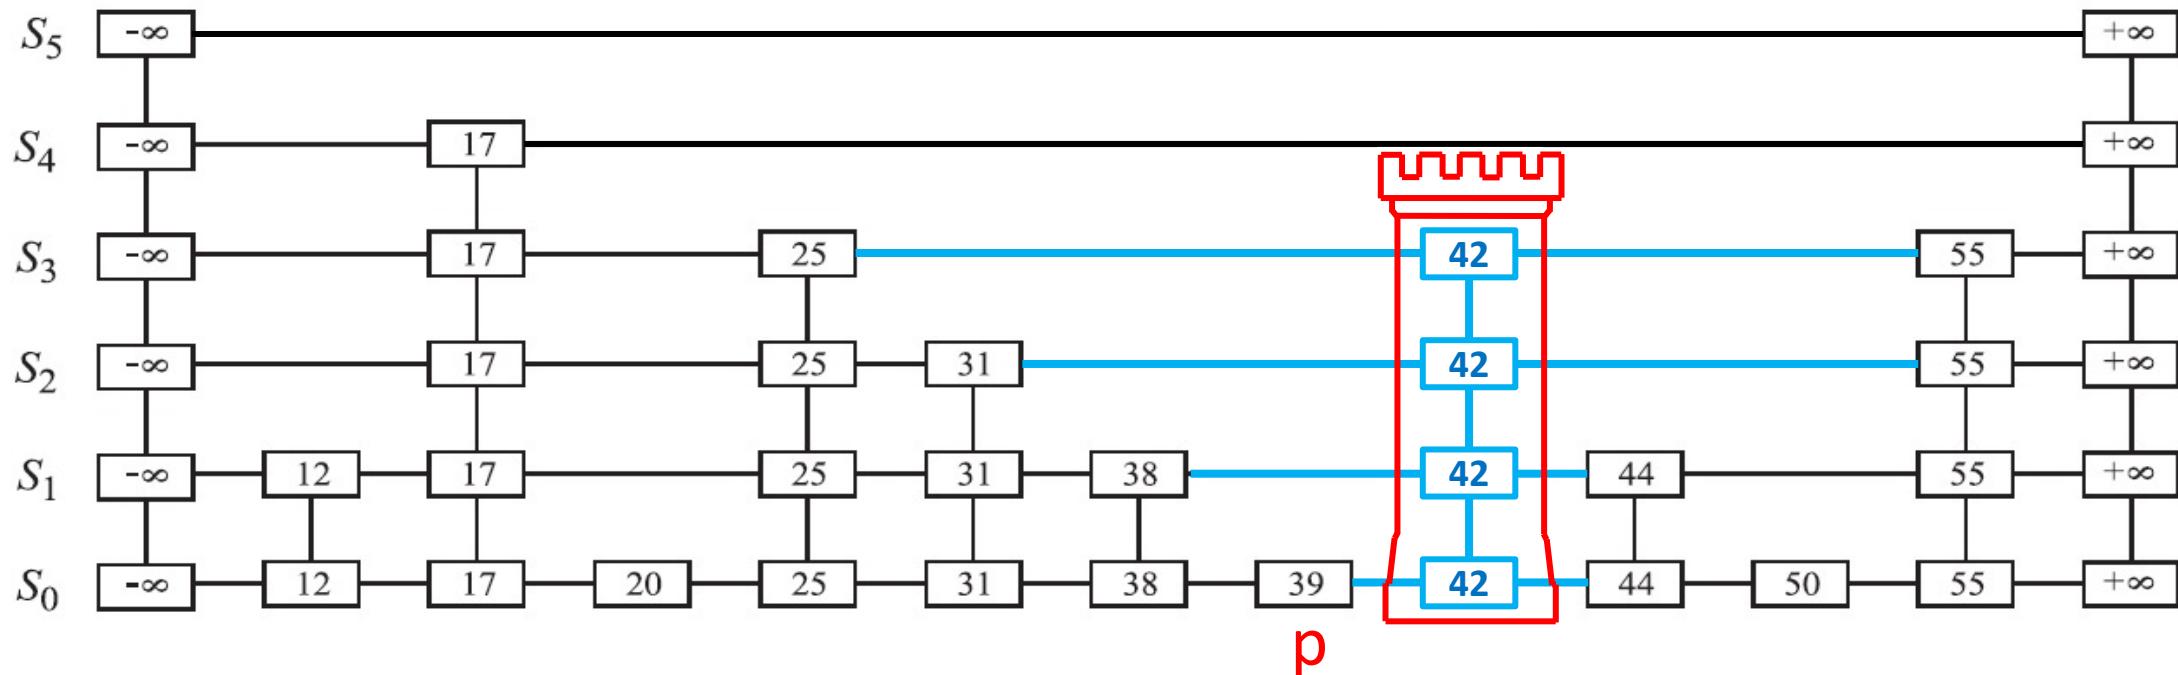
- **Search** for the largest key that is ≤ 42 , and set **p** to be its position
- **Insert** the entry right after **p** in the same level
- **Build a tower** at **p** by repeatedly flip a coin to decide whether to go one more level up!



SKIP LIST – INSERTION

Here is an example of how to **insert** an entry with key = **42**:

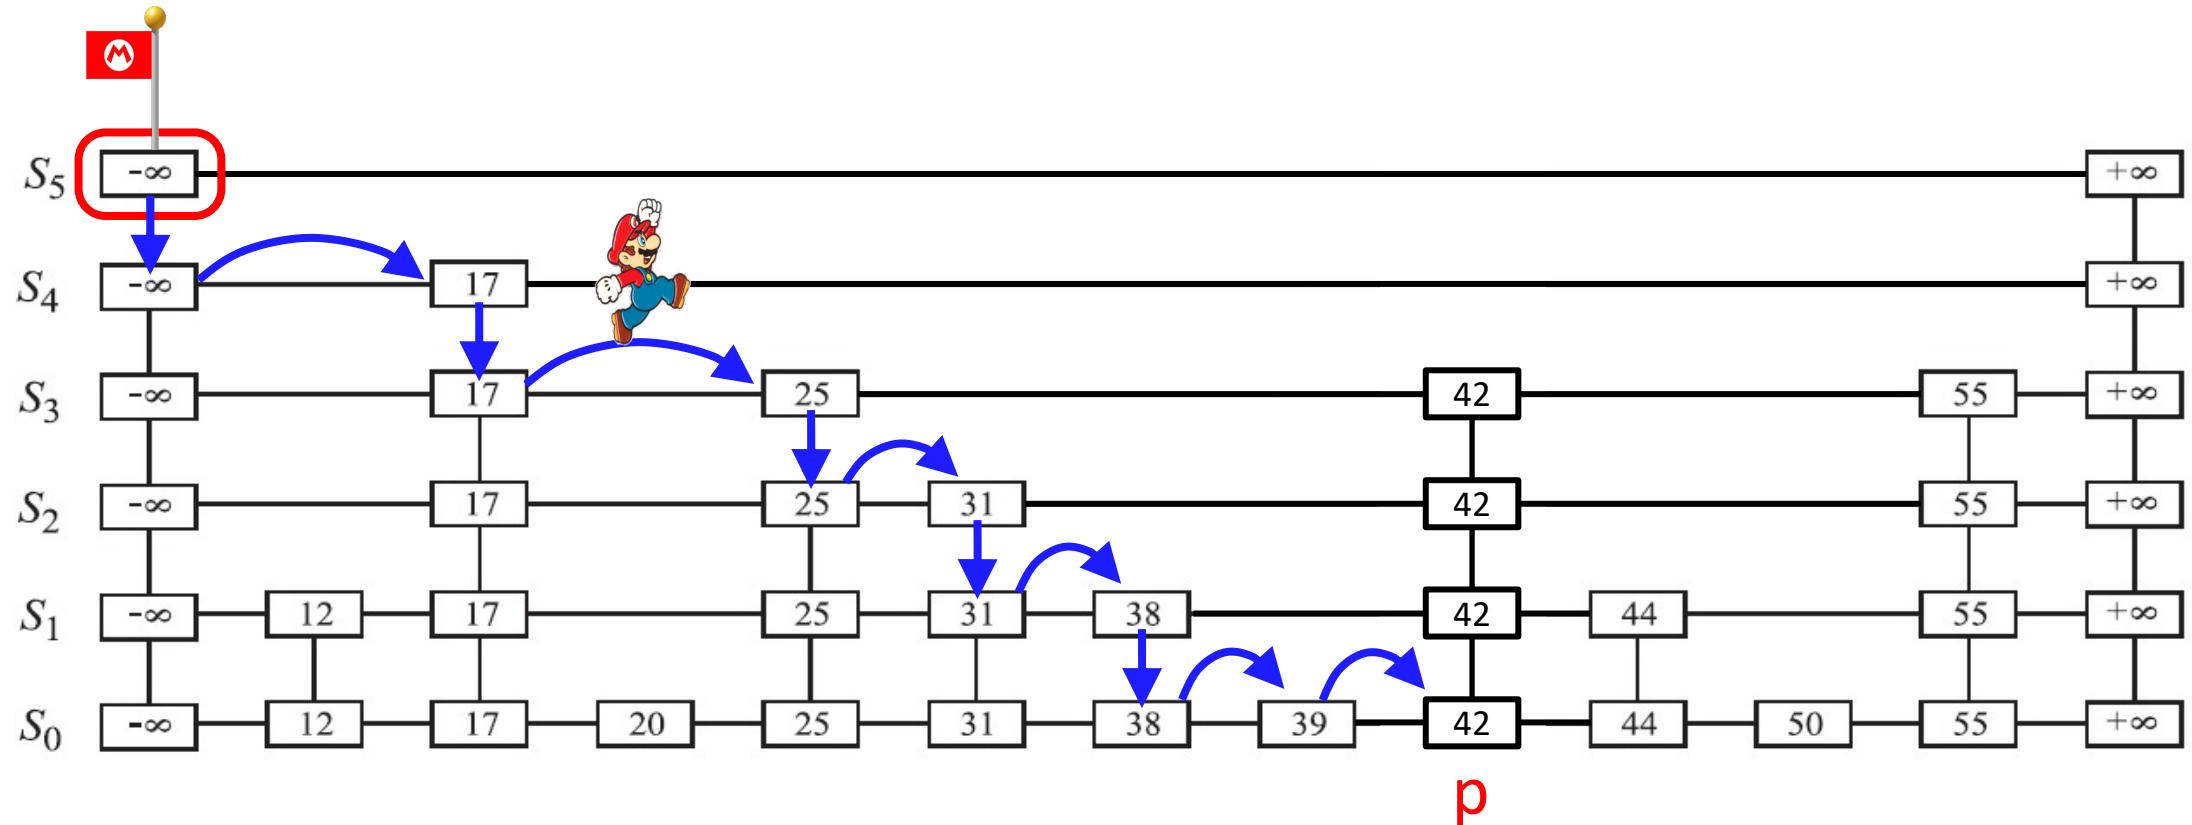
- **Search** for the largest key that is ≤ 42 , and set **p** to be its position
- **Insert** the entry right after **p** in the same level
- **Build a tower** at **p** by repeatedly flip a coin to decide whether to go one more level up!



SKIP LIST – REMOVAL

Here is an example of how to **remove** an entry with key = **42**:

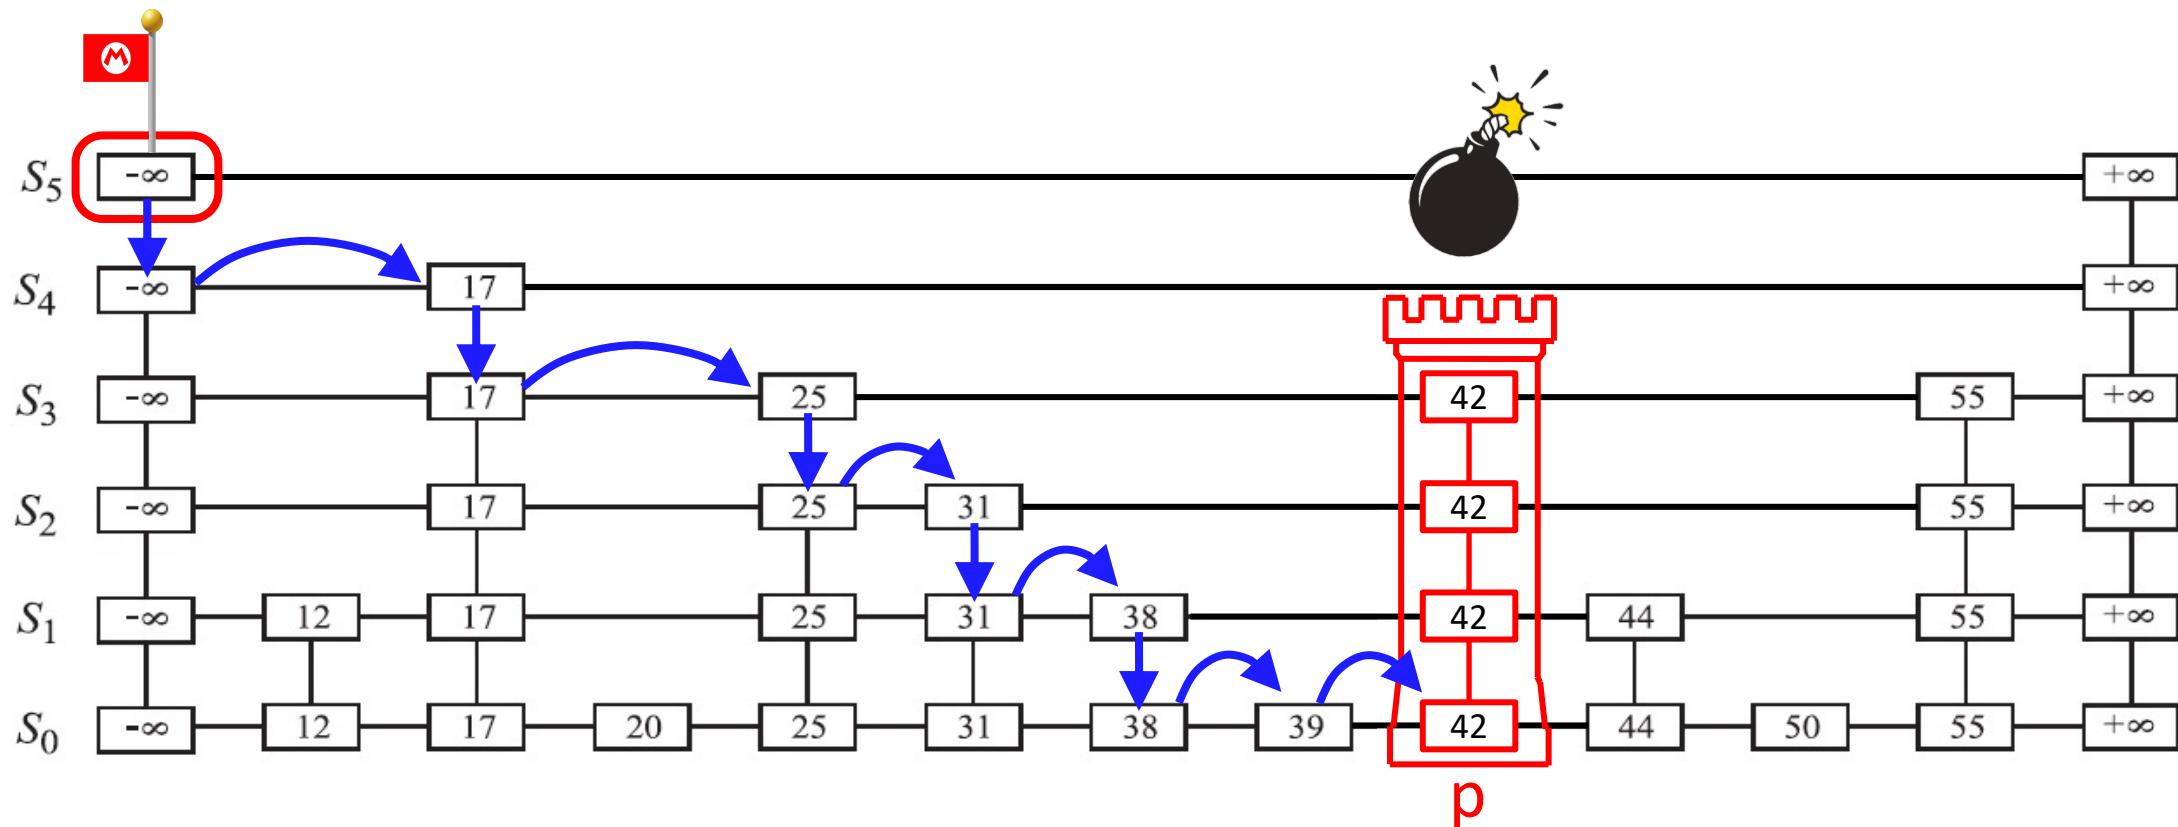
- **Search** for the entry whose key = **42**



SKIP LIST – REMOVAL

Here is an example of how to **remove** an entry with key = **42**:

- **Search** for the entry whose key = **42**
- **Remove** the entire tower at **p**

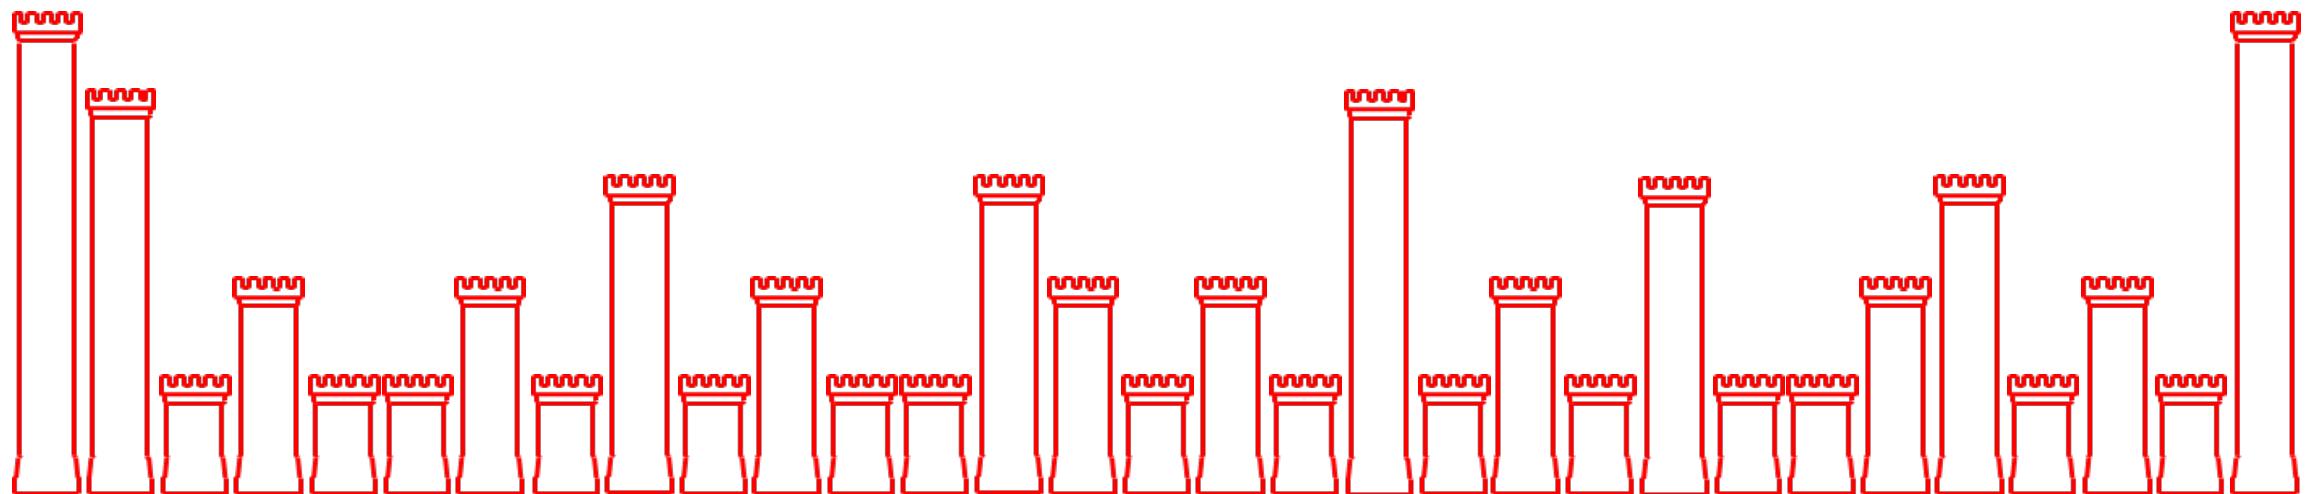
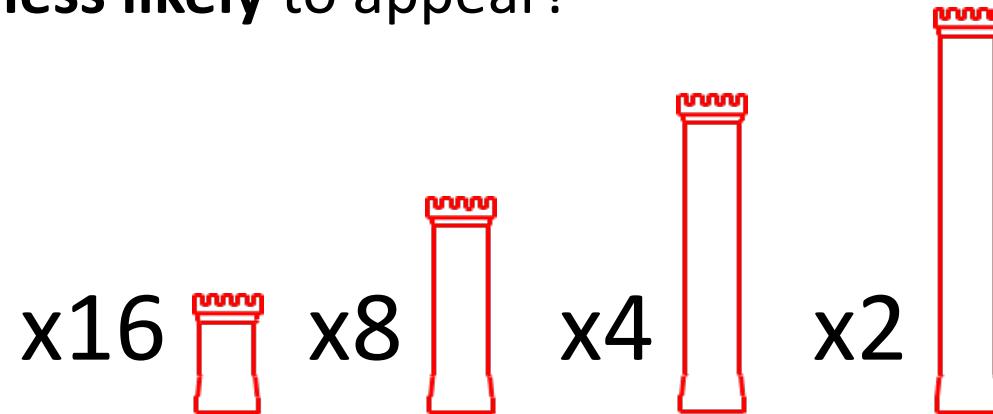


SKIP LISTS – SEARCH ANALYSIS

- Remember, searching for an entry whose key = k takes $O(\log n)$ with **binary search**, because **with each step we slash half** of the number of entries to consider
- But why would it also take $O(\log n)$ to search for an entry in a **skip list**? **Where is the part in which we discard lots of entries to consider?**

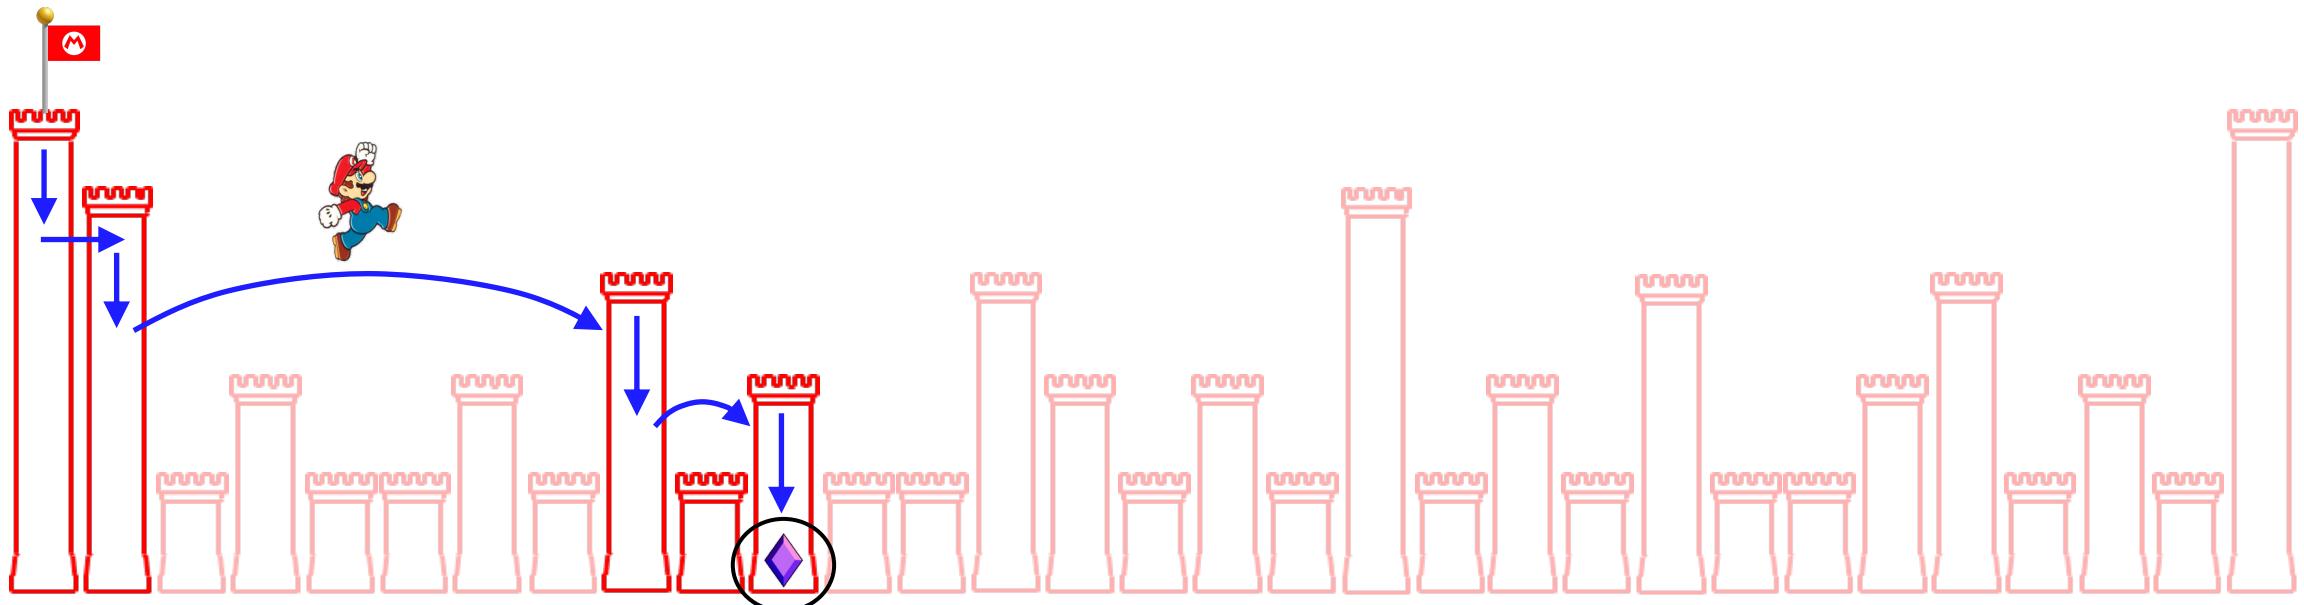
SKIP LISTS – SEARCH ANALYSIS

- A slightly higher tower is **twice less likely** to appear!



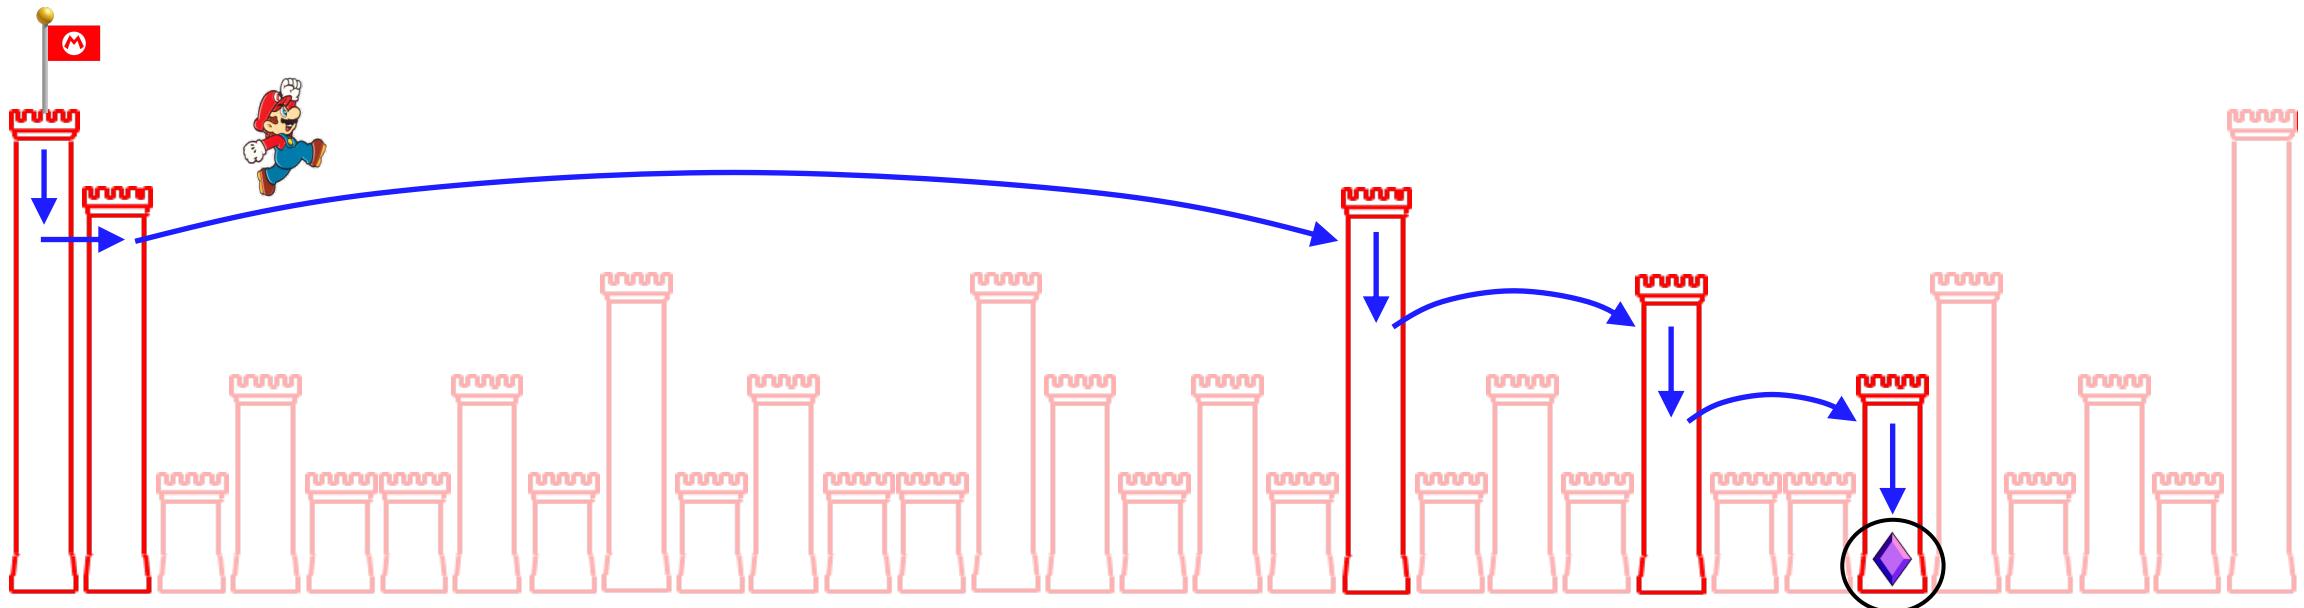
SKIP LISTS – SEARCH ANALYSIS

- A slightly higher tower is **twice less likely** to appear!
- Let us denote **the entry we're searching for** as: ♦
- If ♦ is **close to the start**, we'll find it quickly and disregard everything after it!



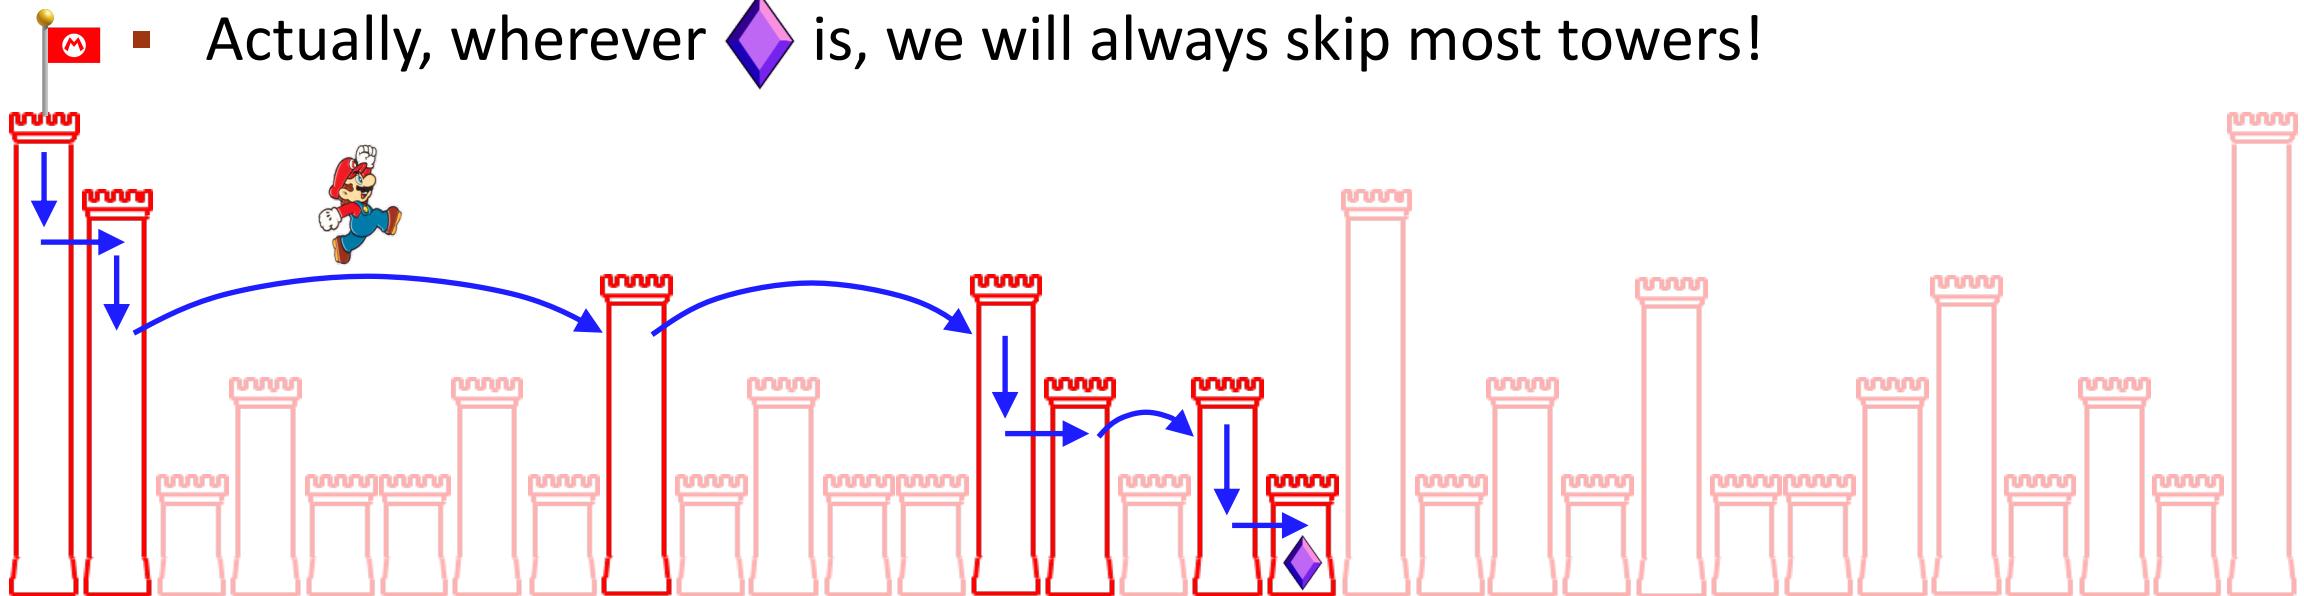
SKIP LISTS – SEARCH ANALYSIS

- A slightly higher tower is **twice less likely** to appear!
- Let us denote **the entry we're searching for** as: ♦
- If ♦ is **close to the start**, we'll find it **quickly** and disregard everything after it!
- If it is **far from the start**, we'll make **large jumps** across higher towers!



SKIP LISTS – SEARCH ANALYSIS

- A slightly higher tower is **twice less likely** to appear!
- Let us denote **the entry we're searching for** as: ♦
- If ♦ is **close to the start**, we'll find it **quickly** and disregard everything after it!
- If it is **far from the start**, we'll make **large jumps** across higher towers!
- Actually, wherever ♦ is, we will always skip most towers!



SKIP LISTS – BOUNDING THE HEIGHT

What is the probability that the height of a tower reaches **h**?

- If you flip a coin **once**, the probability of getting **a head** is: $\frac{1}{2}$

Possible outcomes:

- Outcome 1:

- Outcome 2:


SKIP LISTS – BOUNDING THE HEIGHT

What is the probability that the height of a tower reaches **h**?

- If you flip a coin **once**, the probability of getting **a head** is: $\frac{1}{2}$
- If you flip it **twice**, the probability of getting **two consecutive heads** is: $\frac{1}{2} \times \frac{1}{2} = \frac{1}{4}$

Possible outcomes:

➤ Outcome 1:



➤ Outcome 2:



➤ Outcome 3:



➤ Outcome 4:



SKIP LISTS – BOUNDING THE HEIGHT

What is the probability that the height of a tower reaches **h**?

- If you flip a coin **once**, the probability of getting **a head** is: $\frac{1}{2}$
- If you flip it **twice**, the probability of getting **two consecutive heads** is: $\frac{1}{2} \times \frac{1}{2}$
- If you flip it **three times**, the probability of **three consecutive heads** is: $\frac{1}{2} \times \frac{1}{2} \times \frac{1}{2} = \frac{1}{8}$

Possible outcomes:

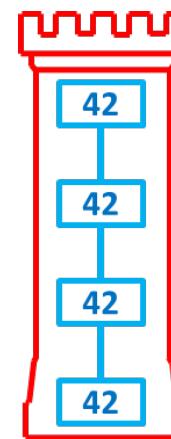
- Outcome 1: (H, H, H)
- Outcome 2: (T, T, T) **Outcome 2 is circled in red.**
- Outcome 3: (H, T, H)
- Outcome 4: (T, H, T)

- Outcome 5: (H, H, T)
- Outcome 6: (T, T, H)
- Outcome 7: (H, T, T)
- Outcome 8: (T, H, H)

SKIP LISTS – BOUNDING THE HEIGHT

What is the probability that the height of a tower reaches h ?

- If you flip a coin **once**, the probability of getting **a head** is: $\frac{1}{2}$
- If you flip it **twice**, the probability of getting **two consecutive heads** is: $\frac{1}{2} \times \frac{1}{2}$
- If you flip it **three times**, the probability of **three consecutive heads** is: $\frac{1}{2} \times \frac{1}{2} \times \frac{1}{2} = \frac{1}{8}$
- If you flip it i times, the probability of i consecutive heads is: $\frac{1}{2^i}$
- Thus, for any given tower, the probability that its height reaches i is : $\frac{1}{2^i}$



SKIP LISTS – BOUNDING THE HEIGHT

- We found that, for any tower, the probability that its height reaches i is $\frac{1}{2^i}$
- Also, from Probability Theory, we know that the **probability that any one of n different events occurs** is **at most the sum** of the probabilities that each occurs.
- Hence, the **probability that at least one of n different towers reaches** a height of i is:

$$p_i \leq \frac{n}{2^i}$$

- Thus, the probability that the height of the entire skip list, h , reaches, say, $3 \log n$ is:

$$p_{3 \log n} \leq \frac{n}{2^{3 \log n}} = \frac{n}{n^3} = \frac{1}{n^2}$$

- For example, if $n=10,000$, the probability that h reaches $3 \log n$ is just **1 in 100,000,000**
- Therefore, given a constant c , the probability that $h \geq c \log n$ is at most $\frac{1}{n^{c-1}}$
- Thus, with high probability, the height h of the skip list is $O(\log n)$

HASH TABLE VS. SEARCH TABLE VS. SKIP LIST

Method	Hash table	Search Table	Skip List
find	$O(1)$ expected $O(n)$ worst case	$O(\log n)$	$O(\log n)$ expected
insert	$O(1)$	$O(n)$	$O(\log n)$ expected
erase	$O(1)$ expected $O(n)$ worst case	$O(n)$	$O(\log n)$ expected

71

DICTIONARIES

DICTIONARY

- A **dictionary** is similar to a map, except that **multiple entries can have the same key**.

size(): Return the **number of entries** in D .

empty(): Return **true if D is empty** and false otherwise.

begin(): Return an iterator to the **first entry of D** .

end(): Return an iterator to a **position just beyond the end of D** .

insert(k, v): Insert entry **(k, v)** into D , returning an iterator referring to the newly created entry.

erase(p): Remove from D the entry referenced by iterator p ; ERROR if p points to the end sentinel.

erase(k): Remove from D an arbitrary entry whose **key = k** ; ERROR if D has no such entry.

find(k): Return **an iterator to an entry whose key = k** ; unless no such entry exists in D , then return the special iterator `end`.

findAll(k): Return a pair of iterators (b, e) , such that **all the entries whose key = k** lie in the range from b up to, but not including, e .

DICTIONARY – EXAMPLE

Example:

- In the **Output** column, we use the notation:

$$p_i : [(k, v)]$$

to mean that the operation returns an iterator p_i that refers to the entry (k, v)

Operation	Output	Dictionary
empty()	true	\emptyset
insert(5,A)	$p_1 : [(5, A)]$	$\{(5, A)\}$
insert(7,B)	$p_2 : [(7, B)]$	$\{(5, A), (7, B)\}$
insert(2,C)	$p_3 : [(2, C)]$	$\{(5, A), (7, B), (2, C)\}$
insert(8,D)	$p_4 : [(8, D)]$	$\{(5, A), (7, B), (2, C), (8, D)\}$
insert(2,E)	$p_5 : [(2, E)]$	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
find(7)	$p_2 : [(7, B)]$	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
find(4)	end	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
find(2)	$p_3 : [(2, C)]$	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
findAll(2)	(p_3, p_4)	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
size()	5	$\{(5, A), (7, B), (2, C), (2, E), (8, D)\}$
erase(5)	–	$\{(7, B), (2, C), (2, E), (8, D)\}$
erase(p_3)	–	$\{(7, B), (2, E), (8, D)\}$
find(2)	$p_5 : [(2, E)]$	$\{(7, B), (2, E), (8, D)\}$

CHAPTER 10: SEARCH TREES

75

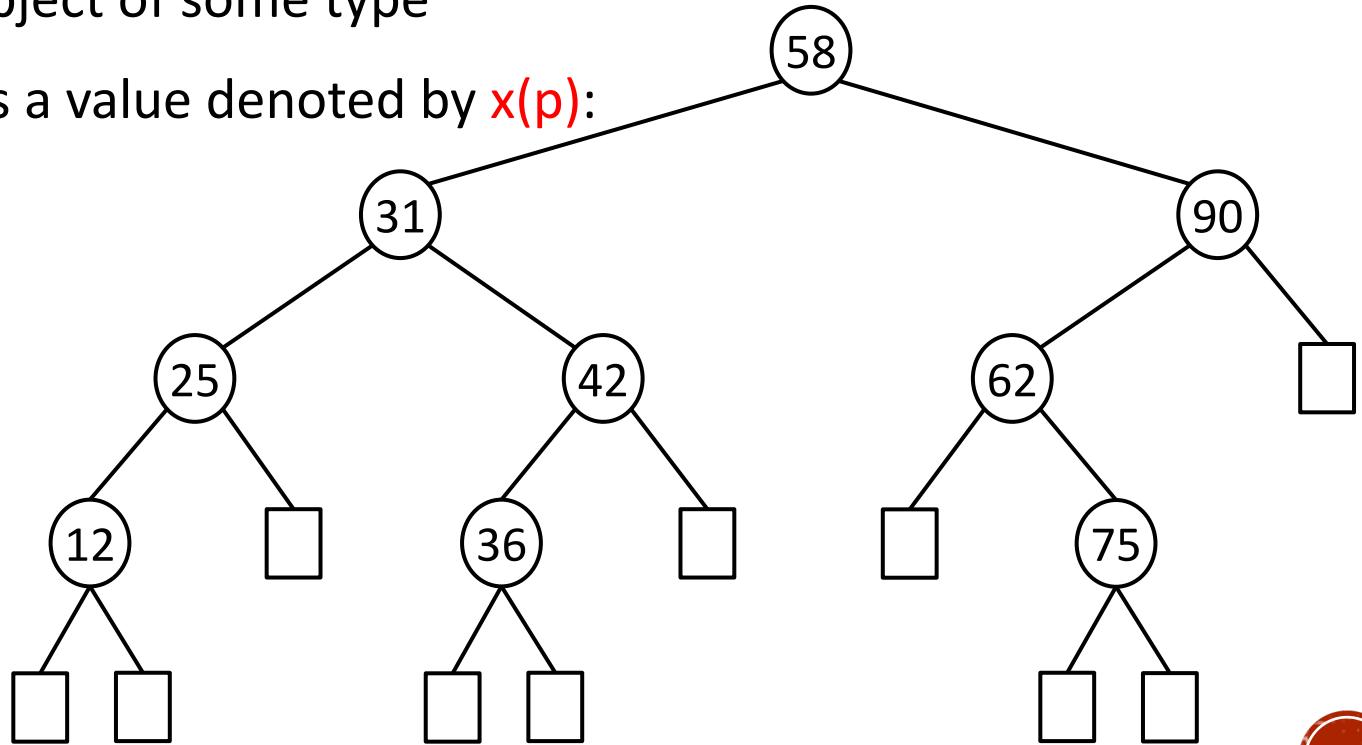
BINARY SEARCH TREES

BINARY SEARCH TREE

An excellent data structure for *storing the entries of a map or a dictionary* is a **binary search tree**, which is defined as a **proper binary tree** where:

- **External nodes** (represented as squares) do not store elements
- Each **internal node p** stores an object of some type
- For each **internal node p**, there is a value denoted by $x(p)$:

- Any object stored in the **left subtree** of p has such a value that is $\leq x(p)$
- Any object stored in the **right subtree** of p has such a value that is $\geq x(p)$

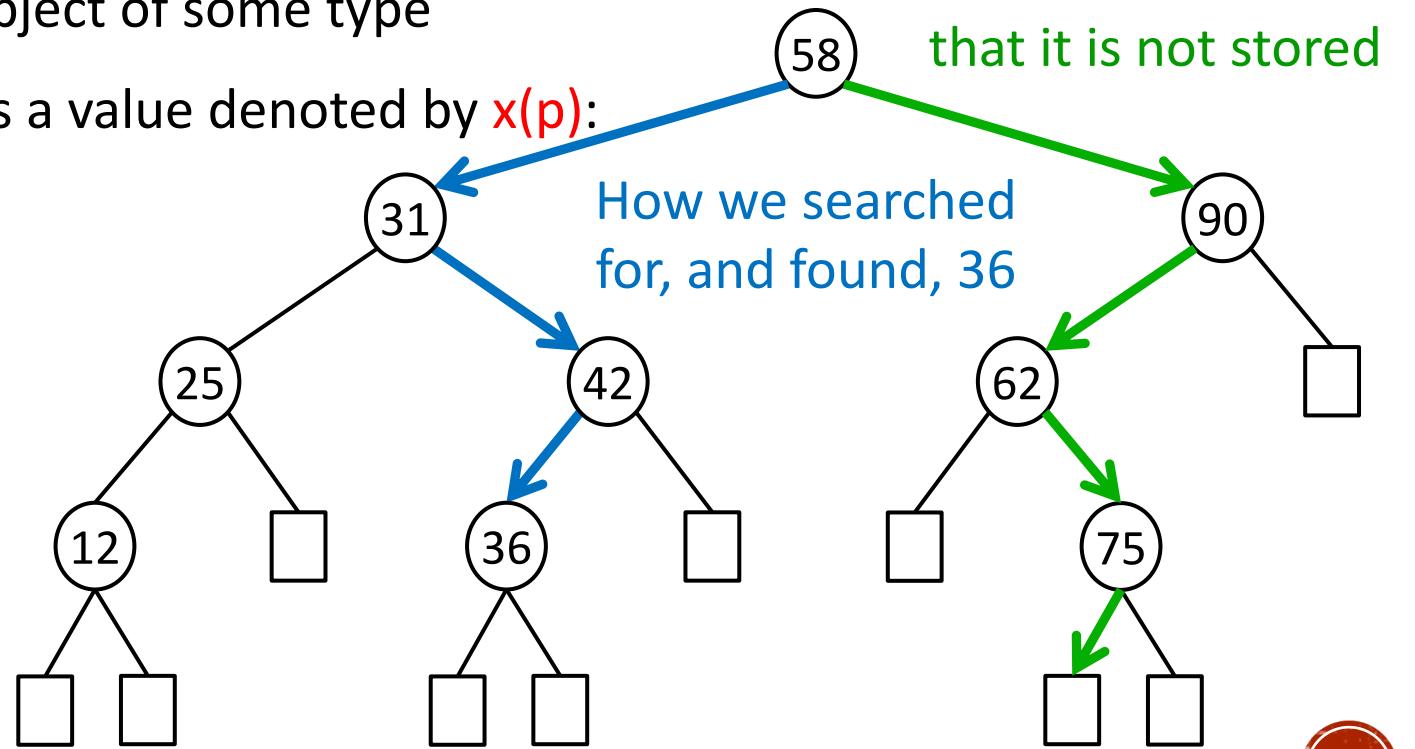


BINARY SEARCH TREE

An excellent data structure for *storing the entries of a map or a dictionary* is a **binary search tree**, which is defined as a **proper binary tree** where:

- **External nodes** (represented as squares) do not store elements
- Each **internal node p** stores an object of some type
- For each **internal node p**, there is a value denoted by $x(p)$:

- Any object stored in the **left subtree** of p has such a value that is $\leq x(p)$
- Any object stored in the **right subtree** of p has such a value that is $\geq x(p)$



BINARY SEARCH TREE

Here is a **recursive algorithm**:

- **TreeSearch(k, v)**: searches for a key k in the subtree rooted at node v
- To search for k in the entire tree, T , simply call **TreeSearch($k, T.root()$)**

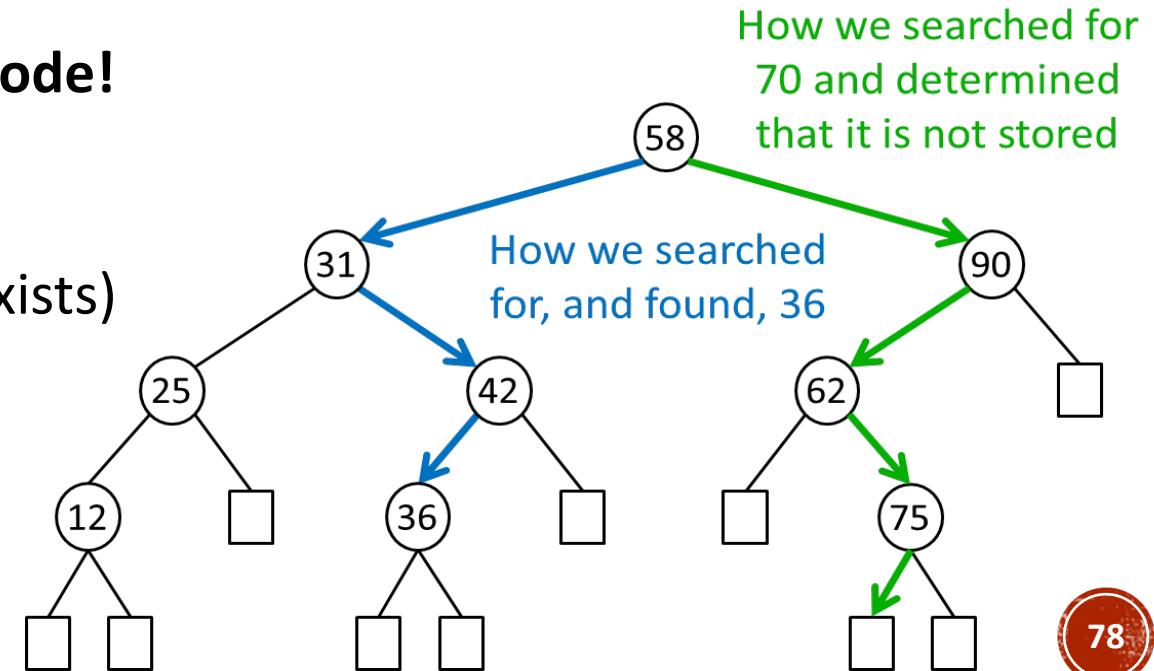
Note that this algorithm always returns a node!

Specifically, it return one of the following:

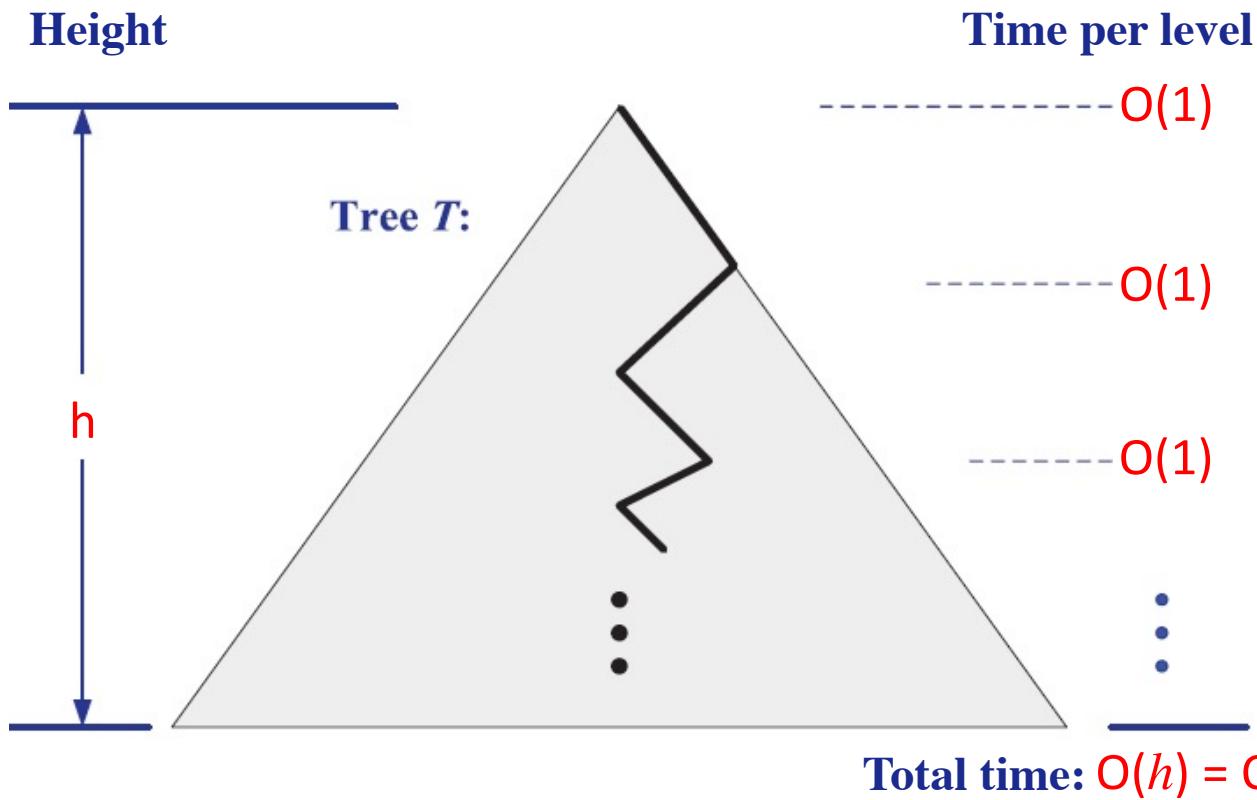
1. A nodes whose key = k (if such a node exists)
2. An external node where it would be suitable to **insert a node whose key = k**

This second output will be handy later on when implementing “**insert**”

```
Algorithm TreeSearch(  $k, v$  ):  
    if  $T.isExternal(v)$  then  
        return  $v$  {return external node}  
    if  $k < v.key()$  then  
        return TreeSearch(  $k, T.left(v)$  )  
    else if  $k > v.key()$  then  
        return TreeSearch(  $k, T.right(v)$  )  
    return  $v$  { $v.key()$  must be equal to  $k$ }
```



What is the complexity of this algorithm?



Since we spend $O(1)$ time per node encountered in the search, the function `find(k)` on a map runs in $O(h)$ time, where h is the height of the binary search tree

```
Algorithm TreeSearch( k, v ):  
    if  $T.\text{isExternal}(v)$  then  
        return  $v$  {return external node}  
    if  $k < v.\text{key}()$  then  
        return TreeSearch(  $k$ ,  $T.\text{left}(v)$  )  
    else if  $k > v.\text{key}()$  then  
        return TreeSearch(  $k$ ,  $T.\text{right}(v)$  )  
    return  $v$  { $v.\text{key}()$  must be equal to  $k$ }
```

