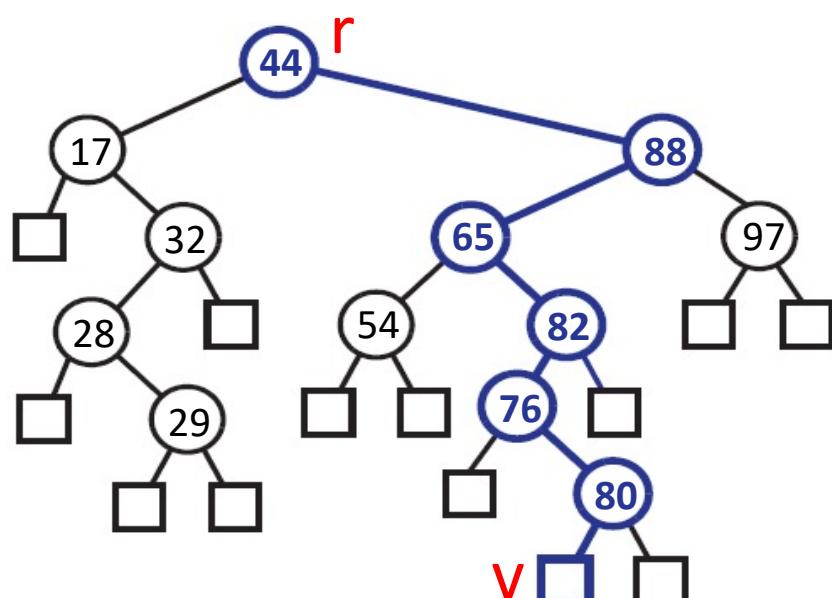
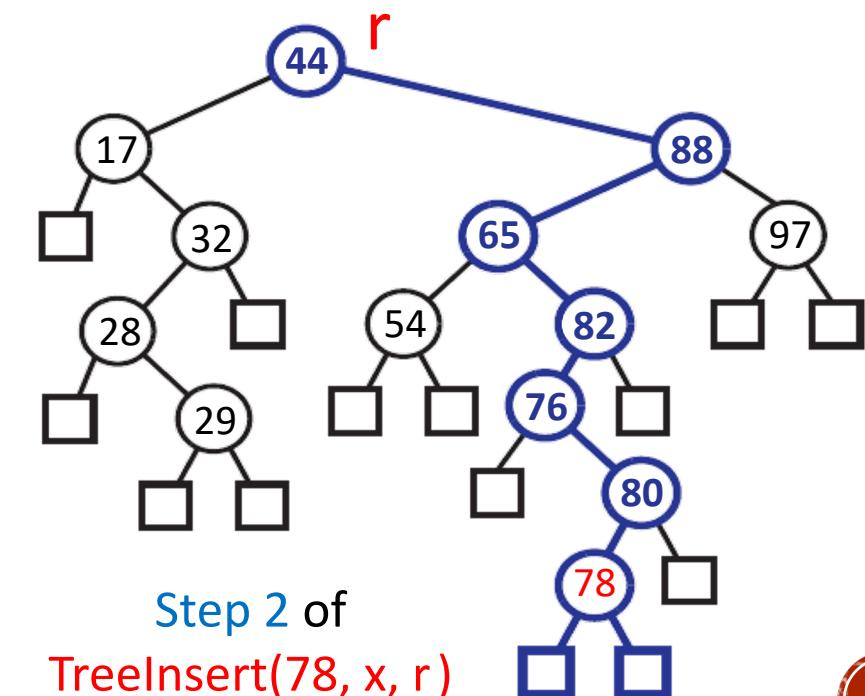
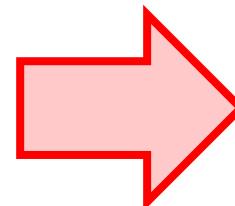


BINARY SEARCH TREE – INSERTION

- Consider this operation: `insertAtExternal(v, (k,x))`, which **inserts** entry **(k,x)** at the **external** node **v**, and **expand** **v** to become **internal**, having two external children
- This operation is used in `Treelnsert(k,x,r)` which works in two steps:
Step 1. **find** the position **v** to insert entry **(k,x)** Step 2. `insertAtExternal(v, (k,x))`



Step 1 of
`Treelnsert(78, x, r)`

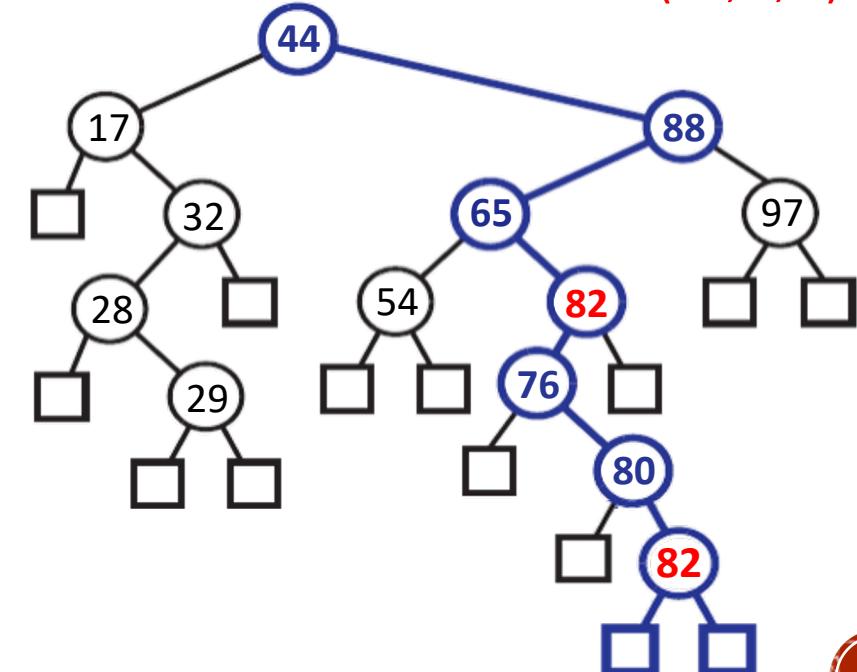


Step 2 of
`Treelnsert(78, x, r)`

BINARY SEARCH TREE – INSERTION

- **Important:** A binary search tree can be used to implement **maps** or **dictionaries**.
- If it will implement a **dictionary**, it should be able to hold **multiple entries** that **have the same key!**
- We said that for any node **r** containing an entry **(k,x)**:
 - Entries in the **left subtree** of **r** have keys $\leq k$
 - Entries in the **right subtree** of **r** have keys $\geq k$
- Where do we insert entries whose key $= k$?
 - We either insert them **always into the left subtree**, or **always into the right subtree**.
- Let's choose the **left subtree**, thus:
 - Entries in the **right subtree** of **r** have keys $> k$

Example:
TreeInsert(82, x, r)

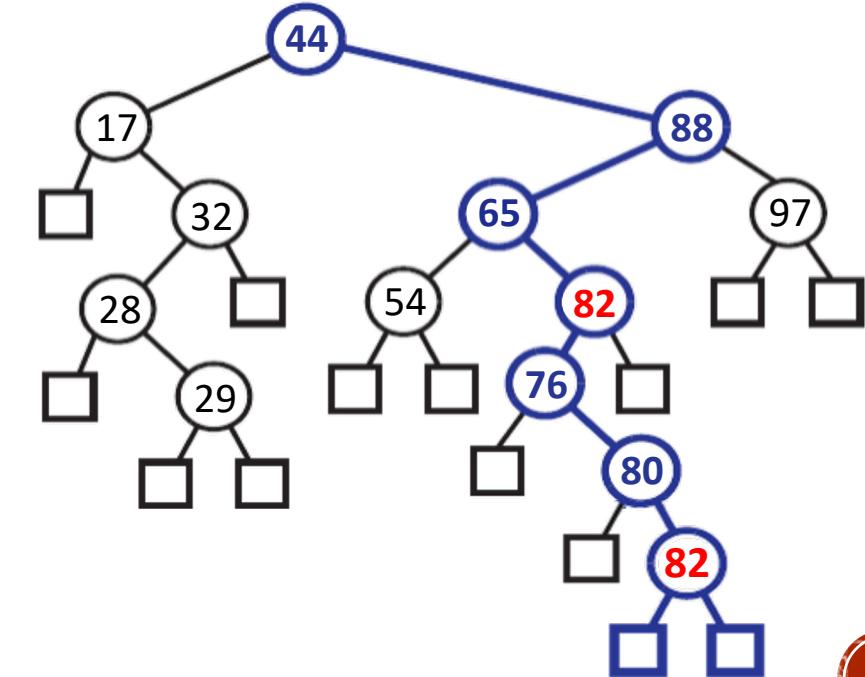


BINARY SEARCH TREE – INSERTION

- `Treelnsert(k, x, r)` is a recursive algorithm that inserts (k, x) in the **subtree rooted at node r**
- `Treelnsert($k, x, T.root()$)` inserts (k, x) in a tree T
- Remember: `TreeSearch(k, r)` returns either:
 - an **internal** node whose key = k , or
 - an **external** node where we can insert (k, x)
- The if-statement checks if w is **internal**. If so, the tree already contains at least one entry whose key = k , in which case we recursively search the subtree rooted at the **left child** of w
- Note: **Termination** (and **insertion**) only happens at an **external node!**

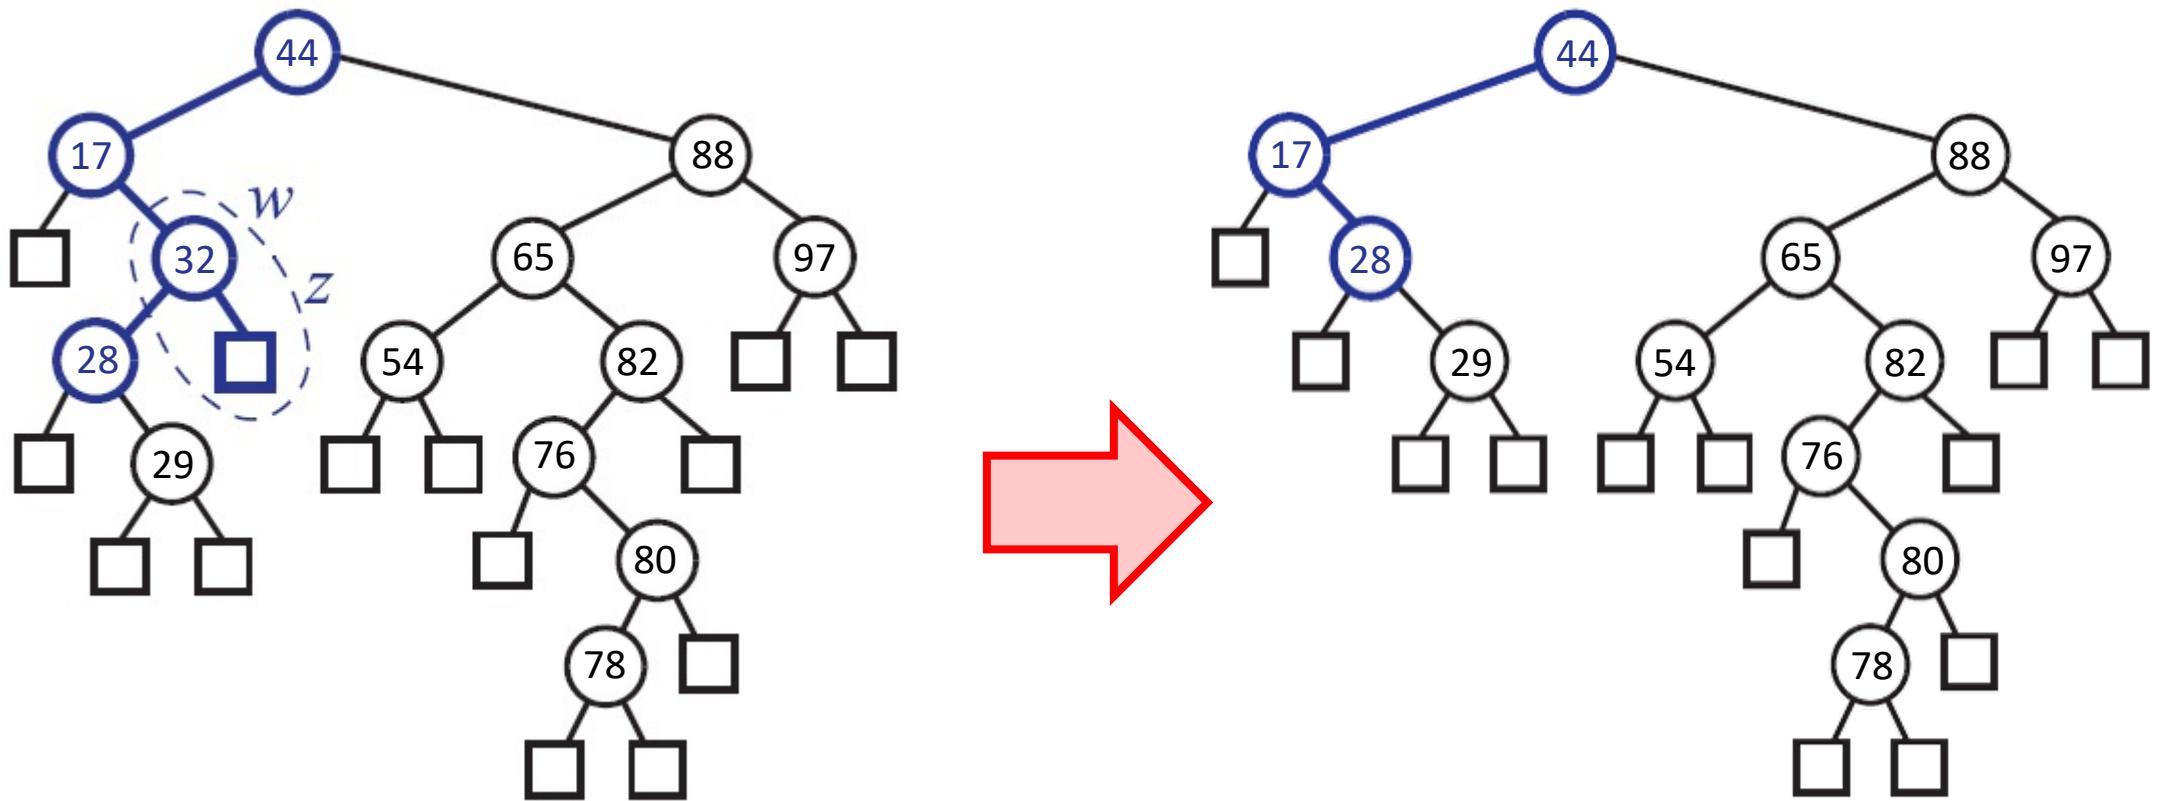
Algorithm `Treelnsert(k, x, r):`

```
 $w \leftarrow \text{TreeSearch}(k, r)$ 
if  $T.\text{isInternal}(w)$  then
    return Treelnsert( $k, x, T.\text{left}(w)$ )
 $T.\text{insertAtExternal}(w, (k, x))$ 
return  $w$ 
```



BINARY SEARCH TREE – REMOVAL

- When **removing nodes**, we'll use the following operation:
 - `RemoveAboveExternal(z)`: Removes the **external node z** along with its **parent**, but after replacing that parent with **z's sibling!**

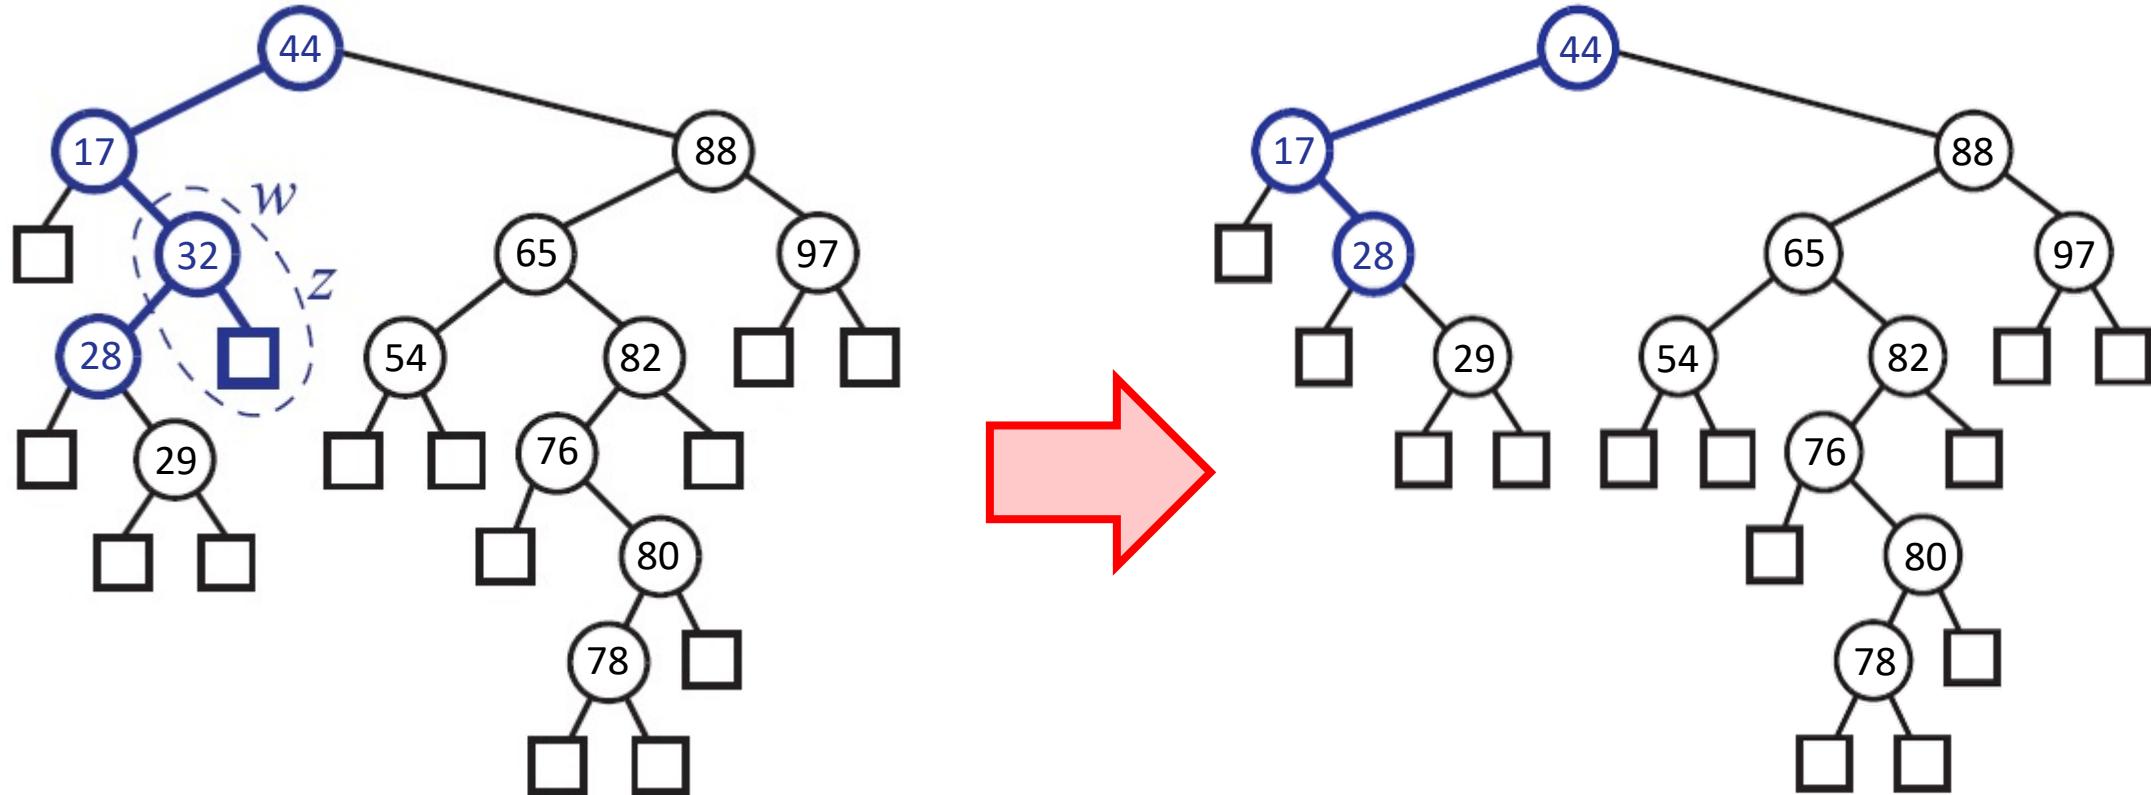


BINARY SEARCH TREE – REMOVAL

- To **remove** a node, **we first search for its key, k , using $\text{TreeSearch}(k, T.\text{root}())$** , which returns one of the following:
 1. An **external node**, implying that the key was **not found**.
 2. An **internal node**, w , **that has a child that is a leaf**
 3. An **internal node**, w , **whose children are both internal**
- In **Case 1**, there is no need to erase anything, and an error is signaled.

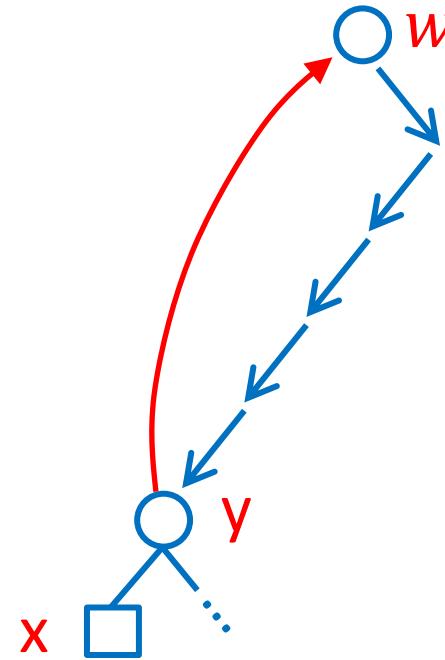
BINARY SEARCH TREE – REMOVAL

- Case 2: `TreeSearch(k , $T.\text{root}()$)` returns an **internal node w** that has a child that is a **leaf**
 - To remove w , we call `removeAboveExternal(z)` where z is an external child of w



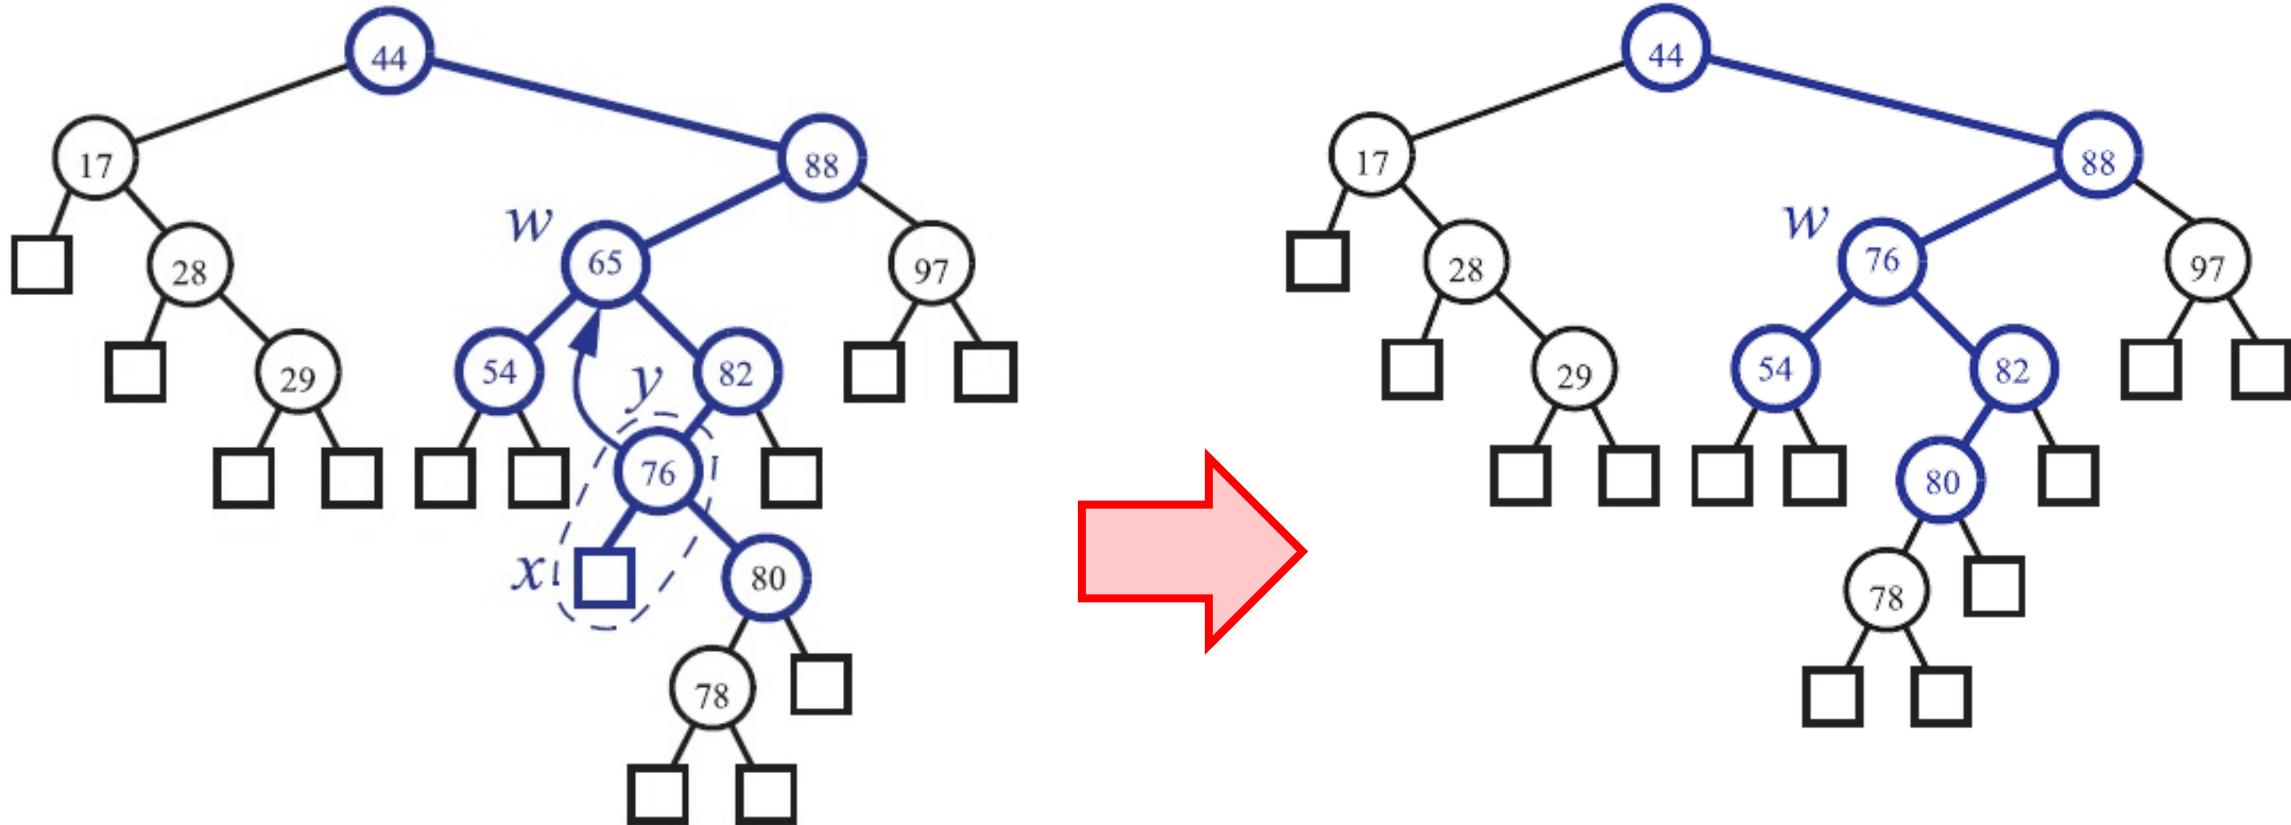
BINARY SEARCH TREE – REMOVAL

- **Case 3:** `TreeSearch(k , $T.root()$)` returns an **internal node w** whose children are **internal**
 - To remove w , we may follow three steps:
 1. Find the last **internal** node y that comes after w by going to the **right child** of w and then move down following the **left children**
 2. Move the entry of y into w (which deletes what was stored in w)
 3. Call `removeAboveExternal(x)` where x is the **left child** of y ; this will remove both x and y



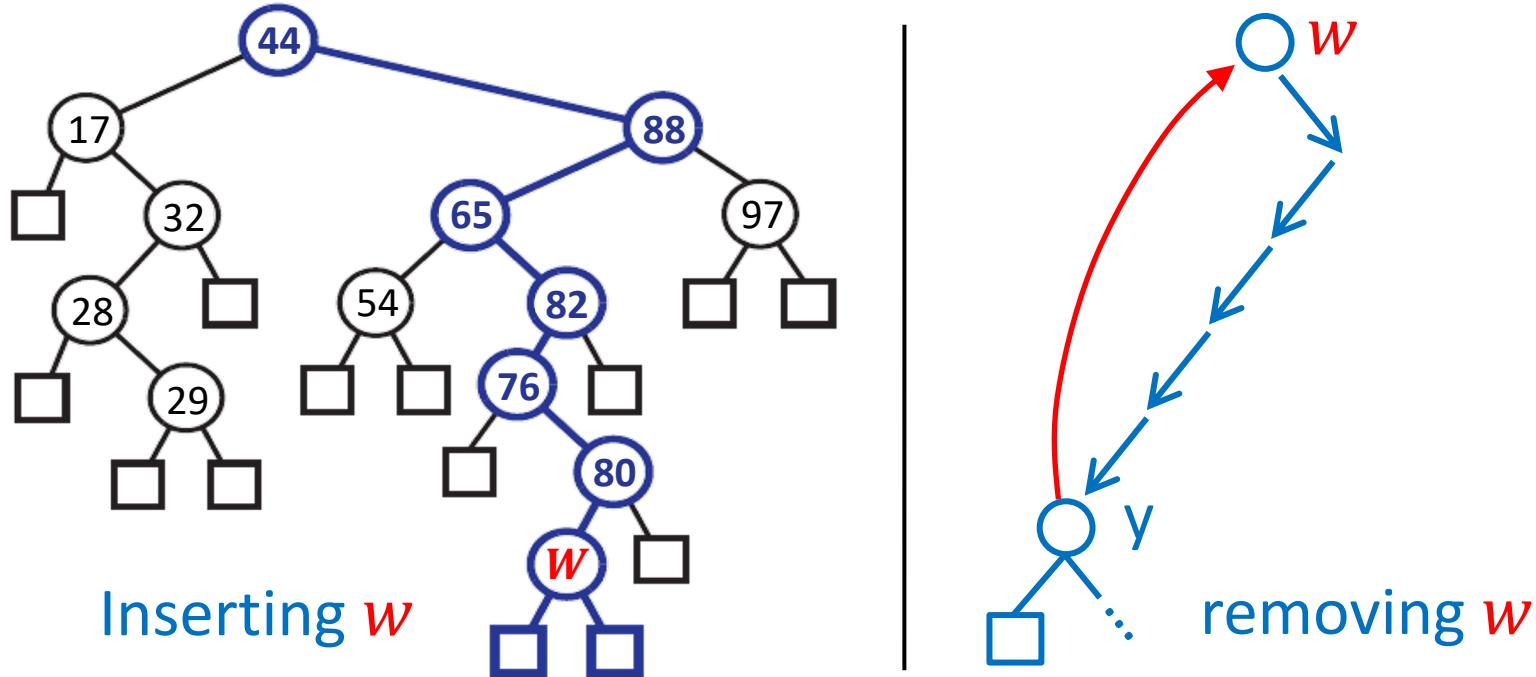
BINARY SEARCH TREE – REMOVAL

- Case 3: `TreeSearch(k , $T.root()$)` returns an **internal node w** whose children are **internal**



BINARY SEARCH TREE – PERFORMANCE

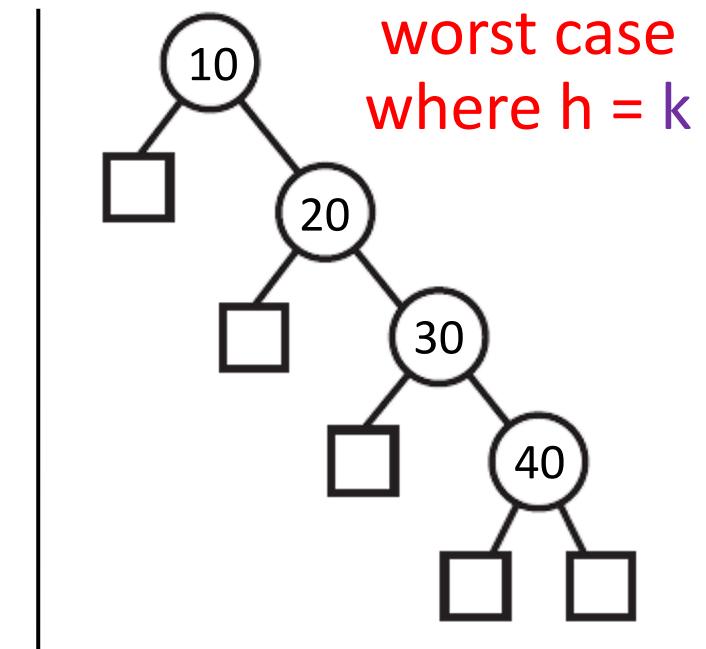
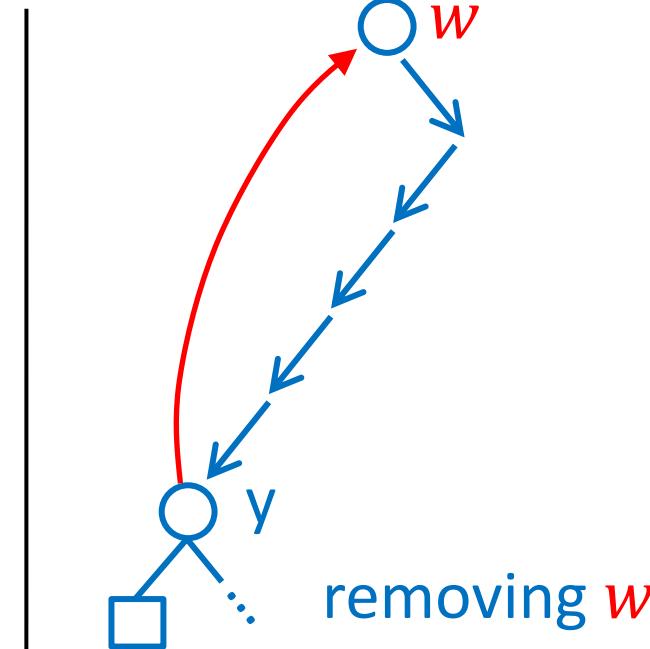
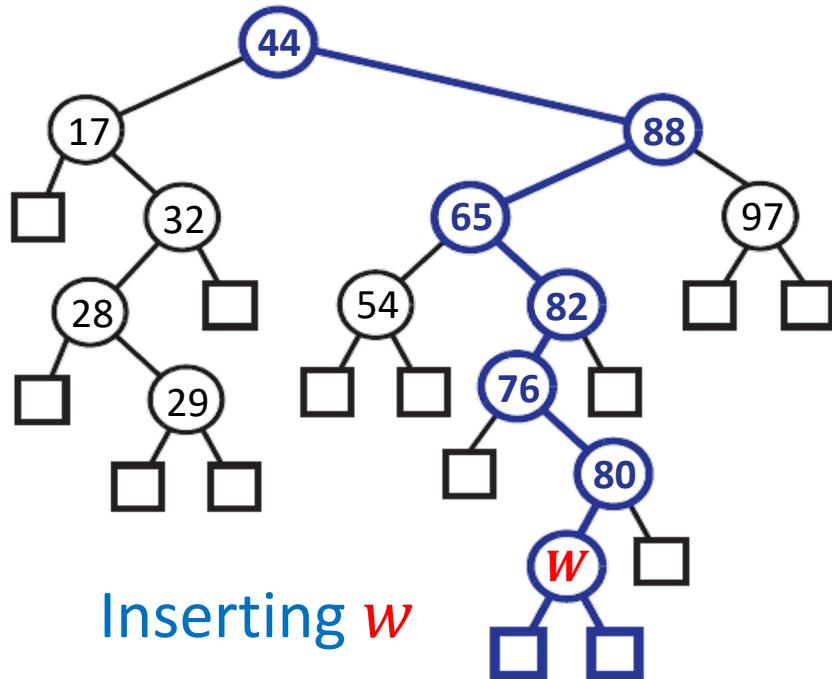
- Insertion and removal take $O(h) = O(\log n)$ time, where h is the **height** of the tree.



- Given k internal nodes, what is the maximum height of a binary search tree, h ?

BINARY SEARCH TREE – PERFORMANCE

- Insertion and removal take $O(h)$ time, where h is the **height** of the tree.

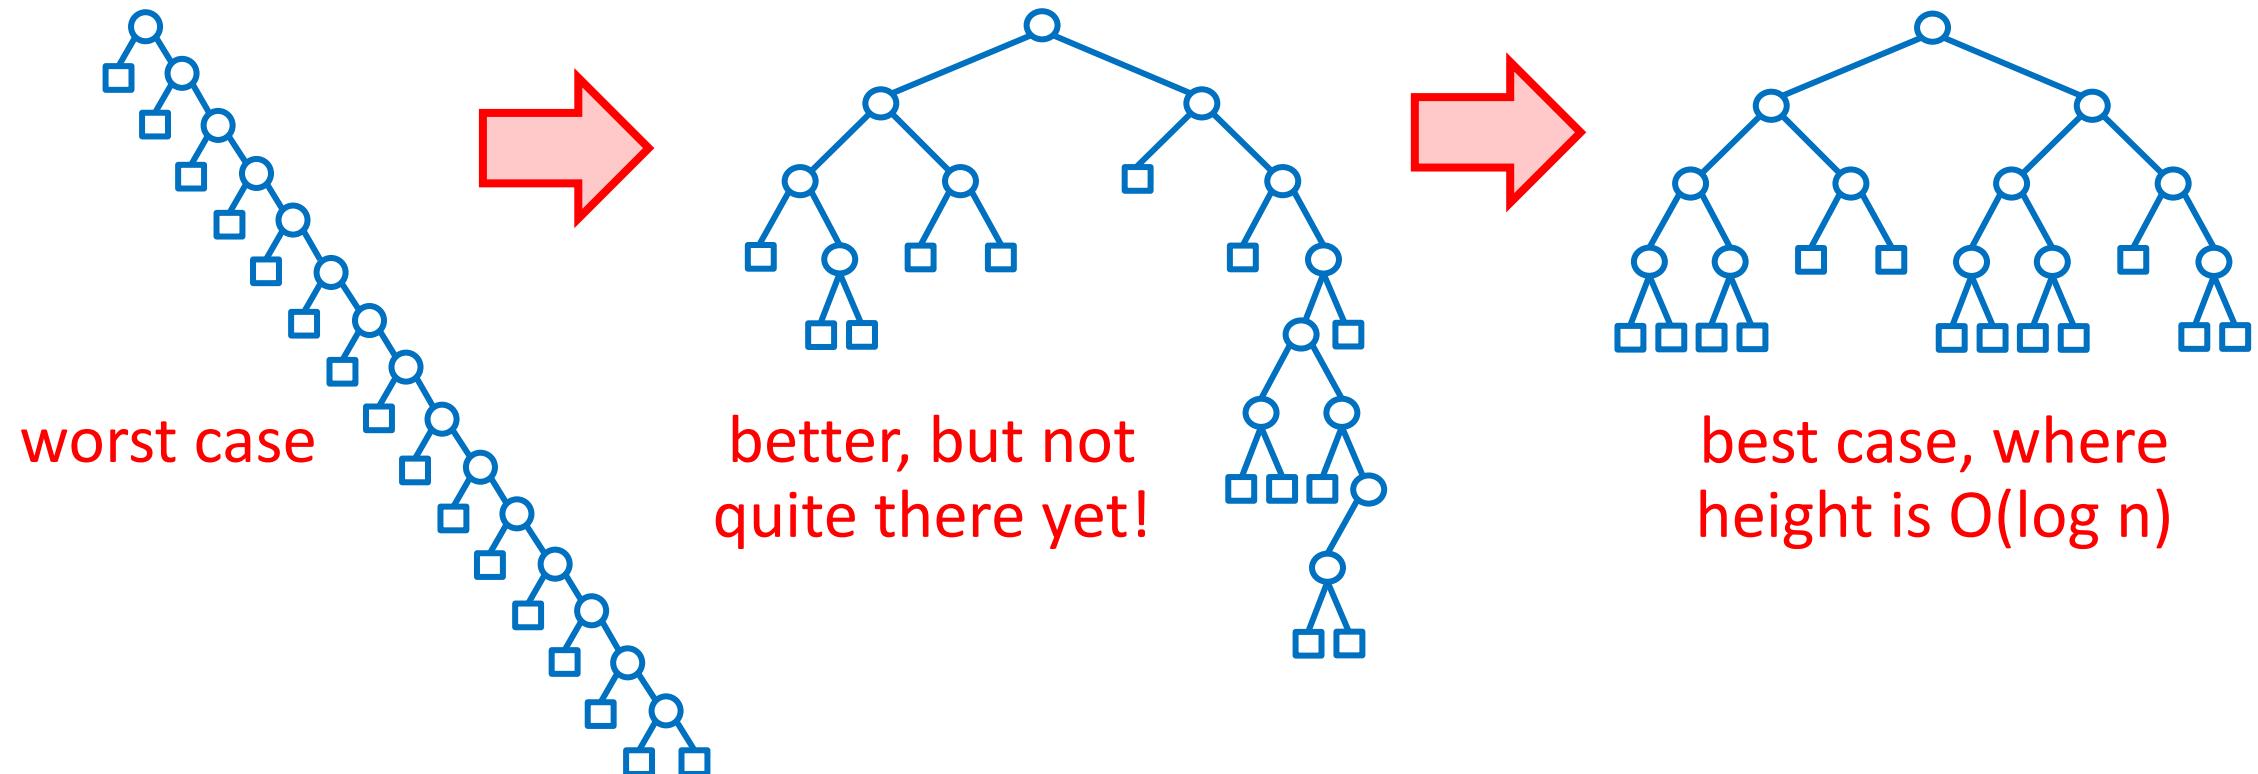


- Given k internal nodes, what is the maximum height of a binary search tree, h ?

In the **worst case**, $h = k$, implying that insertion and removal take $O(k)$ time!

BINARY SEARCH TREE – PERFORMANCE

- The height can be very large if the binary search tree is extremely **unbalanced**, e.g. there are **nodes whose right-subtree is much taller than the left subtree!**



18

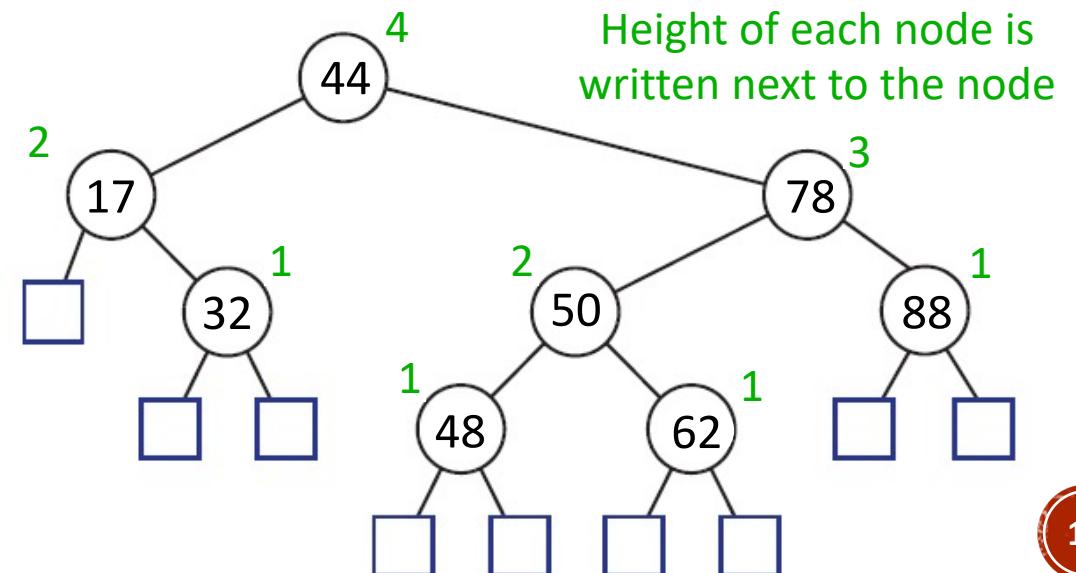
AVL TREES

AVL TREES

- For all the **implementations** of a **map** or **dictionary** we've seen so far (be it a **list**, an **array**, or a **binary search tree (BST)**), some operations run in **O(n)** time still!
- To make all operations run in **O(log n)** time, all we need is to add one more requirement to the structure of a BST **T**:

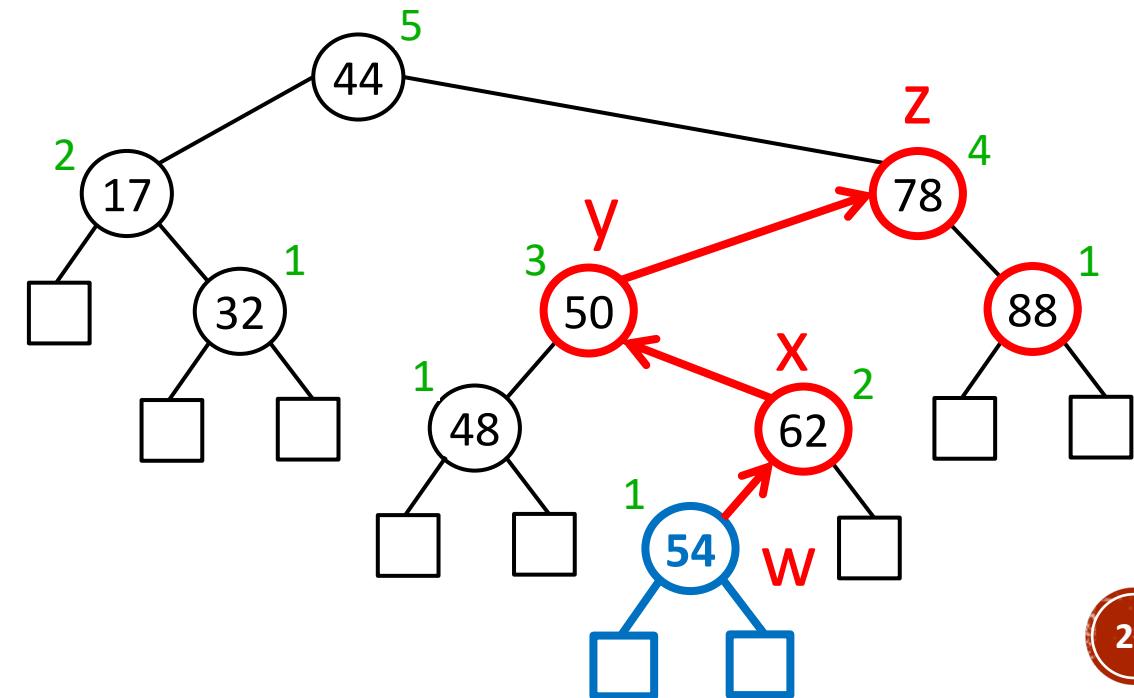
Height-Balance Property: Every internal node **v** of **T** is “**balanced**”, i.e., the heights of the children of **v** differ by **at most 1**.

- Any binary search tree that satisfies this property is an **AVL tree**, named after its inventors (Georgy Adelson-Velsky and Evgenii Landis)!
- This property will assure that the **height** of a BST containing **n** entries is **O(log n)**



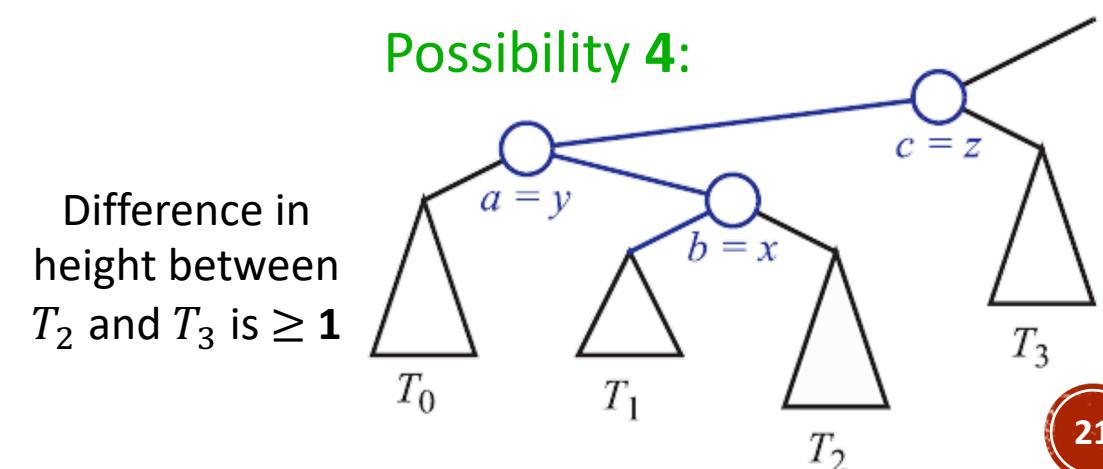
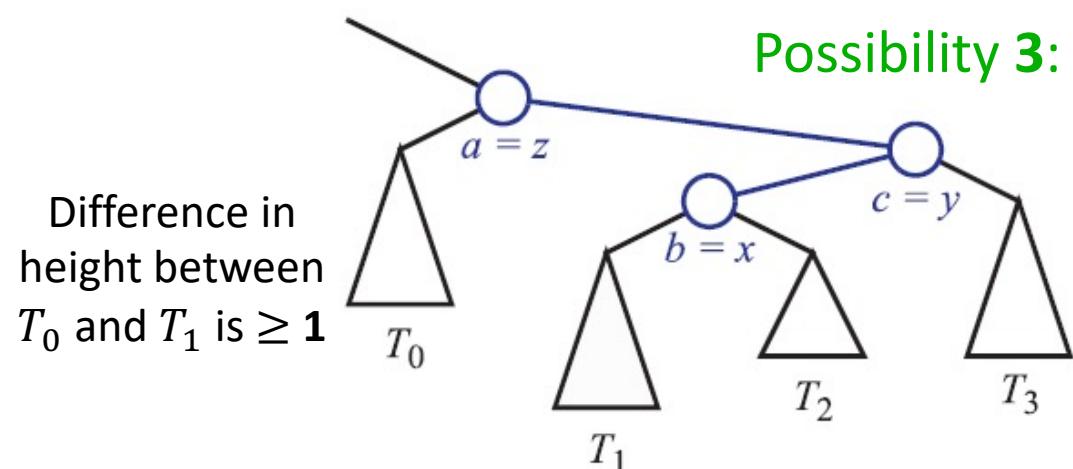
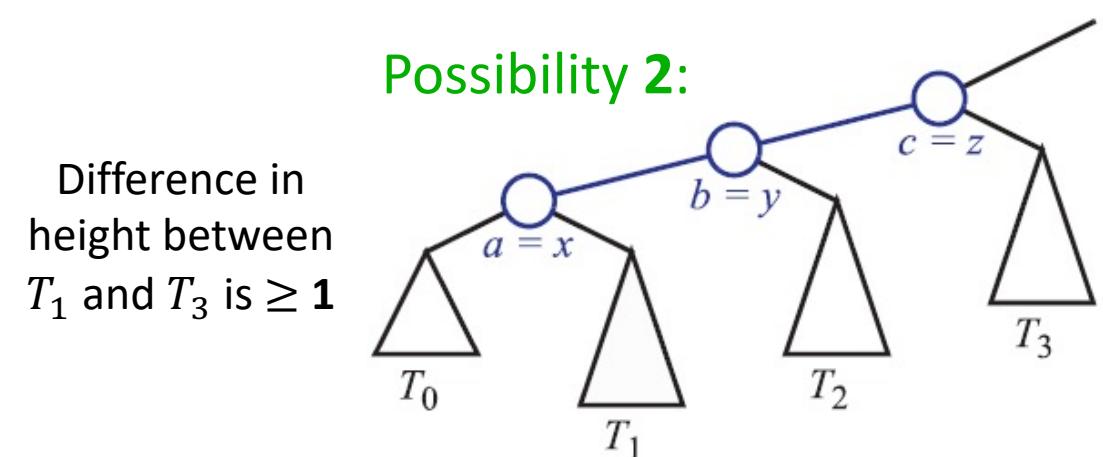
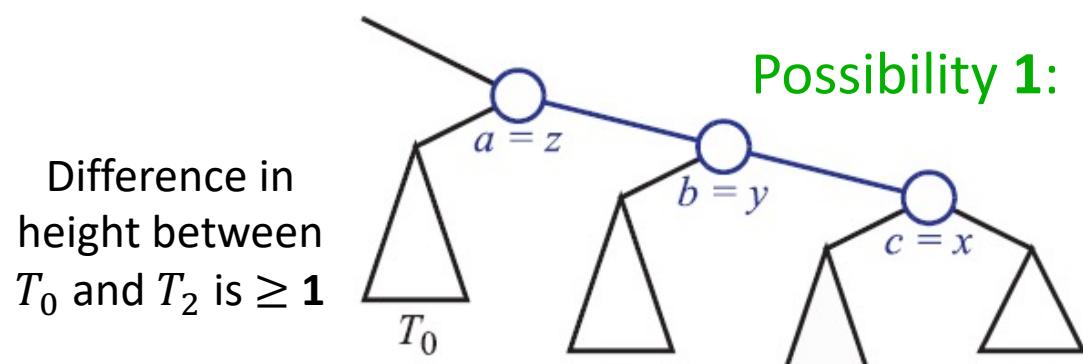
AVL TREES – INSERTION

- Recall that **insertion** always inserts the new entry at a node **w** that was previously an **external node**, and makes **w** an **internal node** with two external children
- Here is an example of an AVL tree that becomes **unbalanced** after inserting a node **w** whose key is **54**, and we want to **determine where to fix the tree**.
 - Let **z** be the **first unbalanced node** we encounter when going from **w** upwards
- The height of **y** is **greater by 2** compared to the height of its sibling, making **z** unbalanced!
- Our **repair strategy** will focus on nodes **x**, **y**, and **z**, using the *trinode restructuring* function



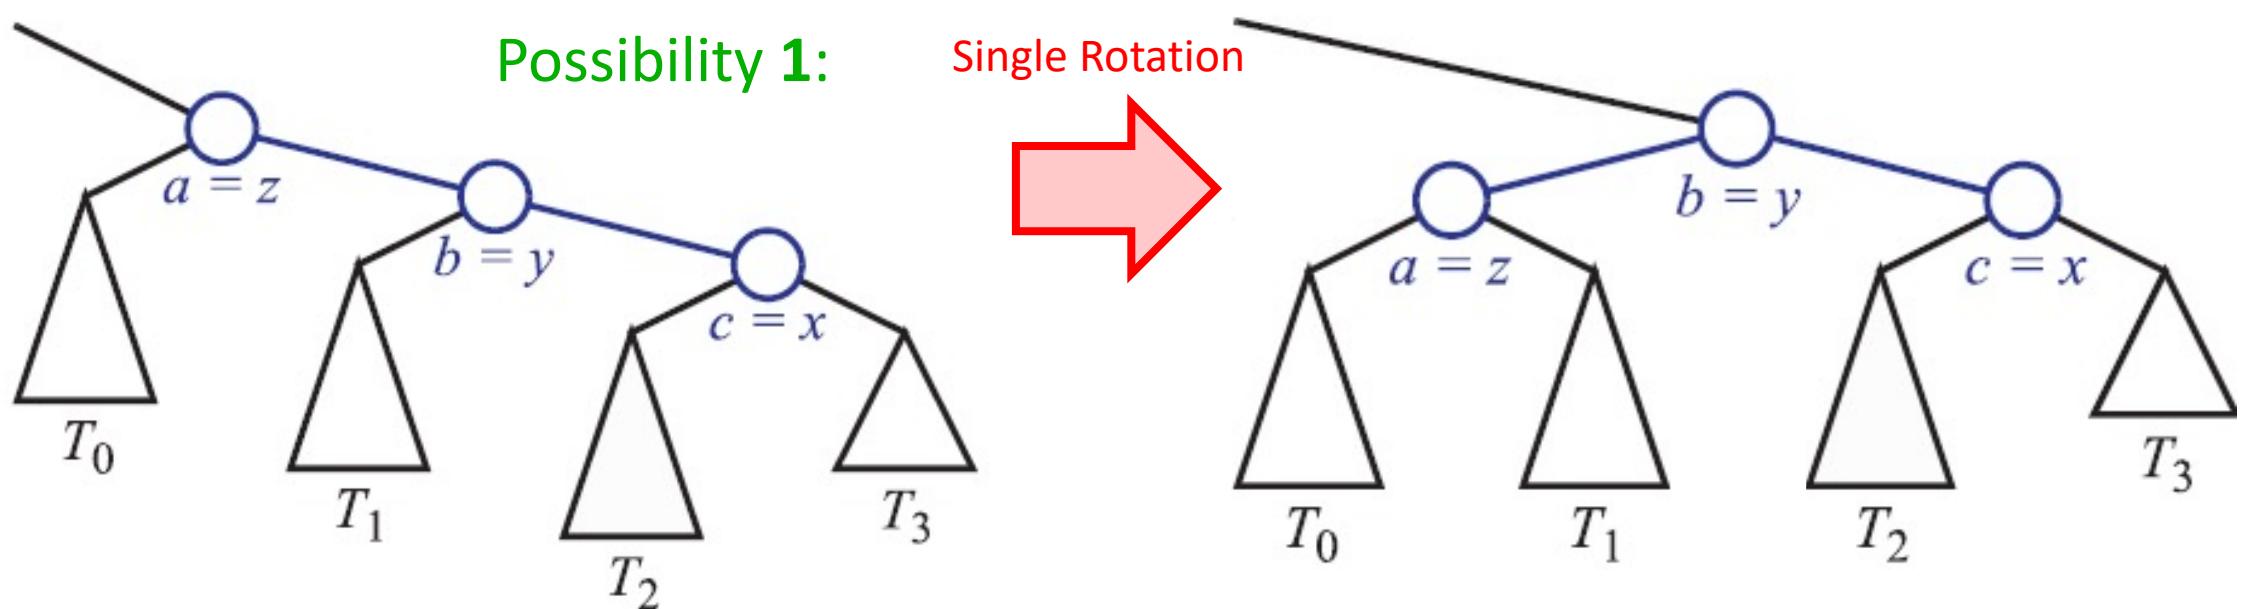
TRINODE RESTRUCTURING

- Here are the **4 possibilities** of how z becomes **unbalanced** after an insertion operation:



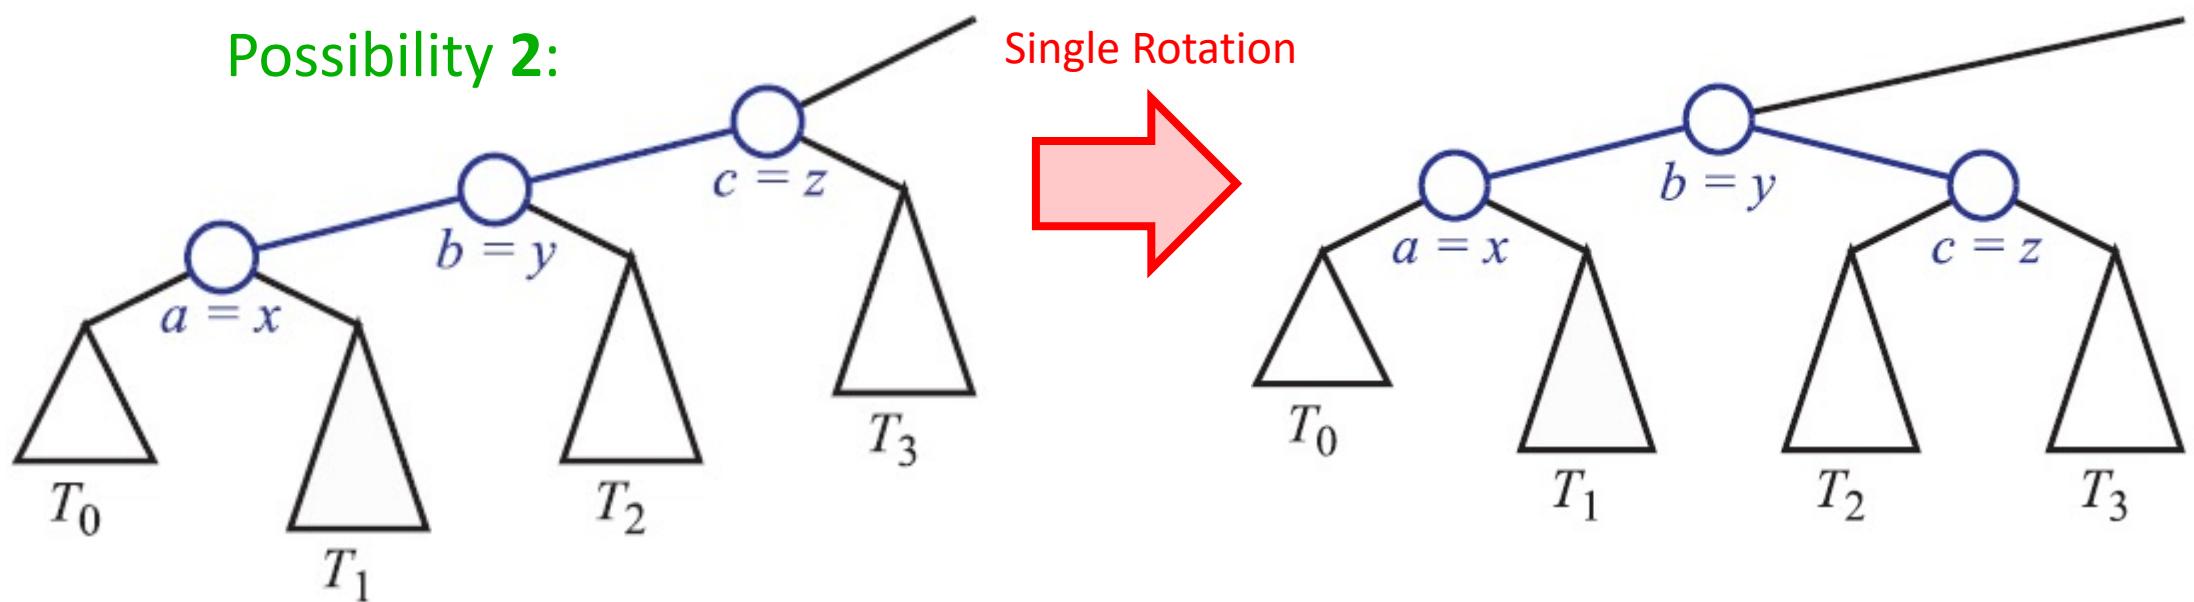
TRINODE RESTRUCTURING

- A **trinode restructuring** does the following:
 1. Make the **b** the parent of **a** and **c**
 2. Make T_0 and T_1 the left and right subtrees of **a**, respectively
 3. Make T_2 and T_3 the left and right subtrees of **c**, respectively



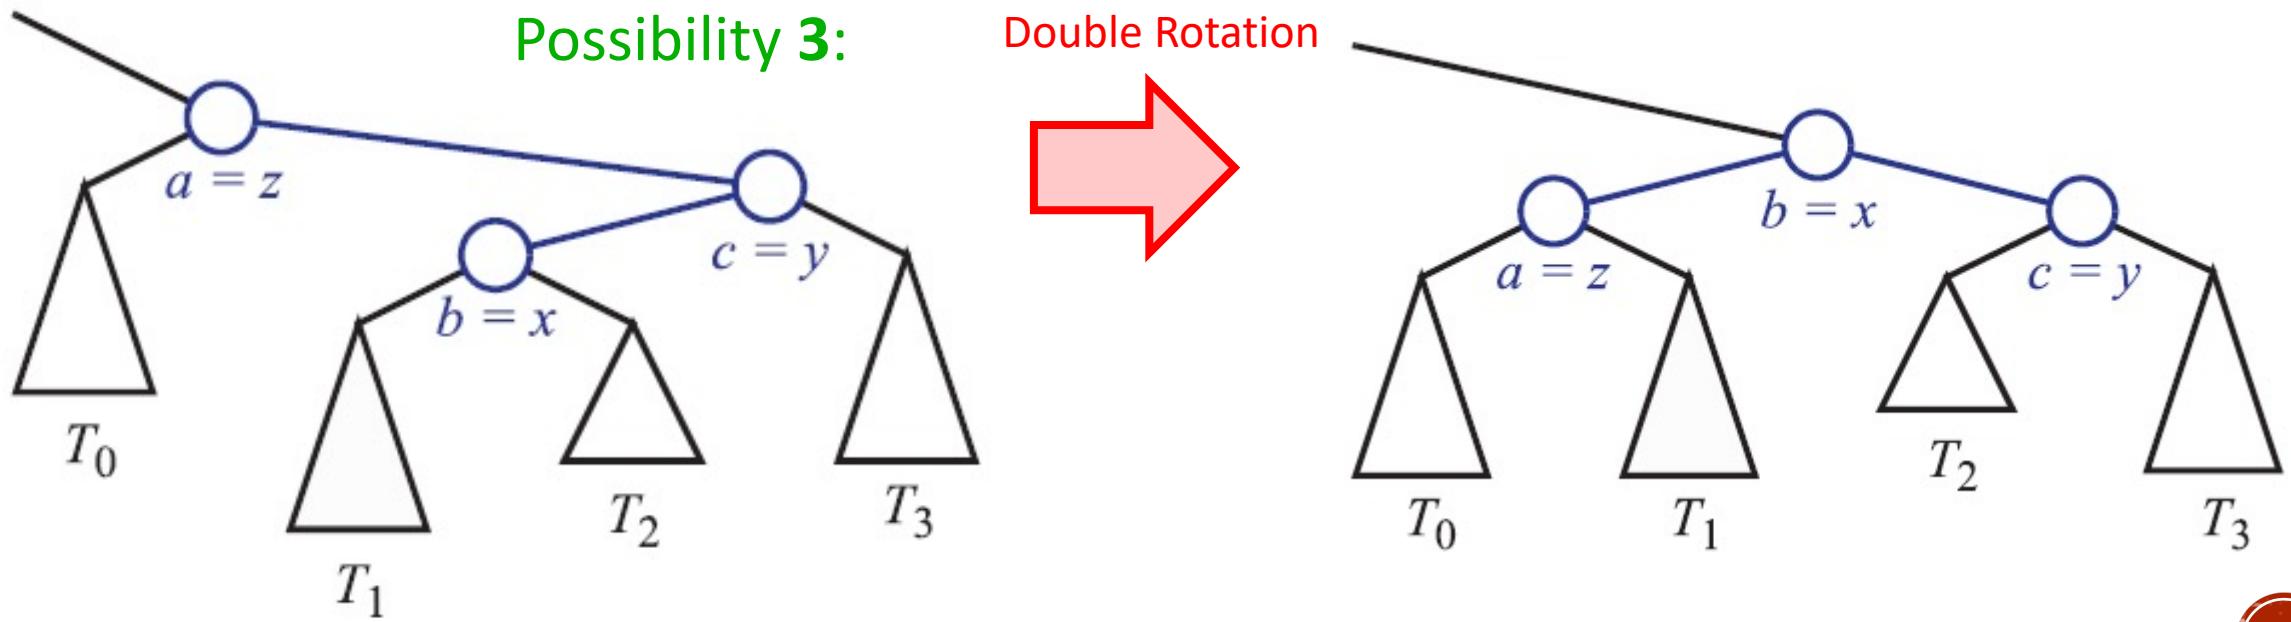
TRINODE RESTRUCTURING

- A **trinode restructuring** does the following:
 1. Make the **b** the parent of **a** and **c**
 2. Make T_0 and T_1 the left and right subtrees of **a**, respectively
 3. Make T_2 and T_3 the left and right subtrees of **c**, respectively



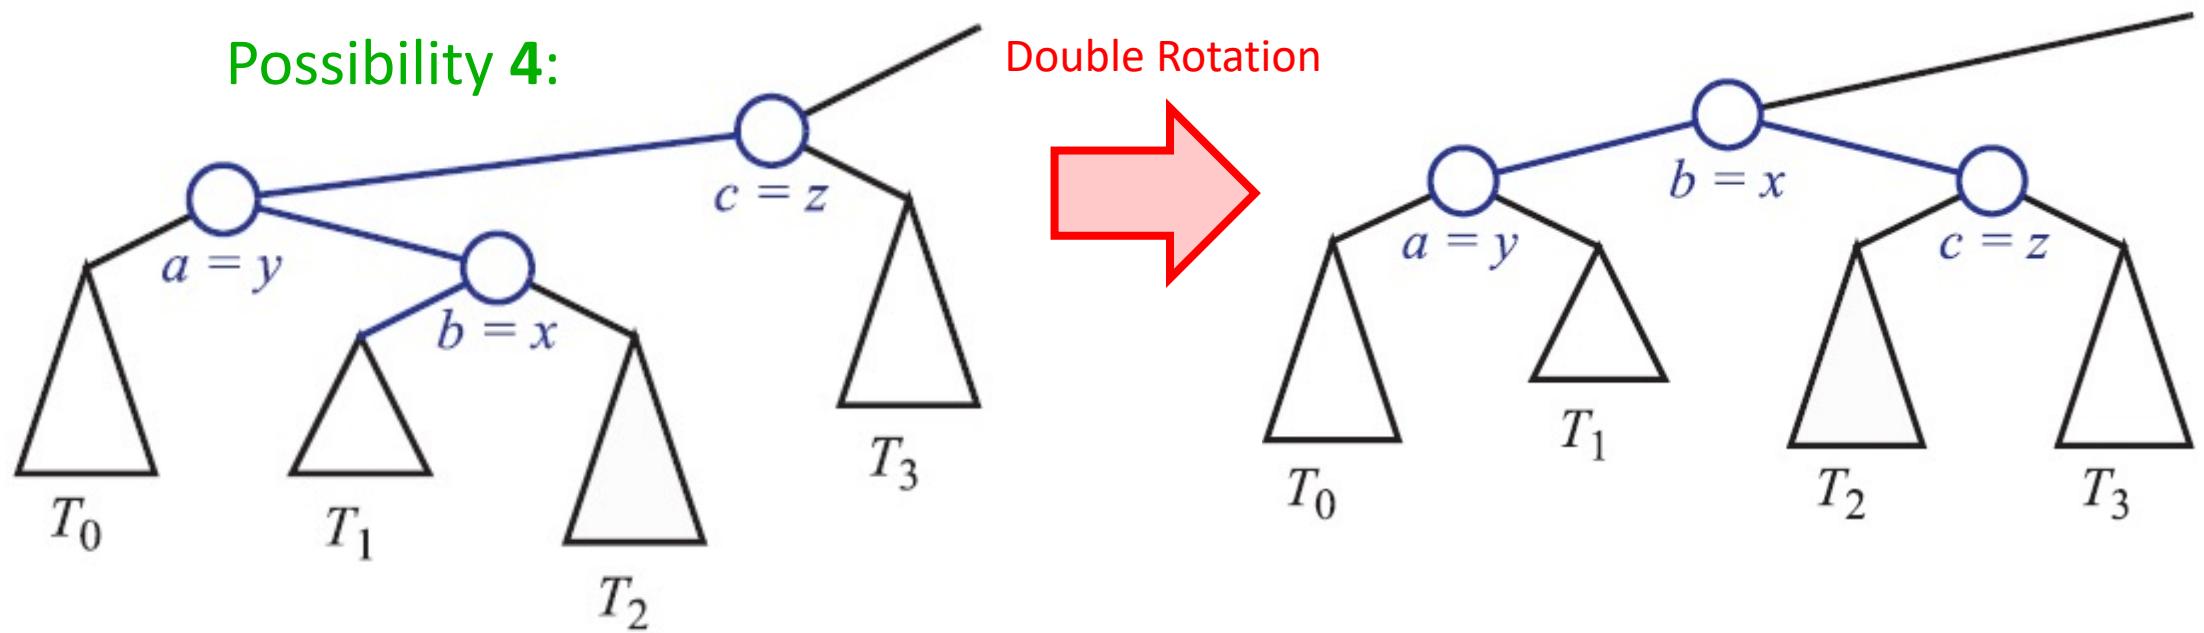
TRINODE RESTRUCTURING

- A **trinode restructuring** does the following:
 1. Make the **b** the parent of **a** and **c**
 2. Make T_0 and T_1 the left and right subtrees of **a**, respectively
 3. Make T_2 and T_3 the left and right subtrees of **c**, respectively



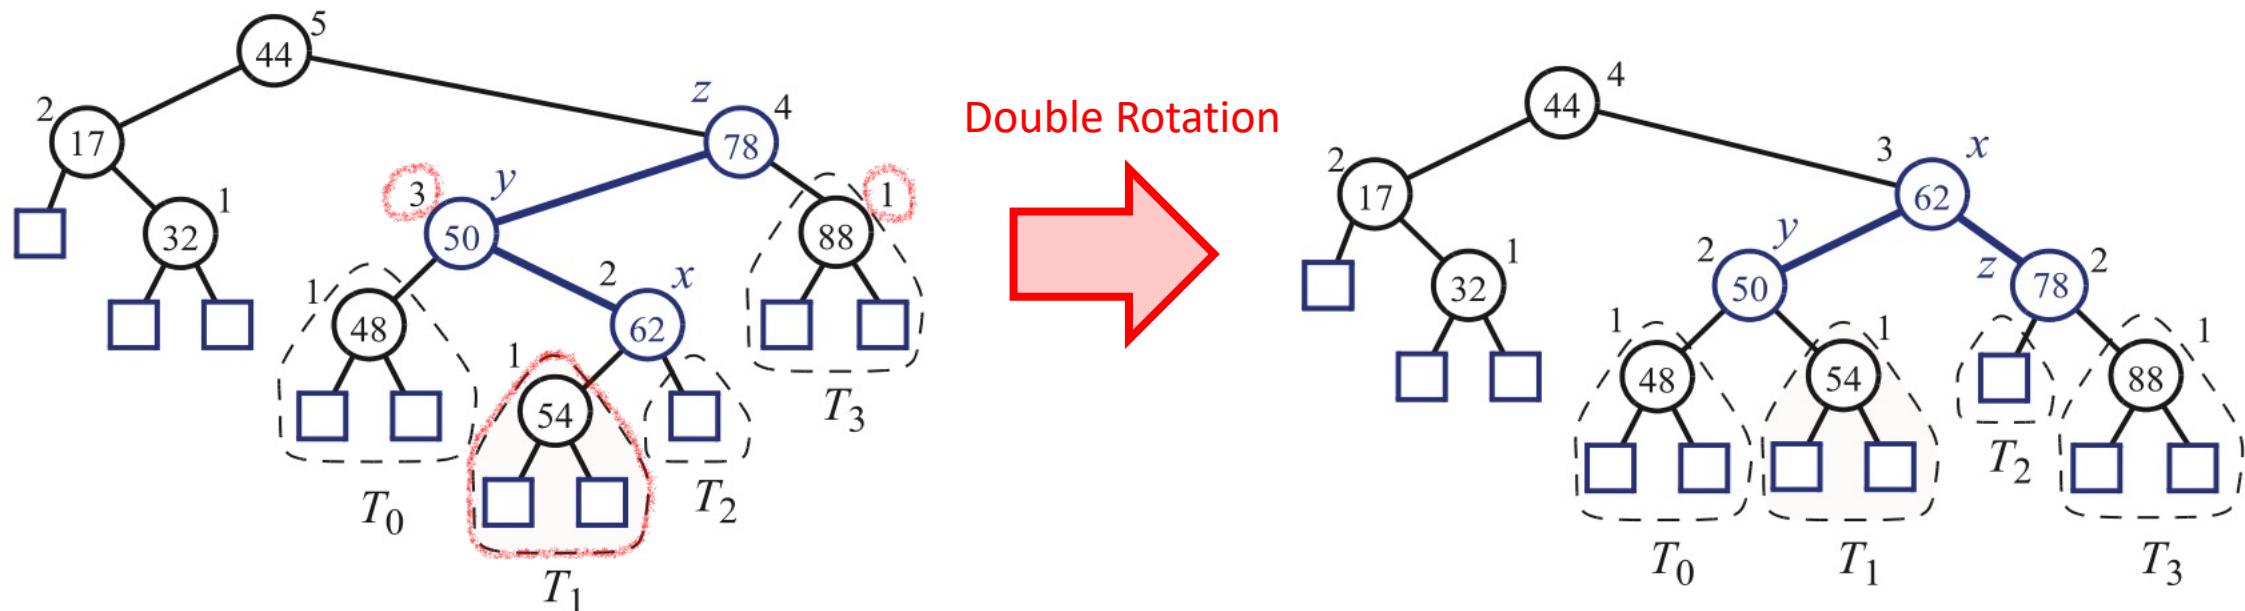
TRINODE RESTRUCTURING

- A **trinode restructuring** does the following:
 1. Make the **b** the parent of **a** and **c**
 2. Make T_0 and T_1 the left and right subtrees of **a**, respectively
 3. Make T_2 and T_3 the left and right subtrees of **c**, respectively



TRINODE RESTRUCTURING

- An example insertion of an entry with key 54 to an AVL tree
 - After adding a new node for key 54, the node storing key 78 becomes unbalanced
 - A **trinode restructuring** restores the height-balance property.

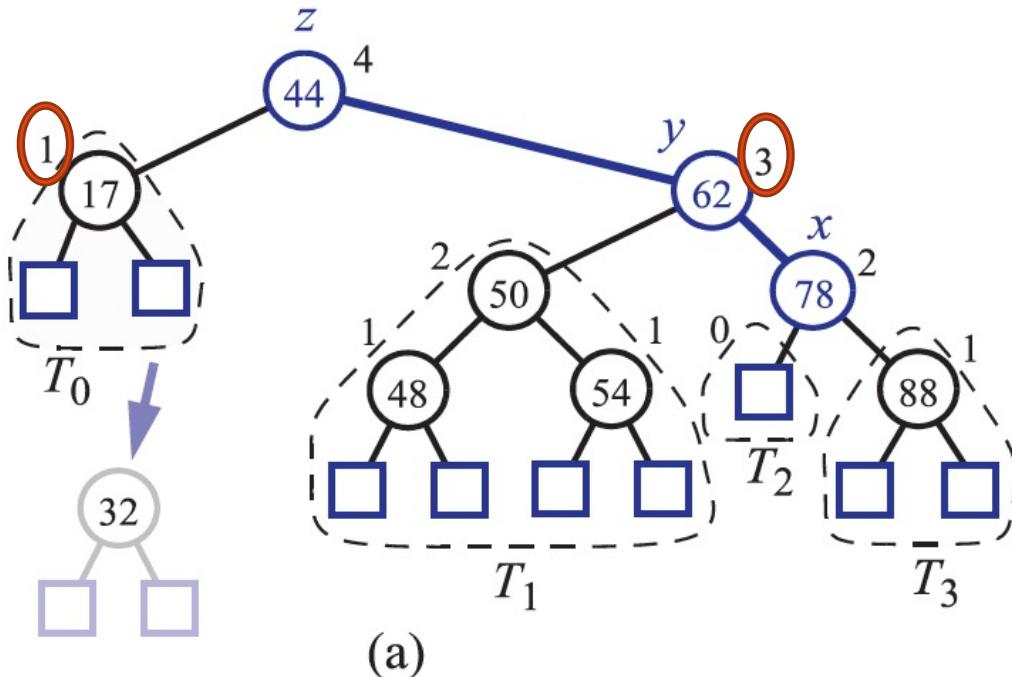


AVL TREES – REMOVAL

- We said **insertion** in AVL trees consists of **two steps**:
 - **Insert the node** as usual, just like in any binary search tree
 - **Fix any unbalanced nodes** using the **trinode restructuring** function
- **Removing** a node is done in the **same way**:
 - **Remove the node** as usual, just like in any binary search tree via **removeAboveExternal()**
 - **Fix any unbalanced nodes** using the **trinode restructuring** function
 - However, fixing an unbalanced node **may make its ancestors unbalanced**, thus we need to repeatedly go up the tree, **looking for (and fixing) unbalanced nodes!**

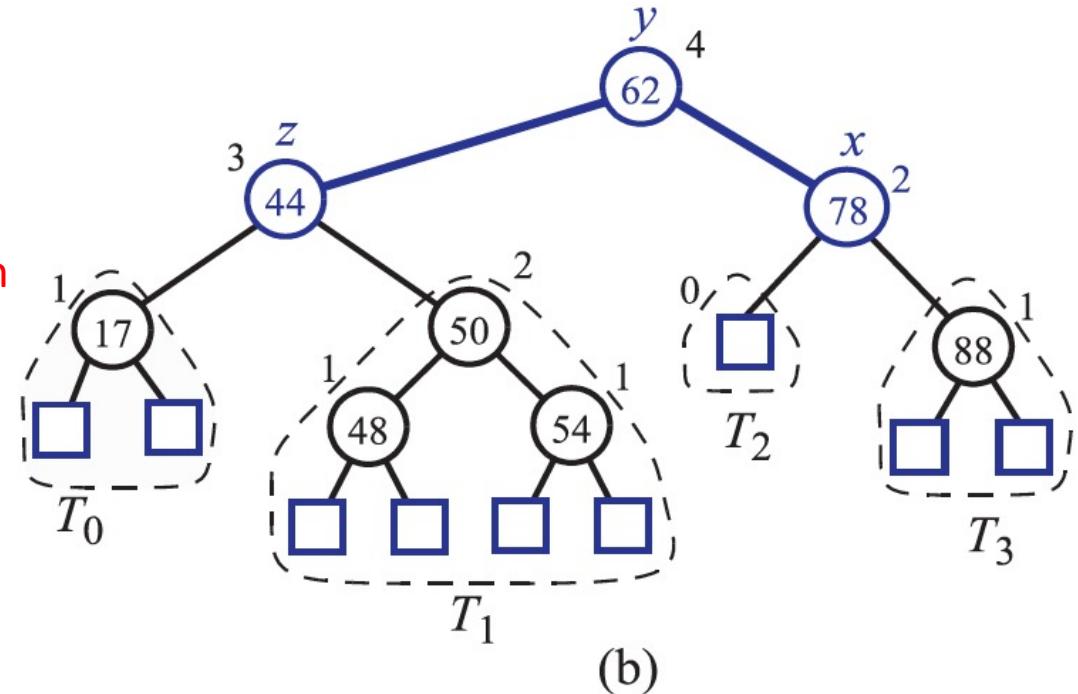
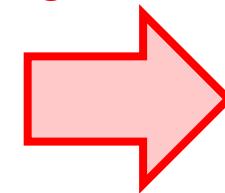
AVL TREES – REMOVAL

Example: we want to remove node with key 32 from the following tree:



Removal of the entry with key 32

Single Rotation



Trinode restructuring restores the height-balance property

AVL TREES – COMPLEXITY

With AVL trees:

- **Searching** for a node takes $O(\log n)$ time, because:
 - **Searching** a binary search tree takes $O(h)$ time, where h is the **height** of the tree
 - The **height** of AVL trees is $O(\log n)$
- **Insertion** and **removal** take $O(\log n)$ time, because:
 - **Insertion** and **removal** require searching, which takes $O(\log n)$, as well as **trinode restructuring**, which takes $O(1)$ (which could be repeated in the case of **removal**)

