

30

MULTI-WAY SEARCH TREES

MULTI-WAY SEARCH TREE

There are data structures that use a **multi-way search tree** where the following holds:

- Each internal node has **at least 2 children** (a node with d children is called a d -node)
- A d -node has $d - 1$ entries, **sorted** by their keys, i.e.:

$$k_1 \leq \dots \leq k_{d-1}$$

- Let us define $k_0 = -\infty$ and $k_d = \infty$.
- Then, for every entry (k, x) in the subtree rooted at the i^{th} child, we have:

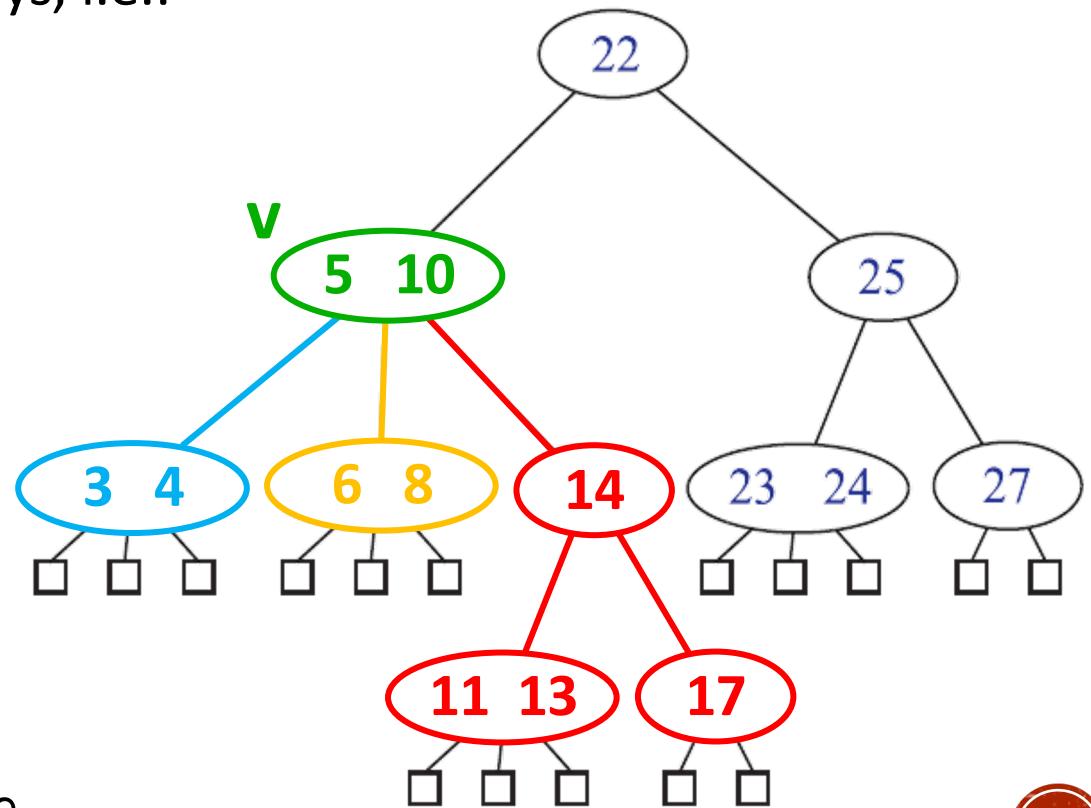
$$k_{i-1} \leq k \leq k_i$$

Example: for node **v**:

$$-\infty \leq (3 \text{ and } 4) \leq 5$$

$$5 \leq (6 \text{ and } 8) \leq 10$$

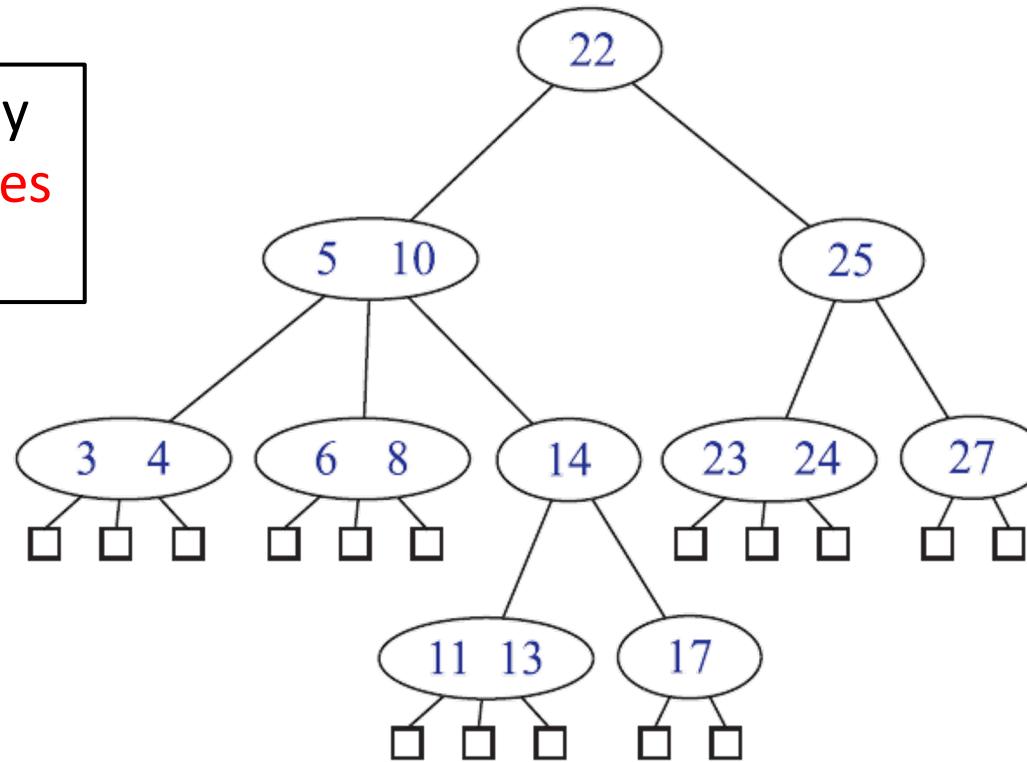
$$10 \leq (11, 13, 14, \text{ and } 17) \leq \infty$$



MULTI-WAY SEARCH TREE

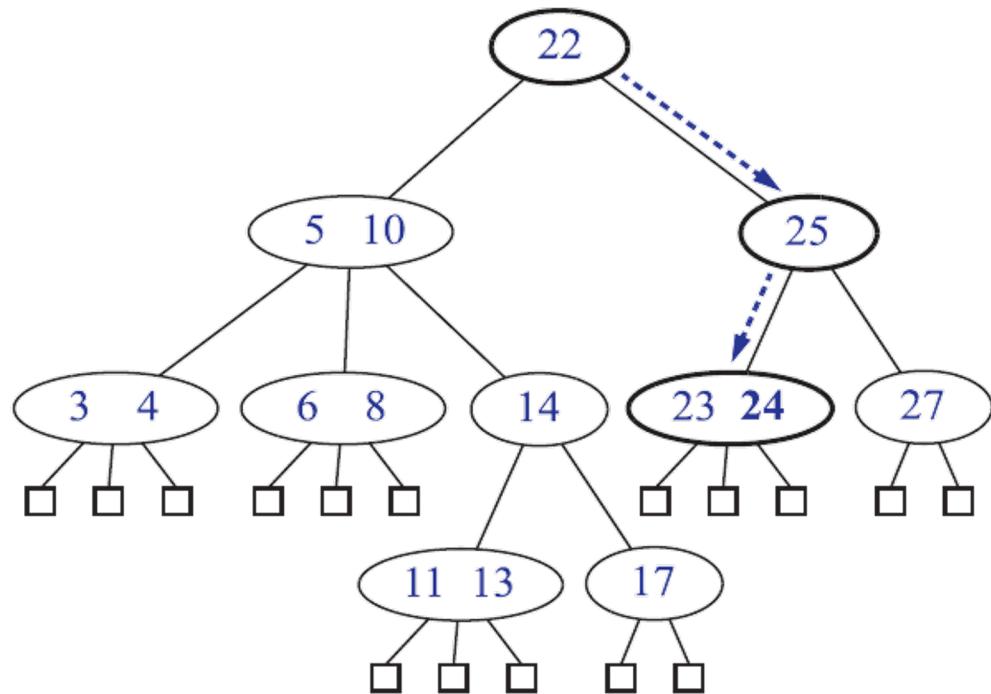
- Notice how, in a **multi-way search tree**:
 - **External nodes** are “*place holders*” that do not contain any entries
 - The node that has j external children has $j - 1$ entries

Proposition: A multi-way search tree with n entries has $n+1$ external nodes

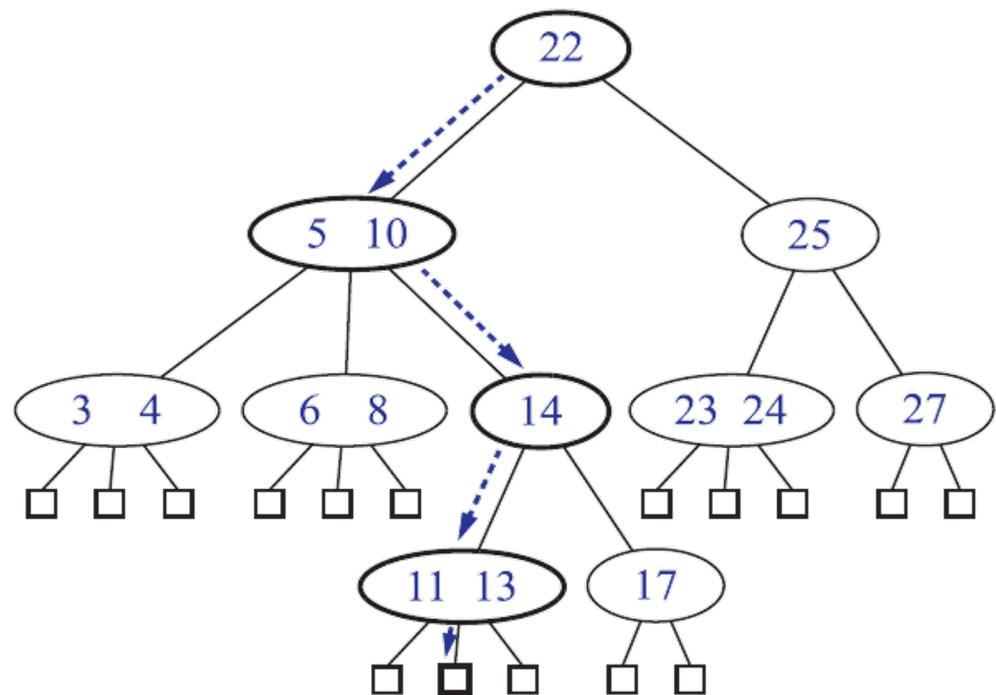


MULTI-WAY SEARCH TREE

- A **binary search tree** is actually a **special case** of a **multi-way search tree!**
- In both types of trees, **search is done in exactly the same way!**



How we searched for, and found, an entry with key = 24



How we searched for an entry with key = 12, and determined that it doesn't exist

IMPLEMENTING EACH NODE

- In any given node, entries are stored in a **map** (remember: a **map** is any data structure that stores **key-value pairs**, called **entries**)

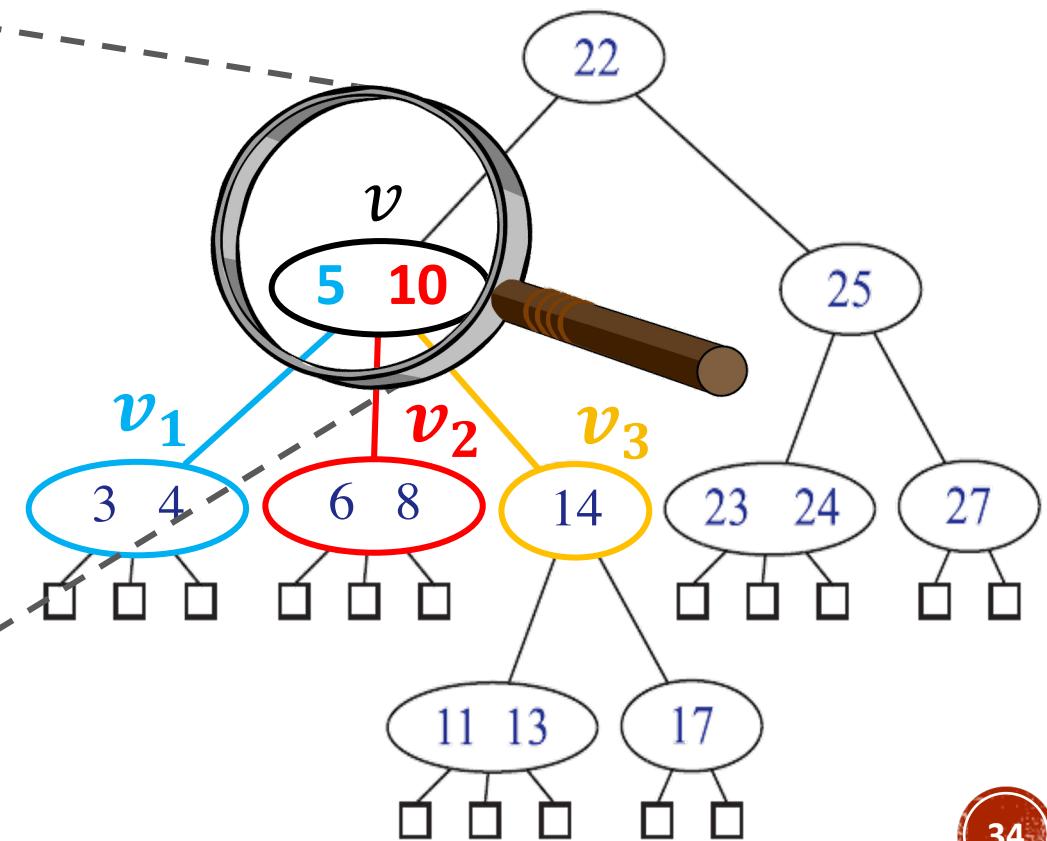
For node v , suppose
the entries were as follows

$(5, x_1)$
 $(10, x_2)$

(recall that the illustration only shows the keys, and omits the values, i.e., omits x_i)

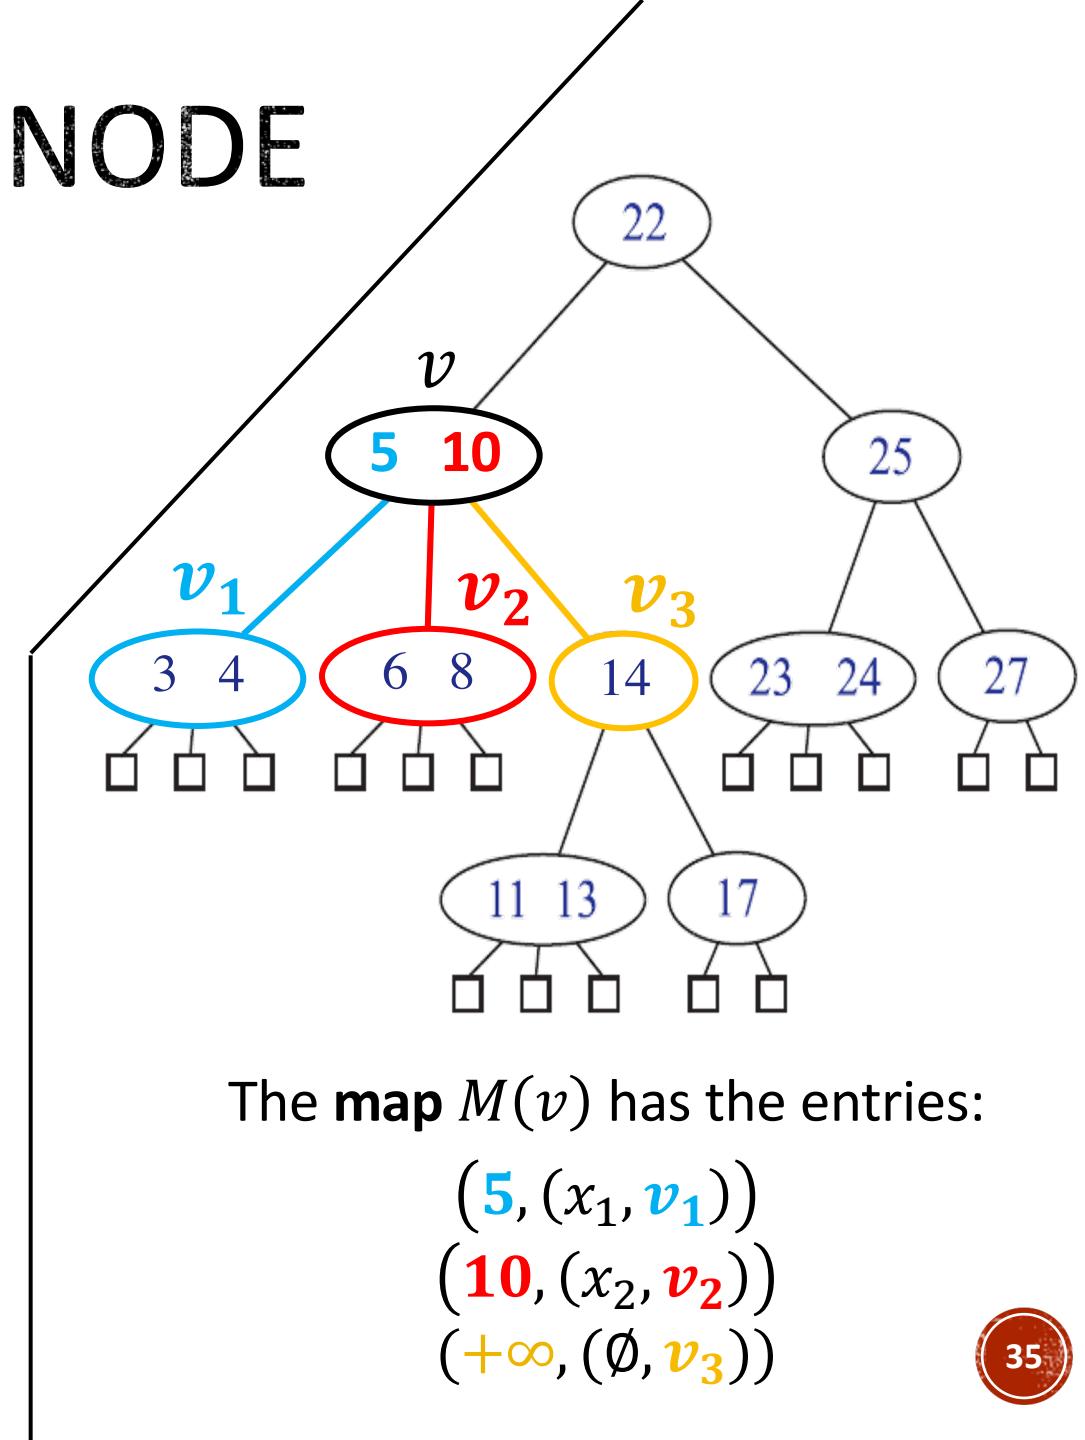
Then the **map** $M(v)$ will have the entries:

$(5, (x_1, v_1))$
 $(10, (x_2, v_2))$
 $(+\infty, (\emptyset, v_3))$



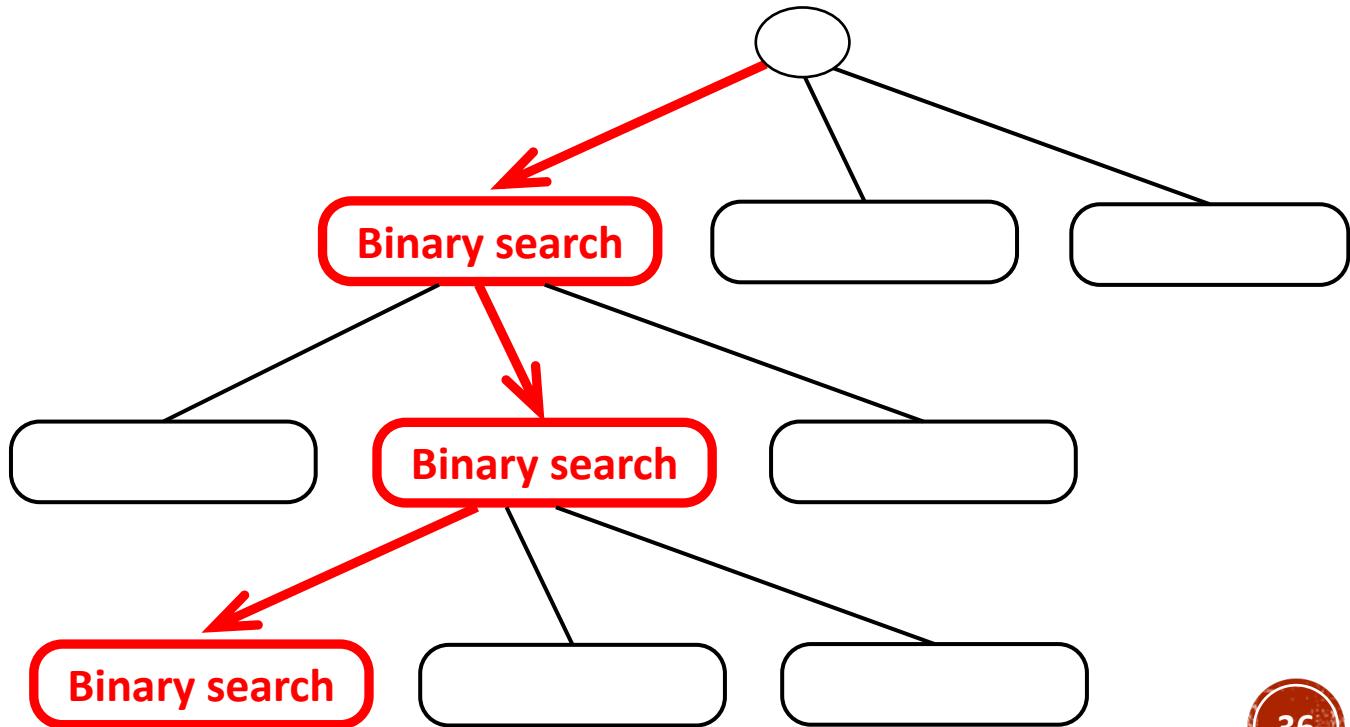
IMPLEMENTING EACH NODE

- Suppose we are **searching the tree** for an entry with key = k .
- Then, if we reach node v , we **search the map** $M(v)$ to find the entry $(k_i, (x_i, v_i))$ with the smallest key that is $k_i \geq k$. Now:
 - If $k_i > k$, we continue by the search, but this time by **searching the map** $M(v_i)$
 - If $k_i = k$ the search **terminates** successfully
- If v is a **d-node** then searching it takes $O(\log d)$ if $M(v)$ is implemented as an **ordered search table**



SEARCH COMPLEXITY

- To summarize, we said that searching the map of a **d-node** takes $O(\log d)$ time, and we may have to proceed to searching one of its children, etc.
- Thus, **searching the entire tree** takes $O(h \log d_{max})$ where:
 - d_{max} is the max number of children of a node in the tree
 - h is the height of the tree
- Thus, to search a multi-way tree efficiently, **we need to keep the height as small as possible**, i.e., a logarithmic function of n .
- A search tree with logarithmic height such as this is called a **balanced search tree**.



37

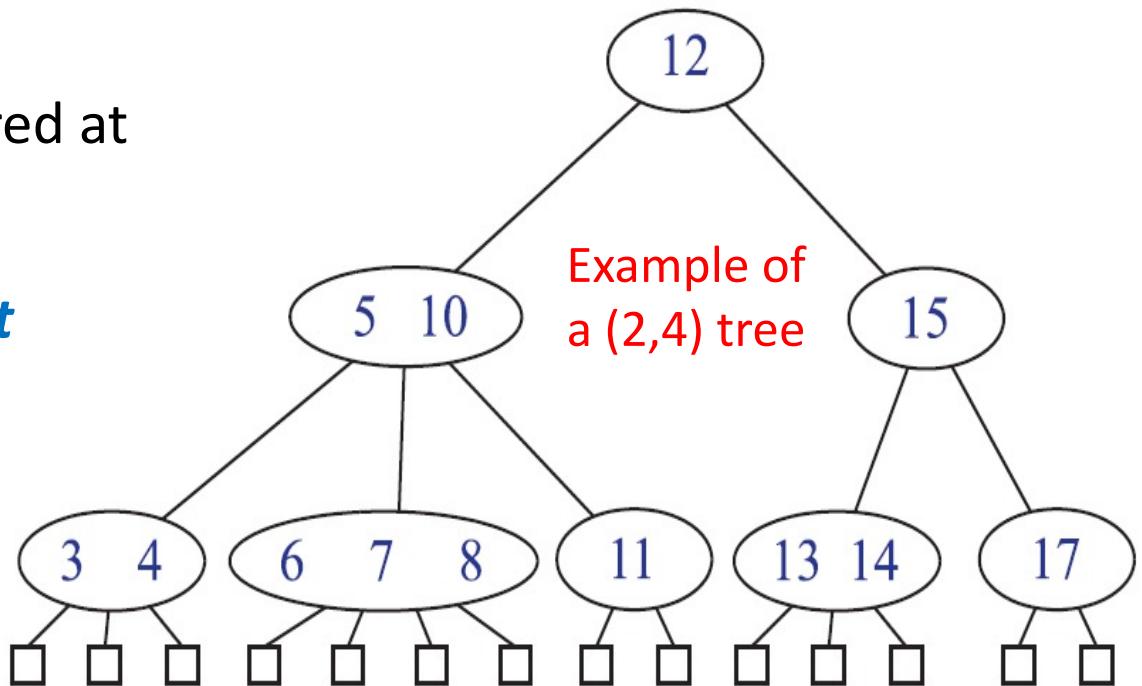
(2,4) TREES

(2,4) TREES

- A “(2,4) tree” is a multi-way search tree that satisfies the following properties:
 - **Depth property:** All external nodes have the same depth
 - **Size property:** Every internal node has at most 4 children
- It is called a “(2,4) tree” because each internal node has:
 - At least 2 children by the definition of multi-way search trees
 - At most 4 children by the **size property**
- Note that the **size property** implies that each internal node has at most 3 entries

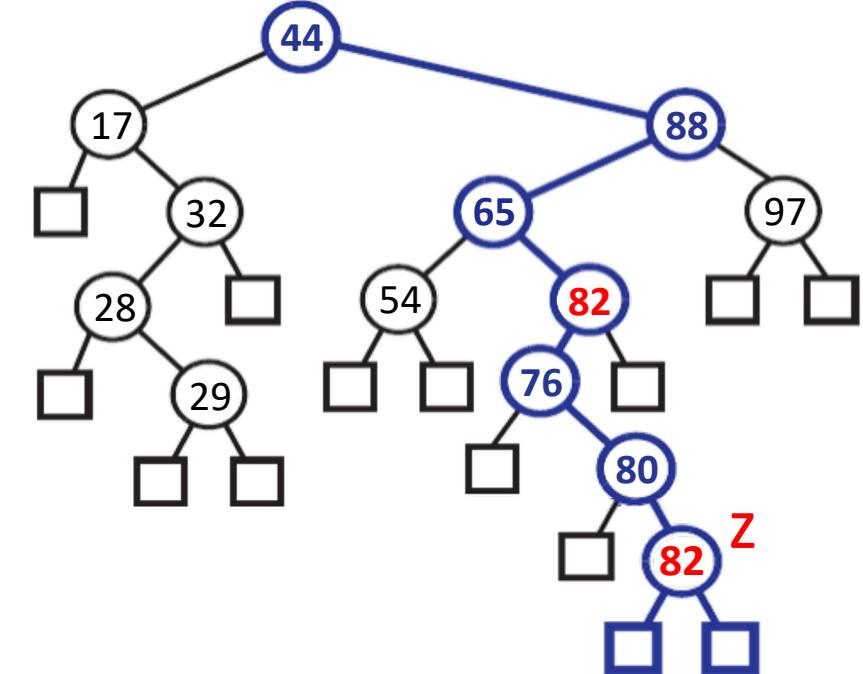
(2,4) TREES

- A “(2,4) tree” is a multi-way search tree that satisfies the following properties:
 - **Depth property:** All external nodes have the same depth
 - **Size property:** Every internal node has at most 4 children
- Note that:
 - The **size property** keeps the map stored at each node small
 - The **depth property** makes the height of the tree logarithmic in n



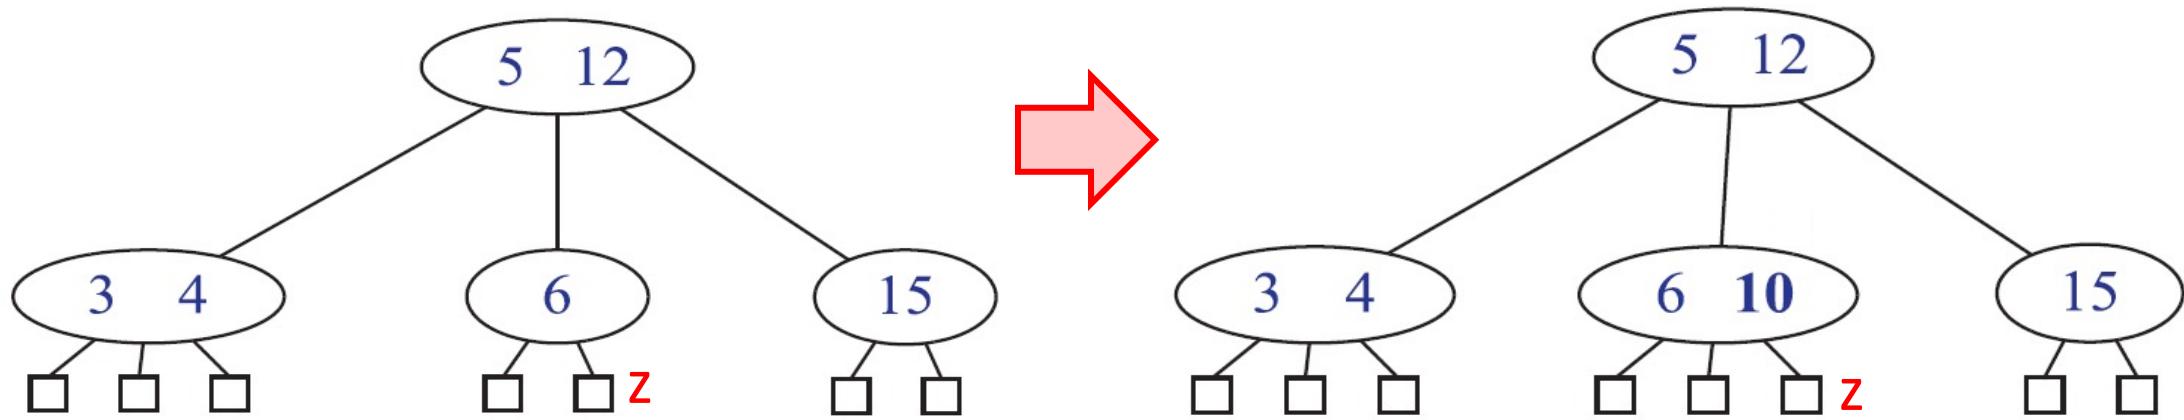
BINARY SEARCH TREE – INSERTION

- Remember: In a binary search tree, to **insert** an entry (k, v) , we first **search** for an entry whose key = k . This search returns either:
 - an **internal** node whose key = k , or
 - an **external** node where we can insert
- If we find an **internal** node, then the **tree already contains an entry whose key = k** , in which case we search further down the tree, **until we reach an external node!**
- Thus, insertion always happens at an **external node!**
- **The same idea applies in (2,4) trees**, i.e., to insert an entry (k, v)
 - We repeatedly **search** for an entry whose key = k , until we reach an external node
 - Then we insert (k, v) at that external node, denoted by **z**



(2,4) TREES – INSERTION

- Here is an example of how to **insert 10**. We **searched** for **10**, and our search did not find it, but **found an external node z** where we can insert **10**.
- We **insert the entry** in the **parent** of **z**, and add an external node to the left of **z**



- The only problem is if the parent of **z** already had **4 children**, which is **the maximum number of permitted children** in a (2,4) tree!
- In this case, after the insertion, the parent of **z** will have **5 children**, which we call an “**overflow**”. To fix this, we need to perform a “**split**” operation.

(2,4) TREES – SPLIT TO FIX AN OVERFLOW

Here is an example where we insert an entry in node v , leaving it with four entries (k_1, k_2, k_3, k_4) and 5 children, resulting in an **overflow**. Here is how we **split v** :

➤ v is split into two nodes, v' and v''

➤ k_1, k_2 go to v'

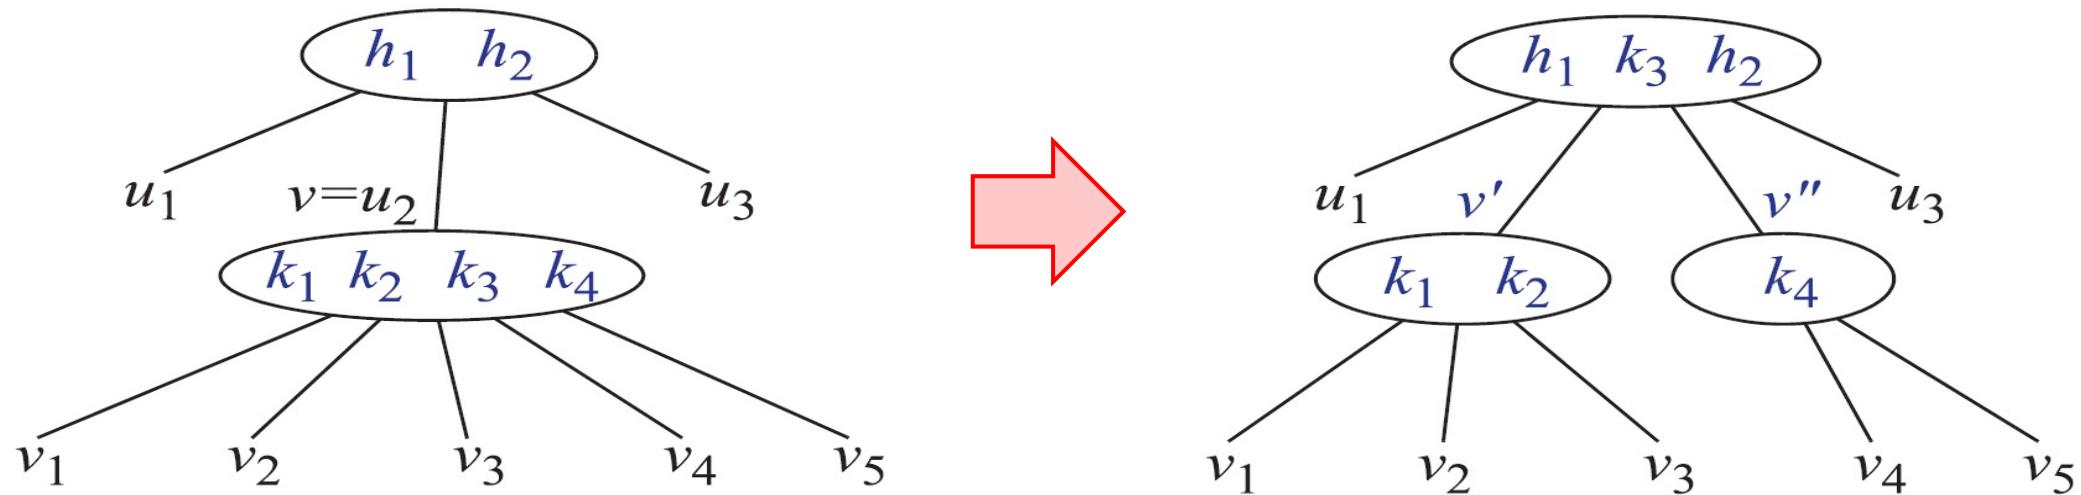
➤ k_3 goes to the **parent** of v

➤ k_4 goes to v''

As for the **children**:

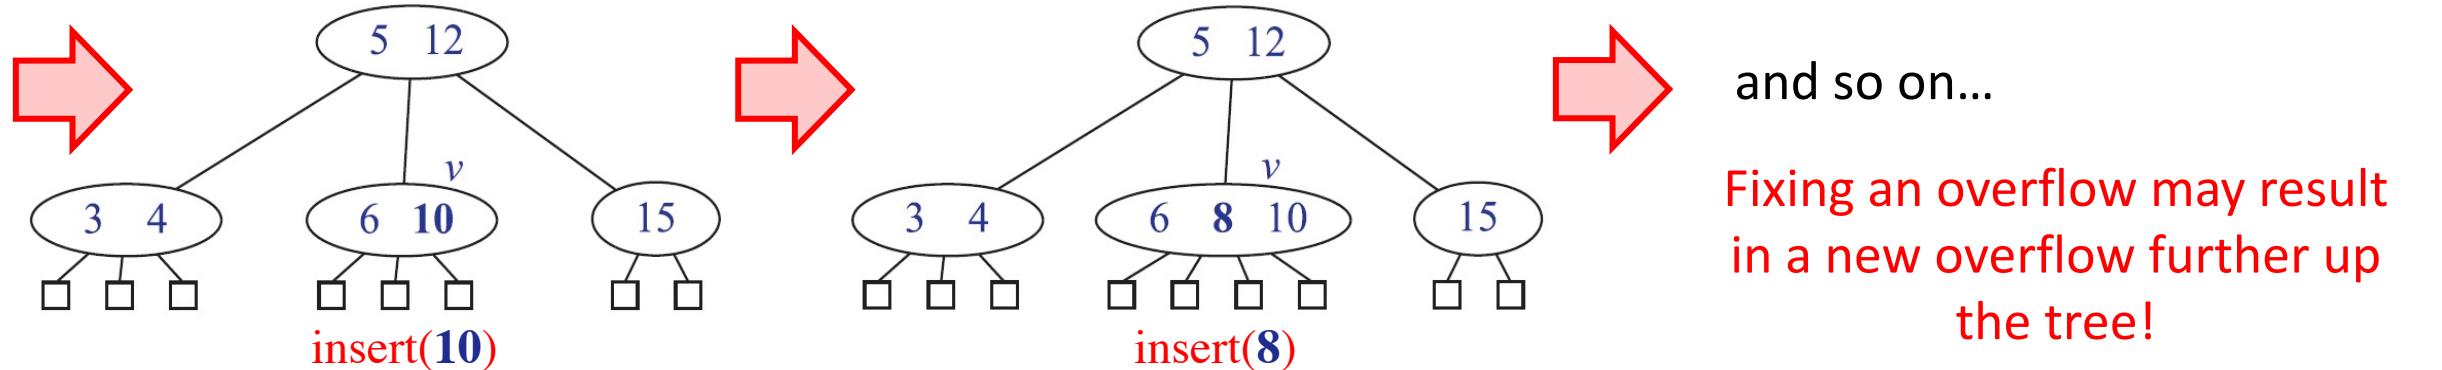
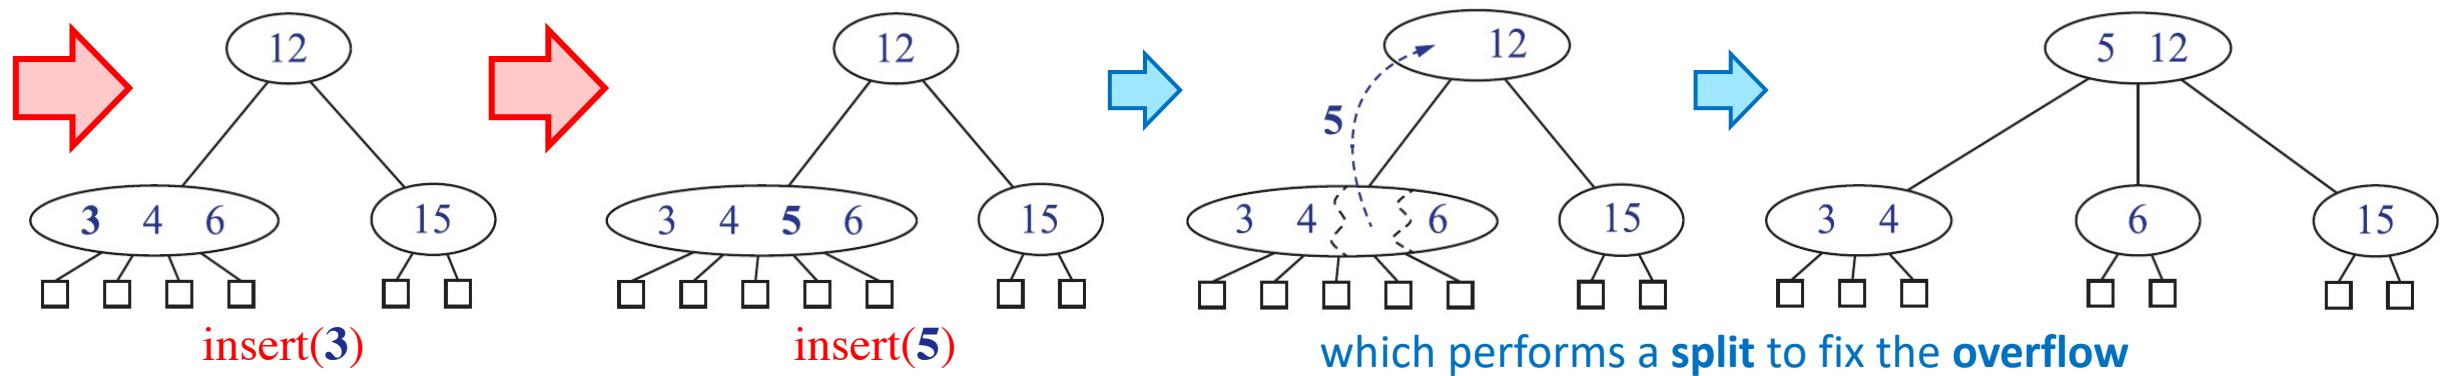
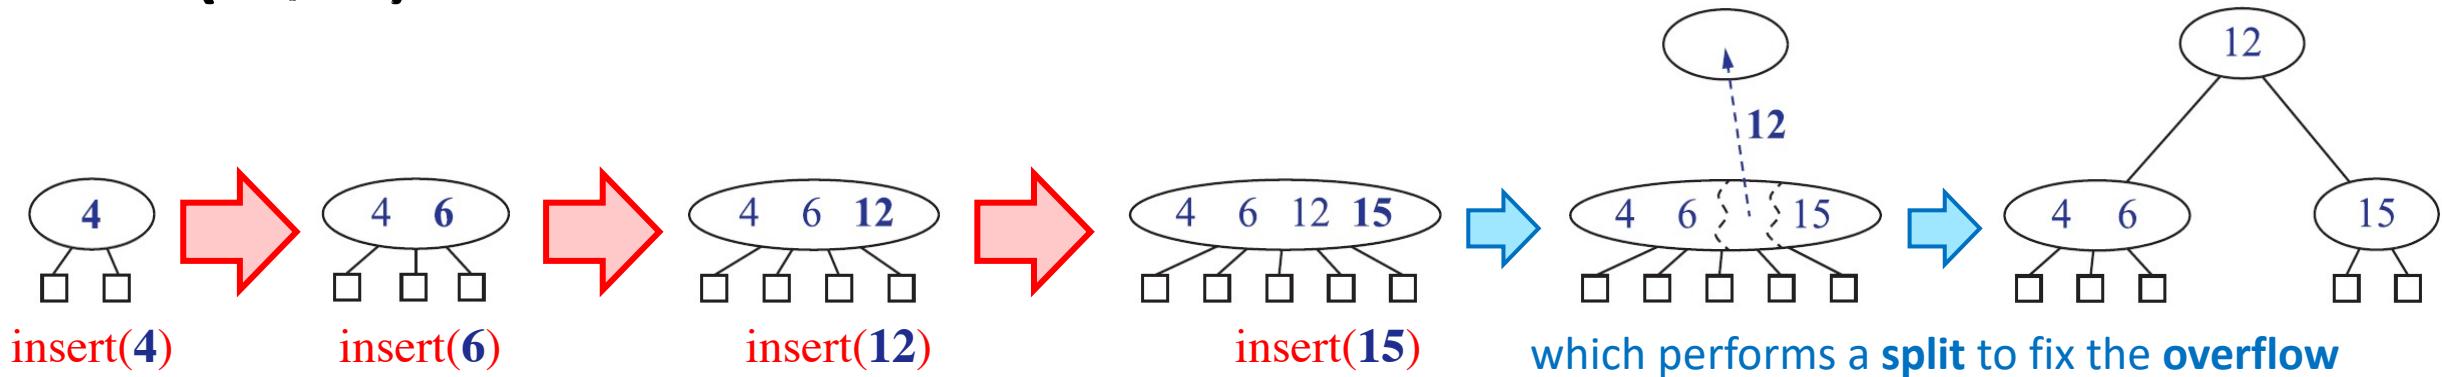
➤ v_1, v_2, v_3 go with v'

➤ v_4, v_5 go with v''

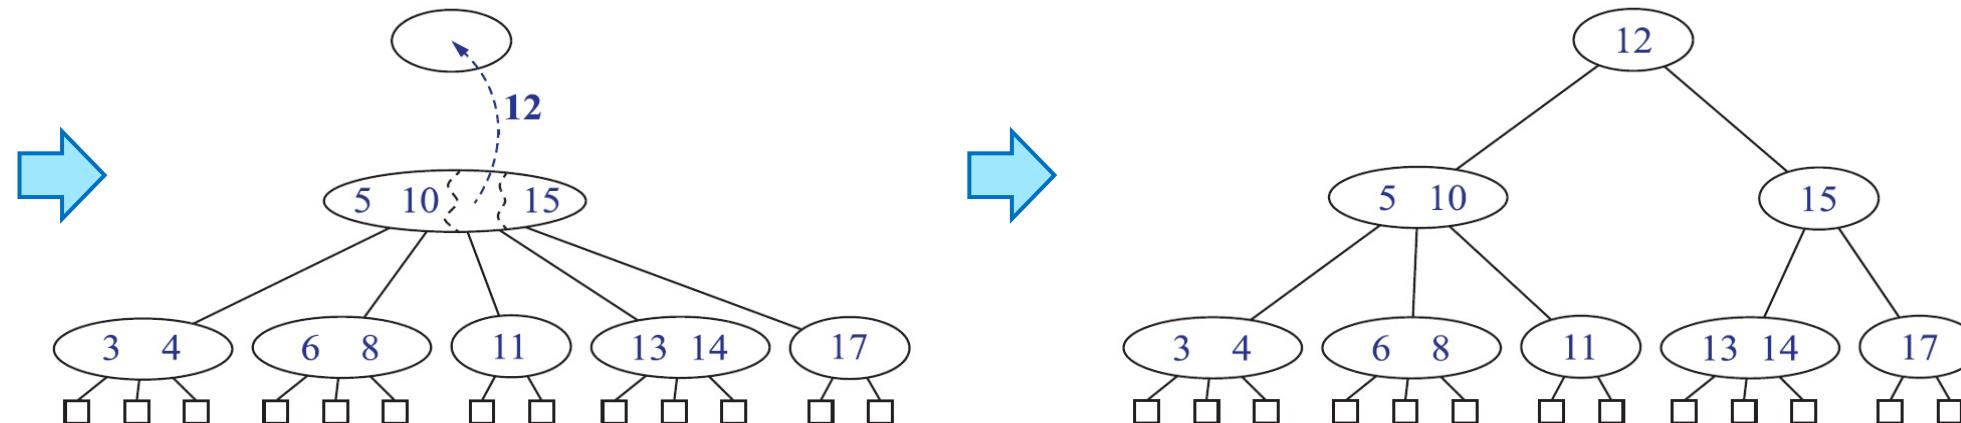
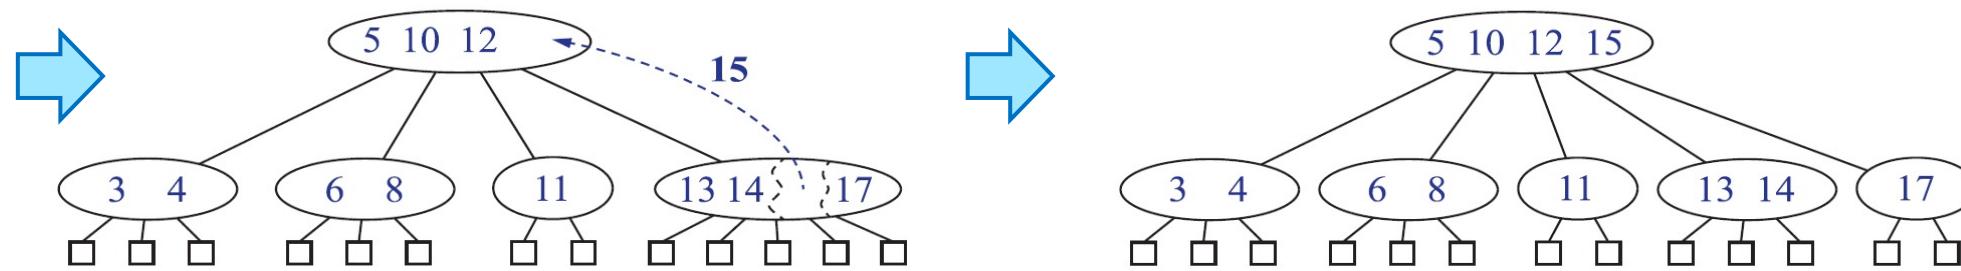
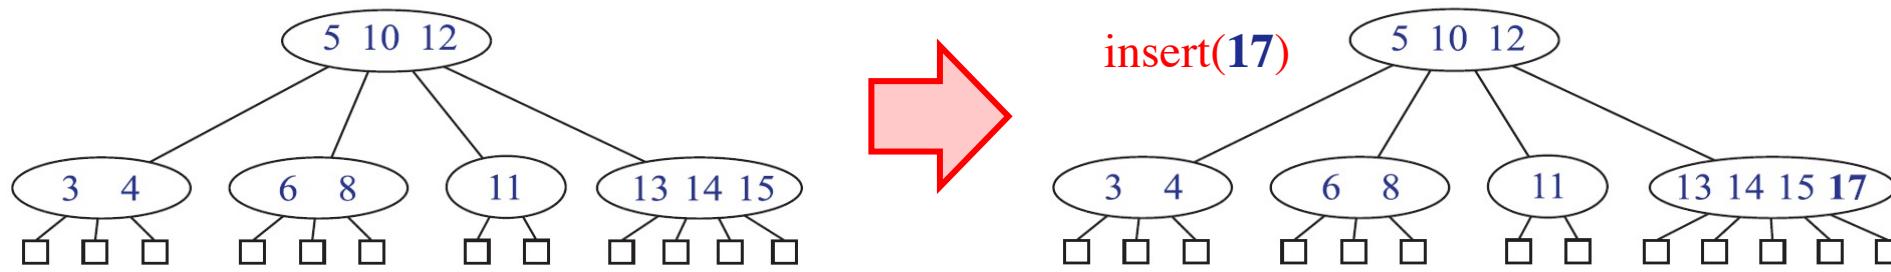


Note: when k_3 goes to the **parent** of v , if v didn't have a parent, we create one for it!

(2,4) TREES – INSERTION EXAMPLE



(2,4) TREES – INSERTION EXAMPLE



(2,4) TREES – INSERTION COMPLEXITY

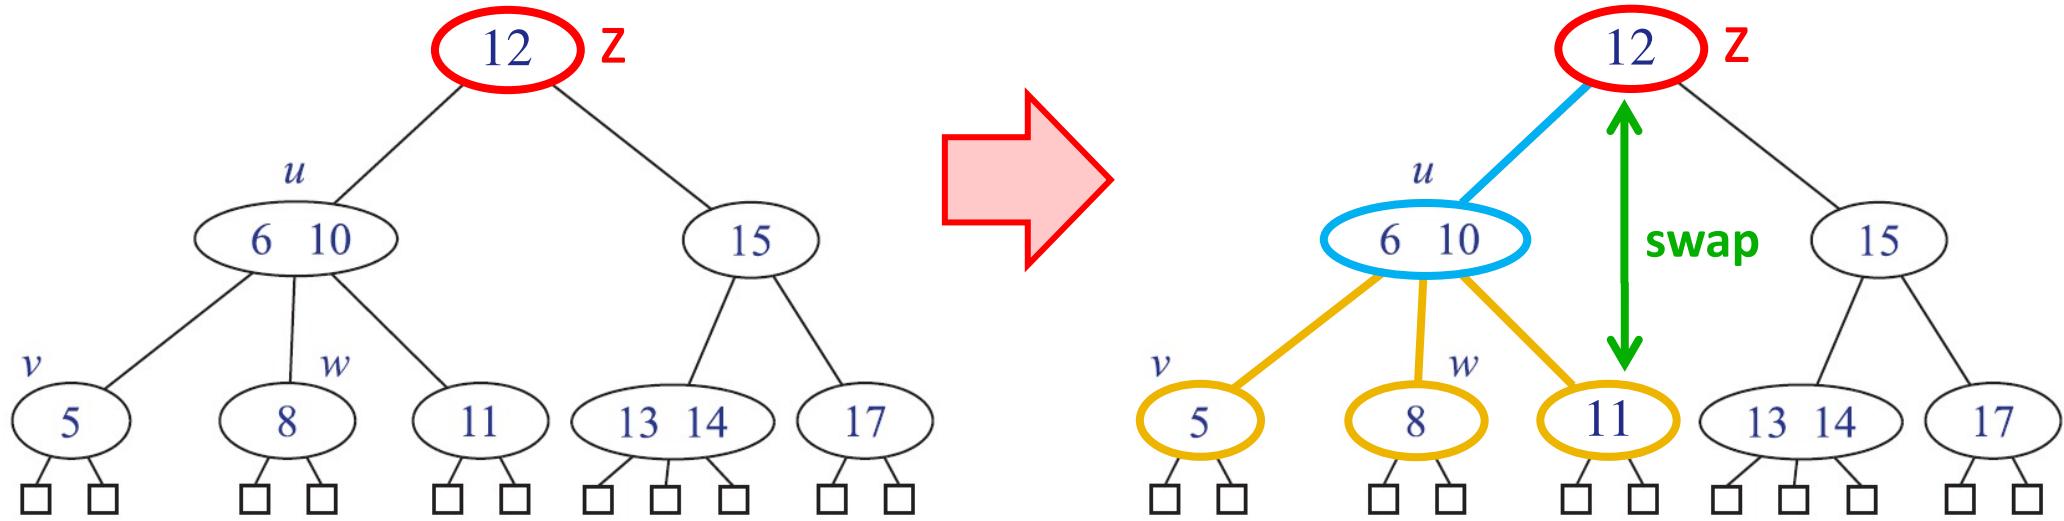
Let's analyze the complexity of inserting an entry (k, v) :

- First, we repeatedly **search** for an entry whose key = k , until we reach an **external node** where it would be suitable to insert (k, v)
 - This takes $O(\log n)$ time, since the **height** of a (2,4) tree is $O(\log n)$
- Then, we may have to repeatedly perform a **split** from that node all the way up to the root
 - We may perform $O(\log n)$ splits, since the **height** of the tree is $O(\log n)$
 - Each split requires a constant no. of operations, i.e., it takes $O(1)$ time

Thus, the **total time** to perform an insertion in a (2,4) tree is $O(\log n)$

(2,4) TREES – REMOVAL

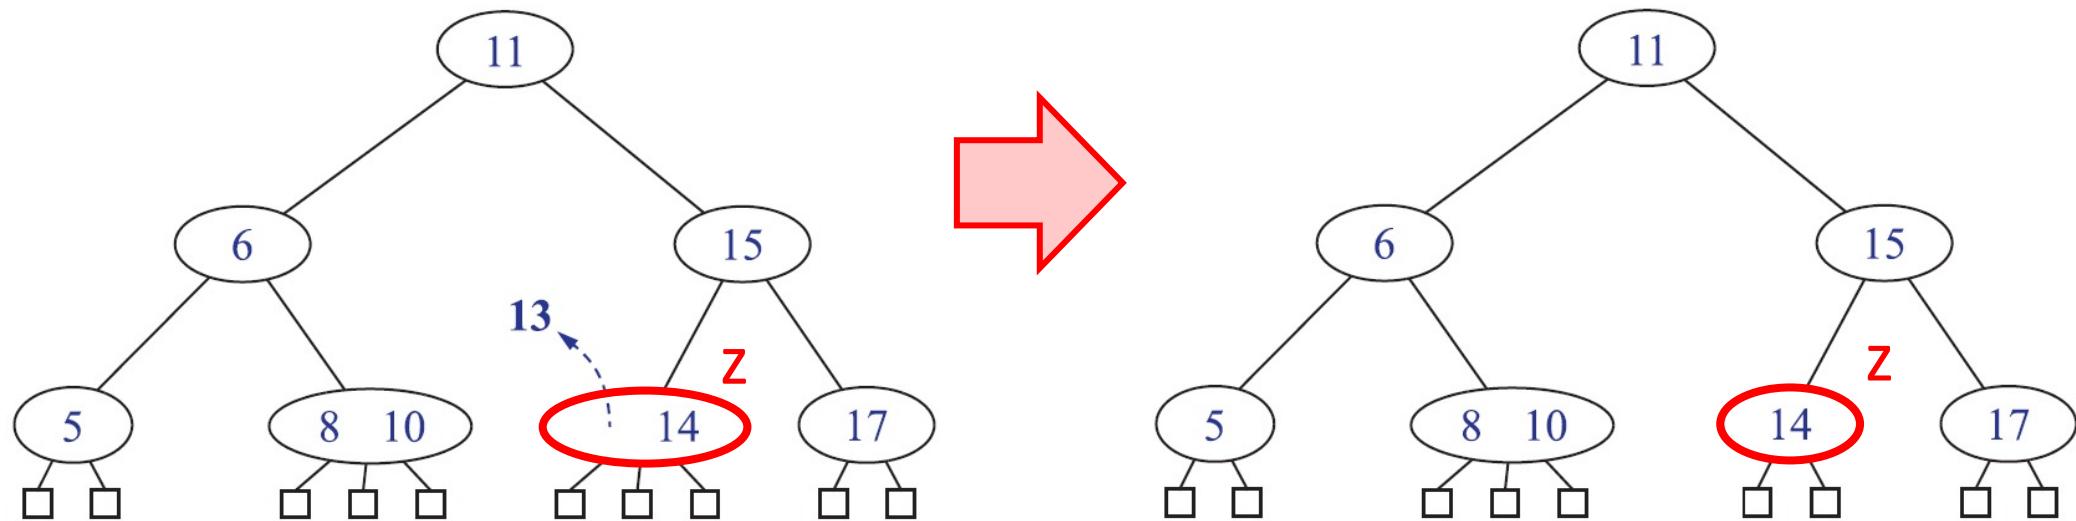
- We will only remove entries from a node whose children are **external nodes**!
- What if we want to remove the i^{th} entry from a node z that has only **internal children**?



- We find the last entry of the **right-most internal node in the subtree** rooted at the i^{th} child of z , and we **swap** it with i^{th} entry of z
- This way, the entry we want to remove is in a **node whose children are external**!

(2,4) TREES – REMOVAL

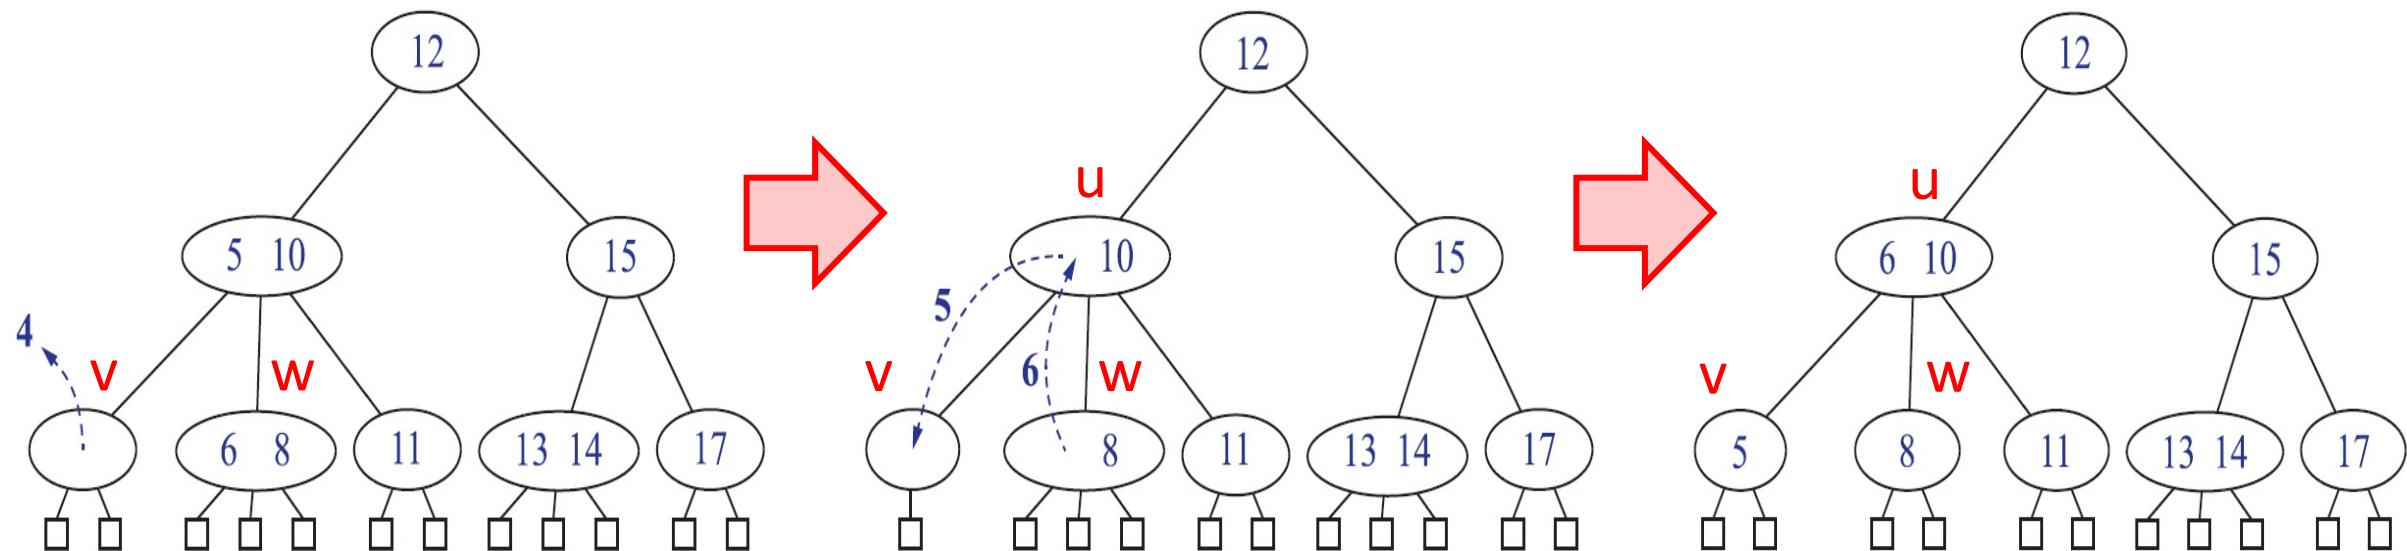
- So far, we ensured that the entry we want to remove is always going to be a node z whose children are external nodes



- Now, when removing an entry from z , we also remove the i^{th} external child of z
- The only problem is if z had only 1 entry (i.e., 2 children). Then, after removing that entry, z will be left with no entries and 1 child, which is not allowed in (2,4) trees! This situation is called an “underflow”

(2,4) TREES – HOW TO FIX AN UNDERFLOW

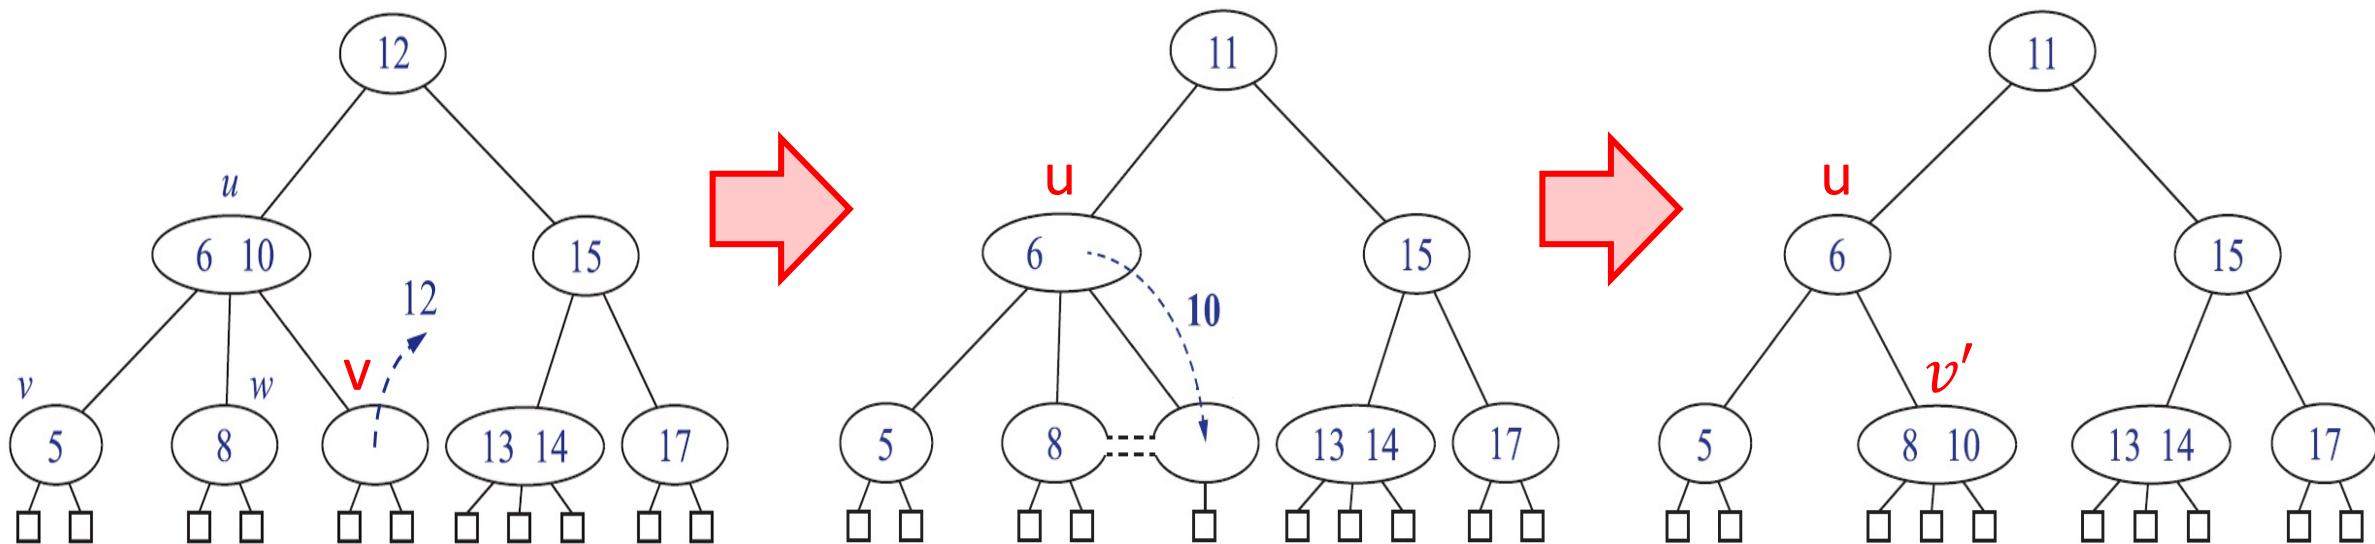
- If v has an **immediate sibling**, w , that has 3 or 4 external nodes, the sibling can afford to lose one of its external nodes (and entries) to help v out. In this case:
 - w gives one of its entries to the parent, u
 - the parent, u , gives one of its entries to v



- This is called a **transfer** operation

(2,4) TREES – HOW TO FIX AN UNDERFLOW

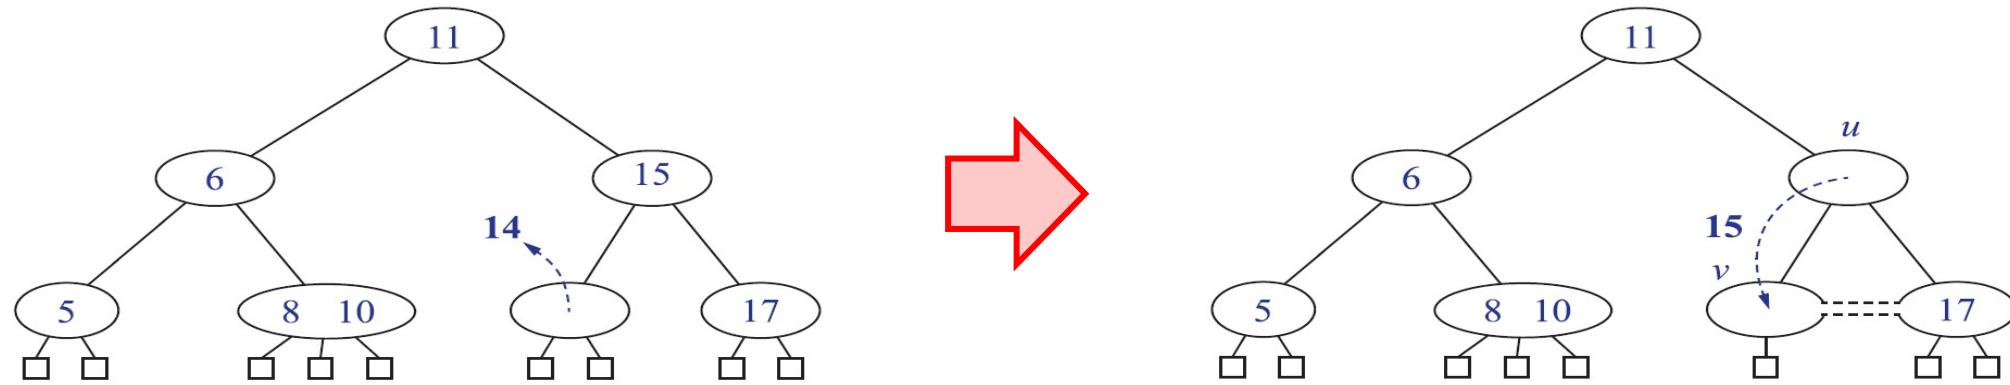
- But what if v does not have an immediate sibling, w , with 3 or 4 external nodes?
 - v is merged with its sibling, creating a new node v'
 - v' takes an entry from its parent, u



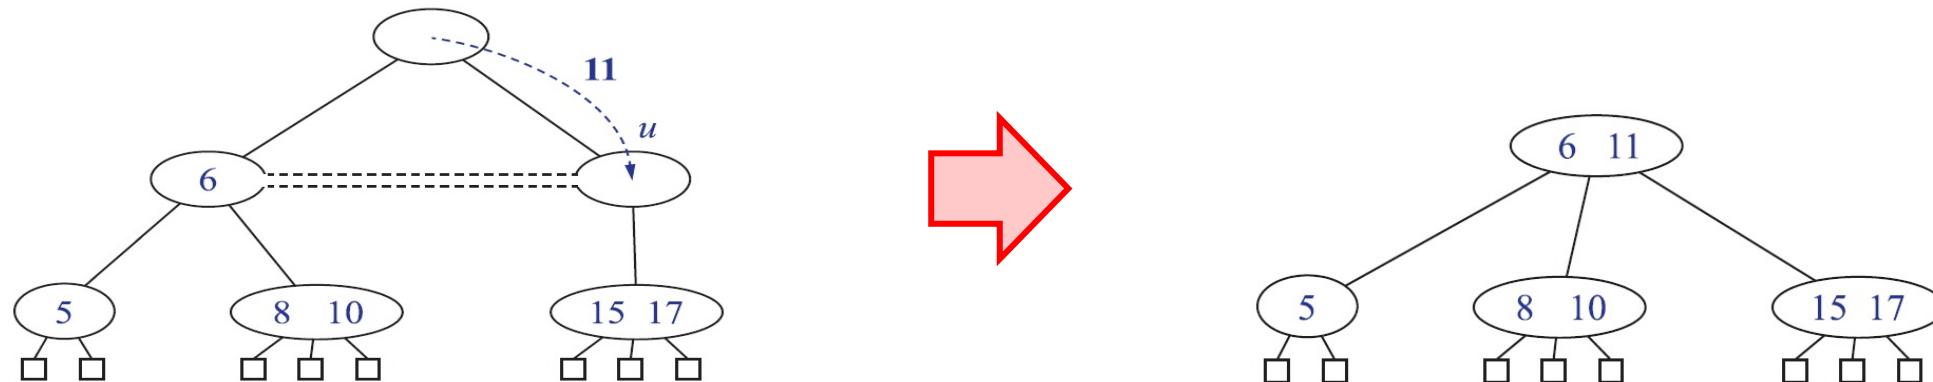
- This is called a **fusion** operation

(2,4) TREES – HOW TO FIX AN UNDERFLOW

- But what if, after v' takes an entry from its parent u in a **fusion operation**, u had no entries left, thereby causing an **underflow**?



- Just **repeat the process** for the parent with either a **transfer** or **fusion** operation (*in the special case where u is the root, just delete u after the fusion!*)



(2,4) TREES – REMOVAL COMPLEXITY

Let us analyze the complexity of removing an entry whose key = k

- First, we **search** for an entry with key = k , and if its children weren't external nodes, **we identify the external node** that should be swapped with it
 - This takes $O(\log n)$ time, since the **height** of a (2,4) tree is $O(\log n)$
- After that, we may have to **repeatedly perform a fusion** operation from that node all the way up to the root (possibly followed by a single **transfer** operation)
 - We may perform $O(\log n)$ such operations, since the tree **height** is $O(\log n)$
 - Each such operation requires a constant No. of steps, i.e., takes $O(1)$ time

Thus, the **total time** to perform a removal in a (2,4) tree is $O(\log n)$

53

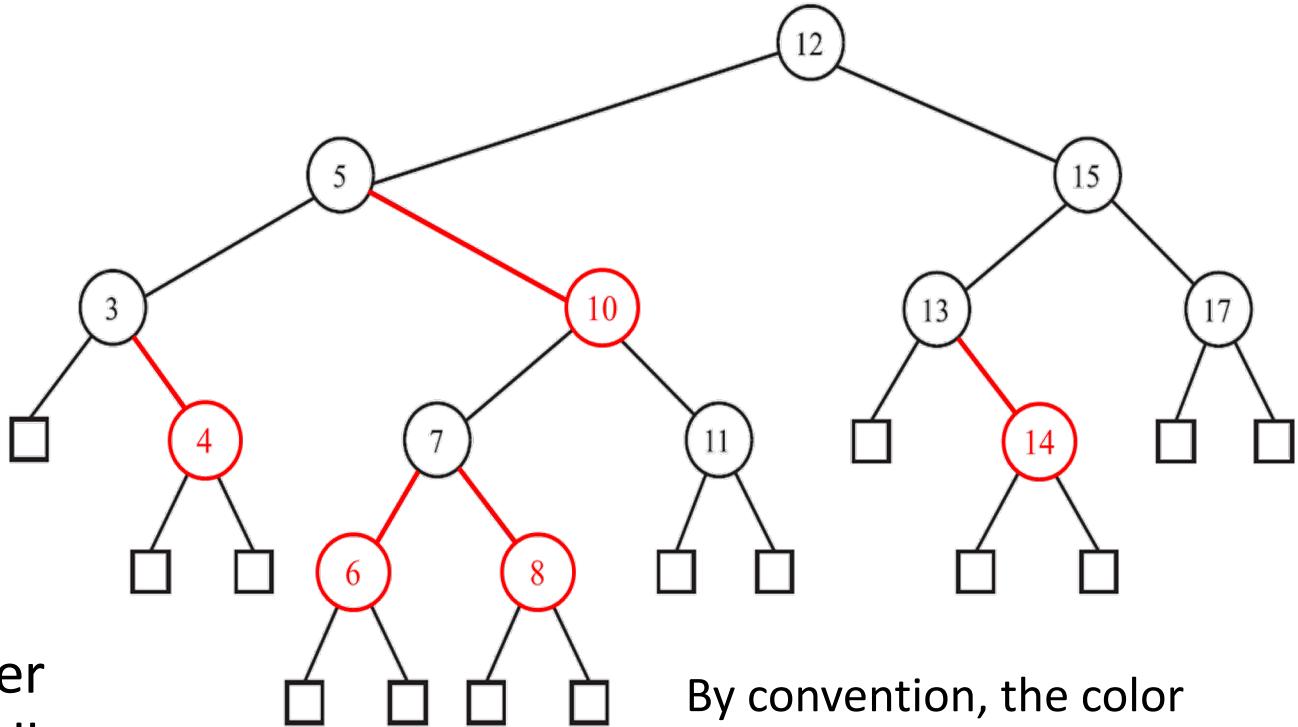
RED-BLACK TREES



RED-BLACK TREES

A **Red-Black** tree is a **BINARY** search tree with nodes colored **Red** or **Black** such that:

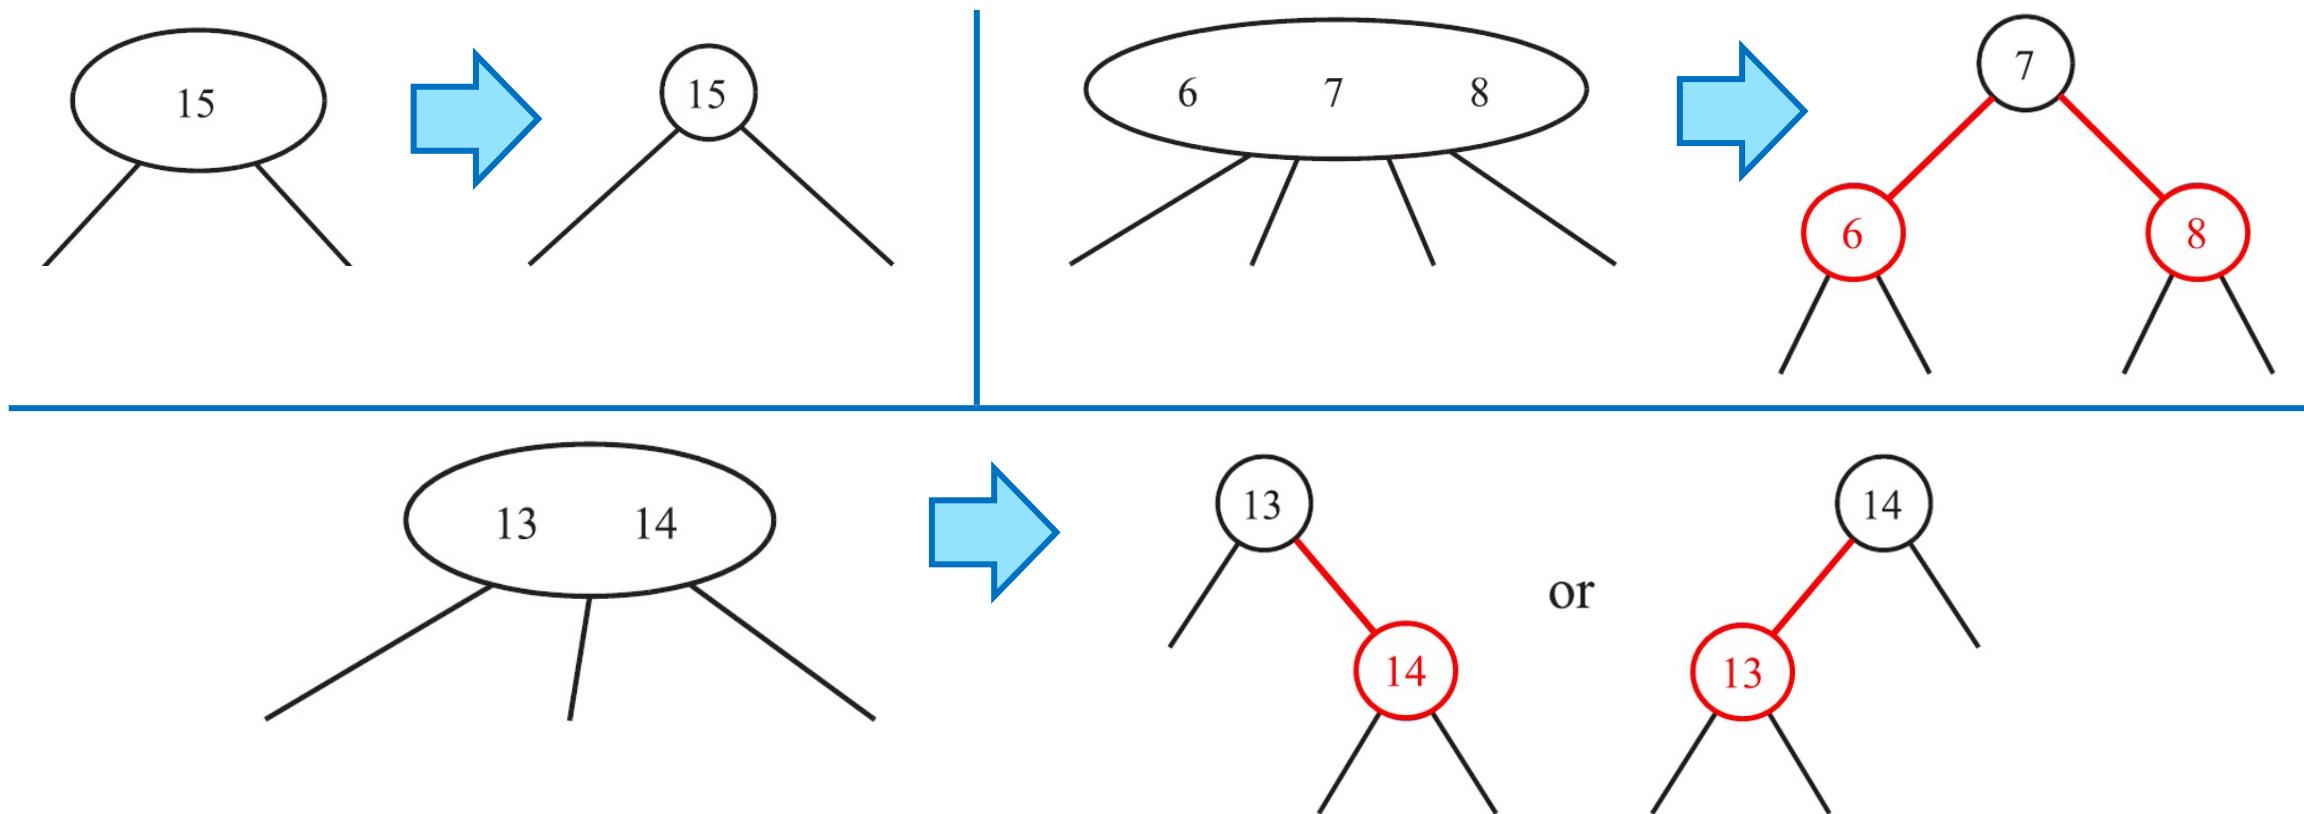
- **Root property:**
The root is **black**
- **External property:**
External nodes are **black**
- **Internal property:**
Children of a **red** node are **black**
- **Depth property:**
All external nodes have the same **black-depth**, defined as the number of **black** ancestors minus one (recall that a node is an ancestor of itself!)



By convention, the color of an edge = color of the node underneath it!

RED-BLACK VS. (2,4) TREES

- You can transform a (2,4) tree into a Red-Black tree, and vice versa as follows:

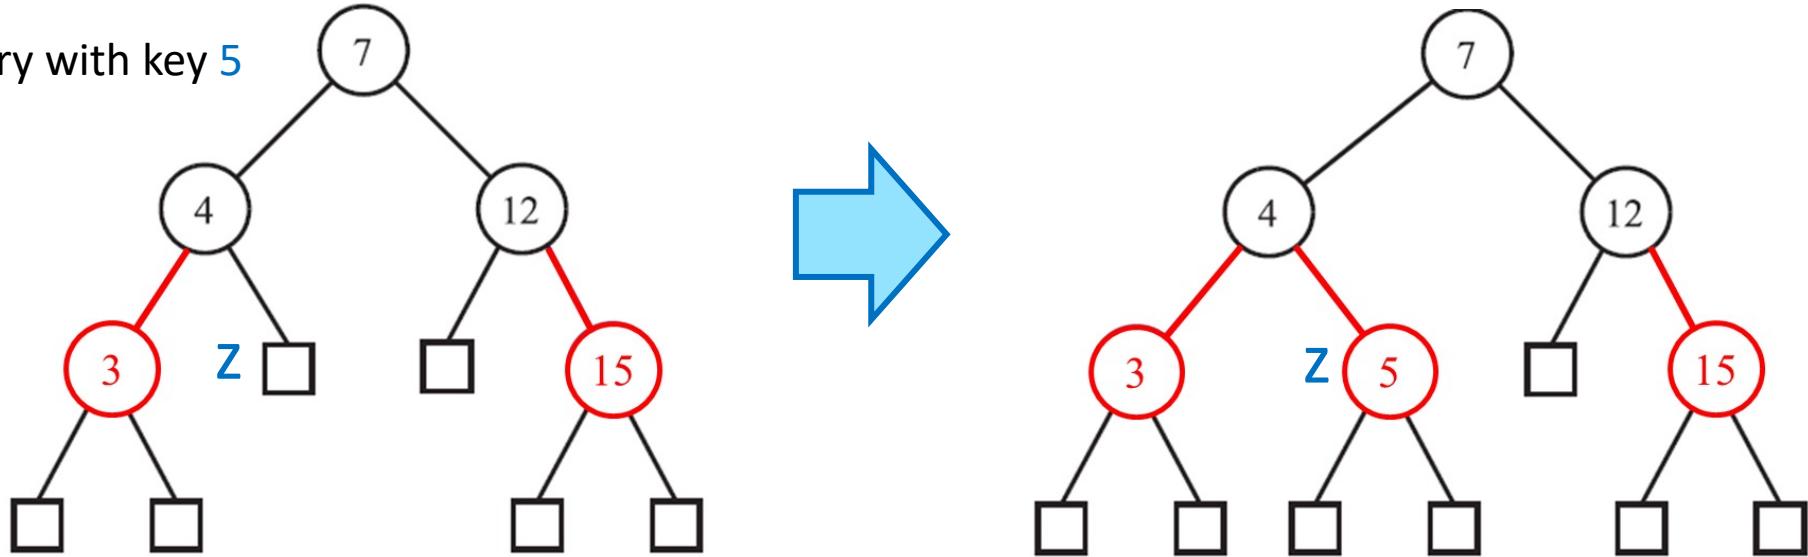


- Also, just like (2,4) trees, the height of a Red-Black tree with n entries is $O(\log n)$

RED-BLACK – INSERTION

- As with any search tree, to insert an entry (k, x) in a **Red-Black** tree, we repeatedly **search** for an entry whose key = k , until we reach an external node z

E.g. Insert an entry with key 5

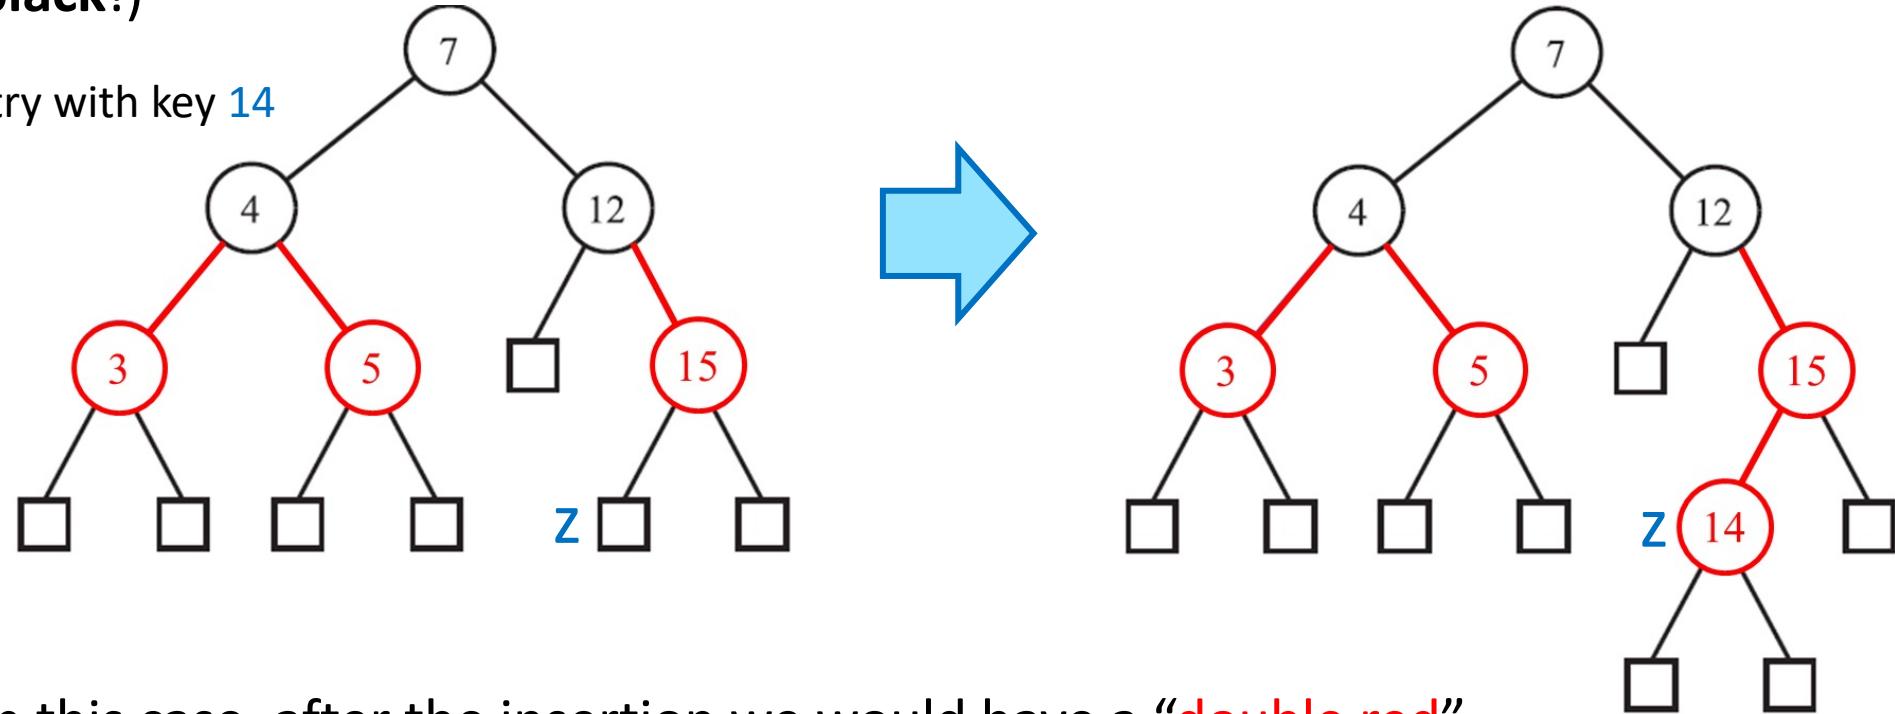


- Insert the entry in z
- Make z **red** (unless it's the root)
- Give z **black** external children

RED-BLACK – INSERTION

- The only problem is if **z** had a **red** parent. In this case, after the insertion we would have a “**double red**” violation (remember, children of **red** nodes must be **black**!)

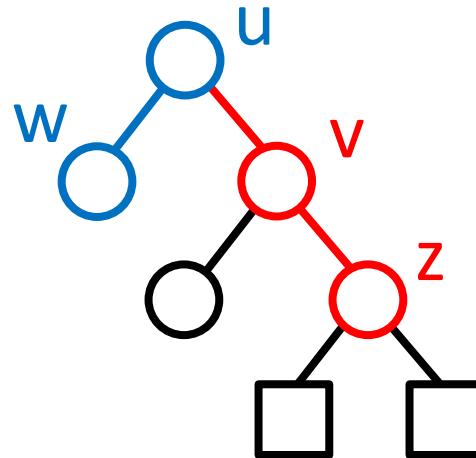
E.g. Insert an entry with key **14**



- In this case, after the insertion we would have a “**double red**” violation (remember, children of **red** nodes must be **black**!)

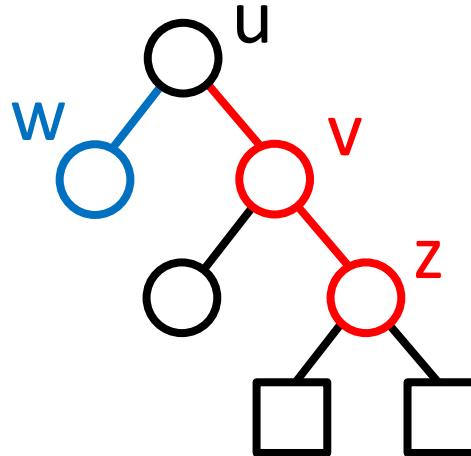
RED-BLACK – RESOLVING DOUBLE RED

- Let's use **blue** for when we haven't yet determined whether to use **Red** or **Black**
- Let's analyze the “**double red**” violation
 - Since v is **red**, it's definitely not the root (because the root of a **Red-Black** tree must be **black**).
 - Thus, v has a parent, u



RED-BLACK – RESOLVING DOUBLE RED

- Let's use **blue** for when we haven't yet determined whether to use **Red** or **Black**
- Let's analyze the “**double red**” violation
 - Since v is **red**, it's definitely not the root (because the root of a **Red-Black** tree must be **black**).
 - Thus, v has a parent, u
 - u must be **black**, because if it wasn't, we would have had a **double red** violation to start with (i.e., before the insertion), and we're assuming that we are inserting in a **Red-Black** tree!



RED-BLACK – RESOLVING DOUBLE RED

- Let's use **blue** for when we haven't yet determined whether to use **Red** or **Black**

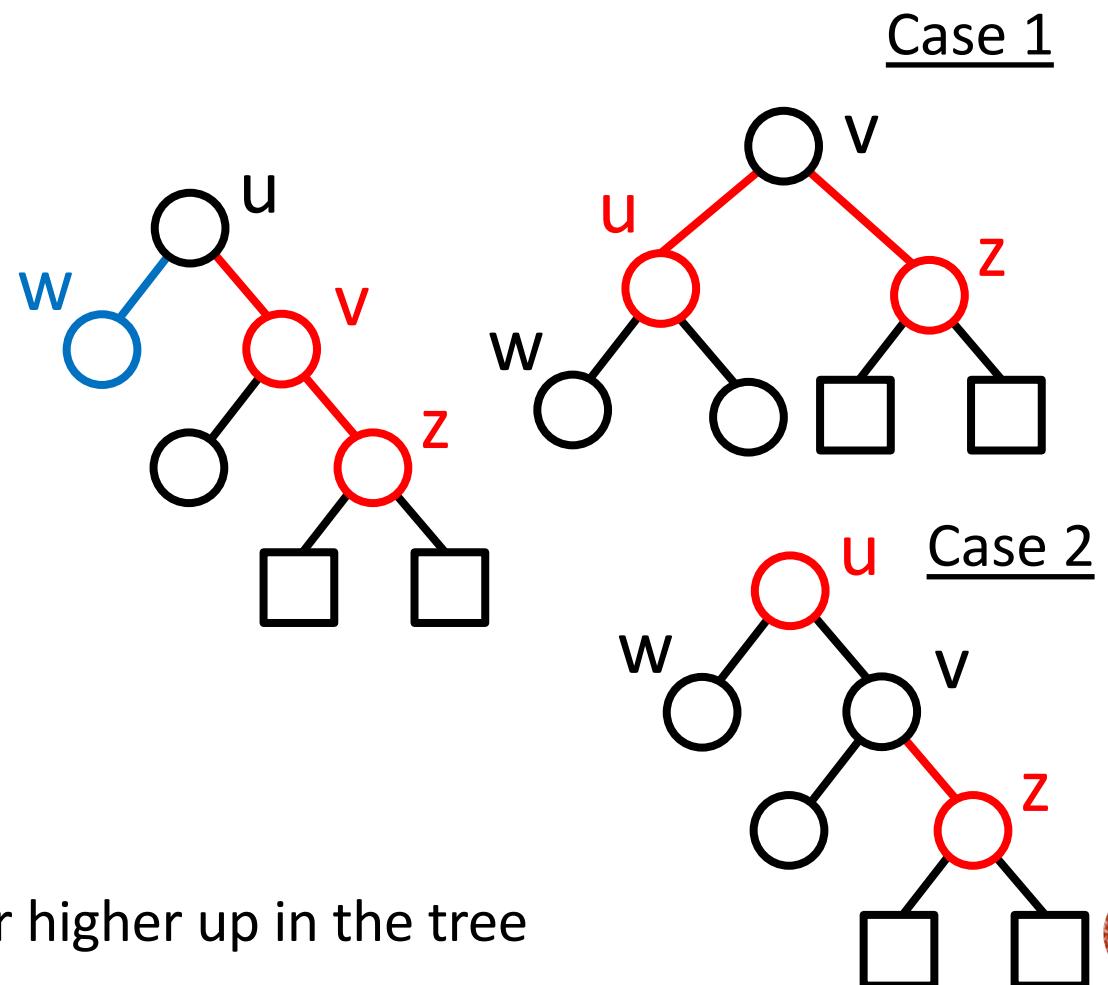
We remedy a **double red** as follows:

- Case 1: if **w** (the sibling of **v**) is **black**

- Rename (u,v,z) as (a,b,c) and use **trinode restructuring**
- Make a and c **red** and b **black**

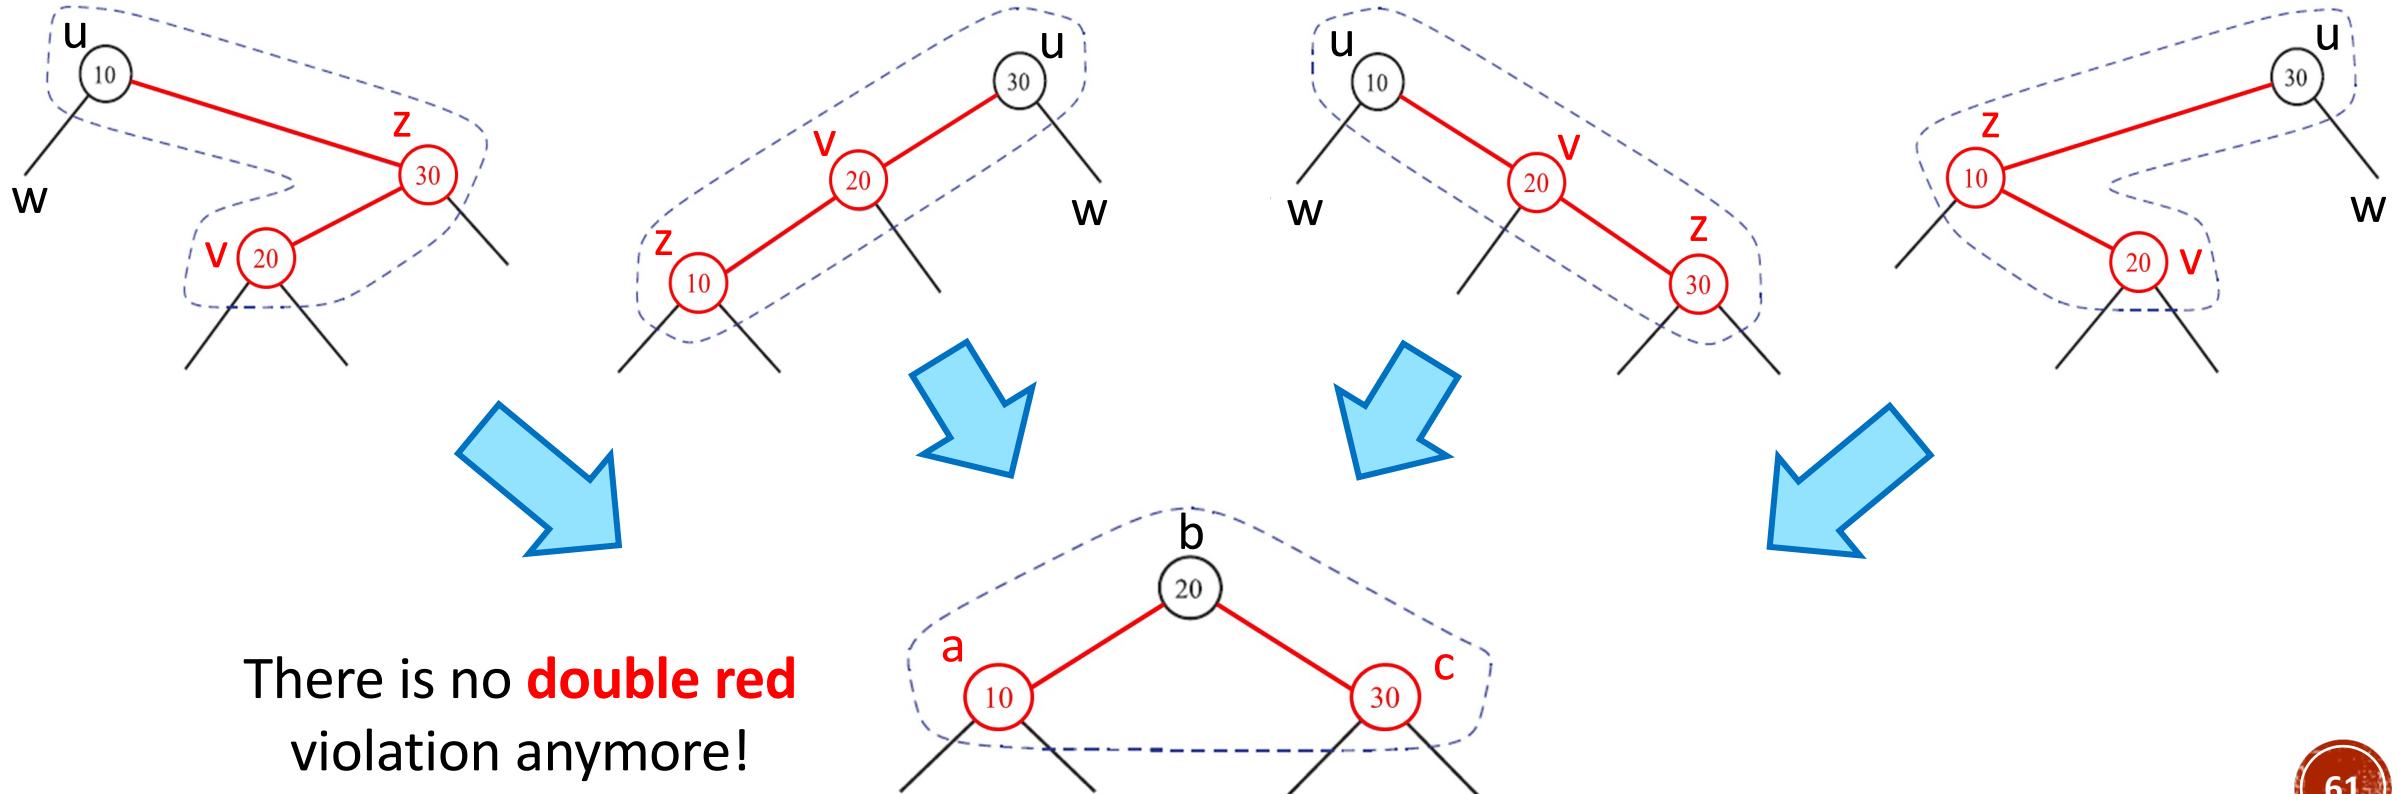
- Case 2: if **w** (the sibling of **v**) is **red**

- Make v and w **black**
- Make u **red** (unless it's the root)
- This is using **recoloring**
- After this, **double red** may re-appear higher up in the tree



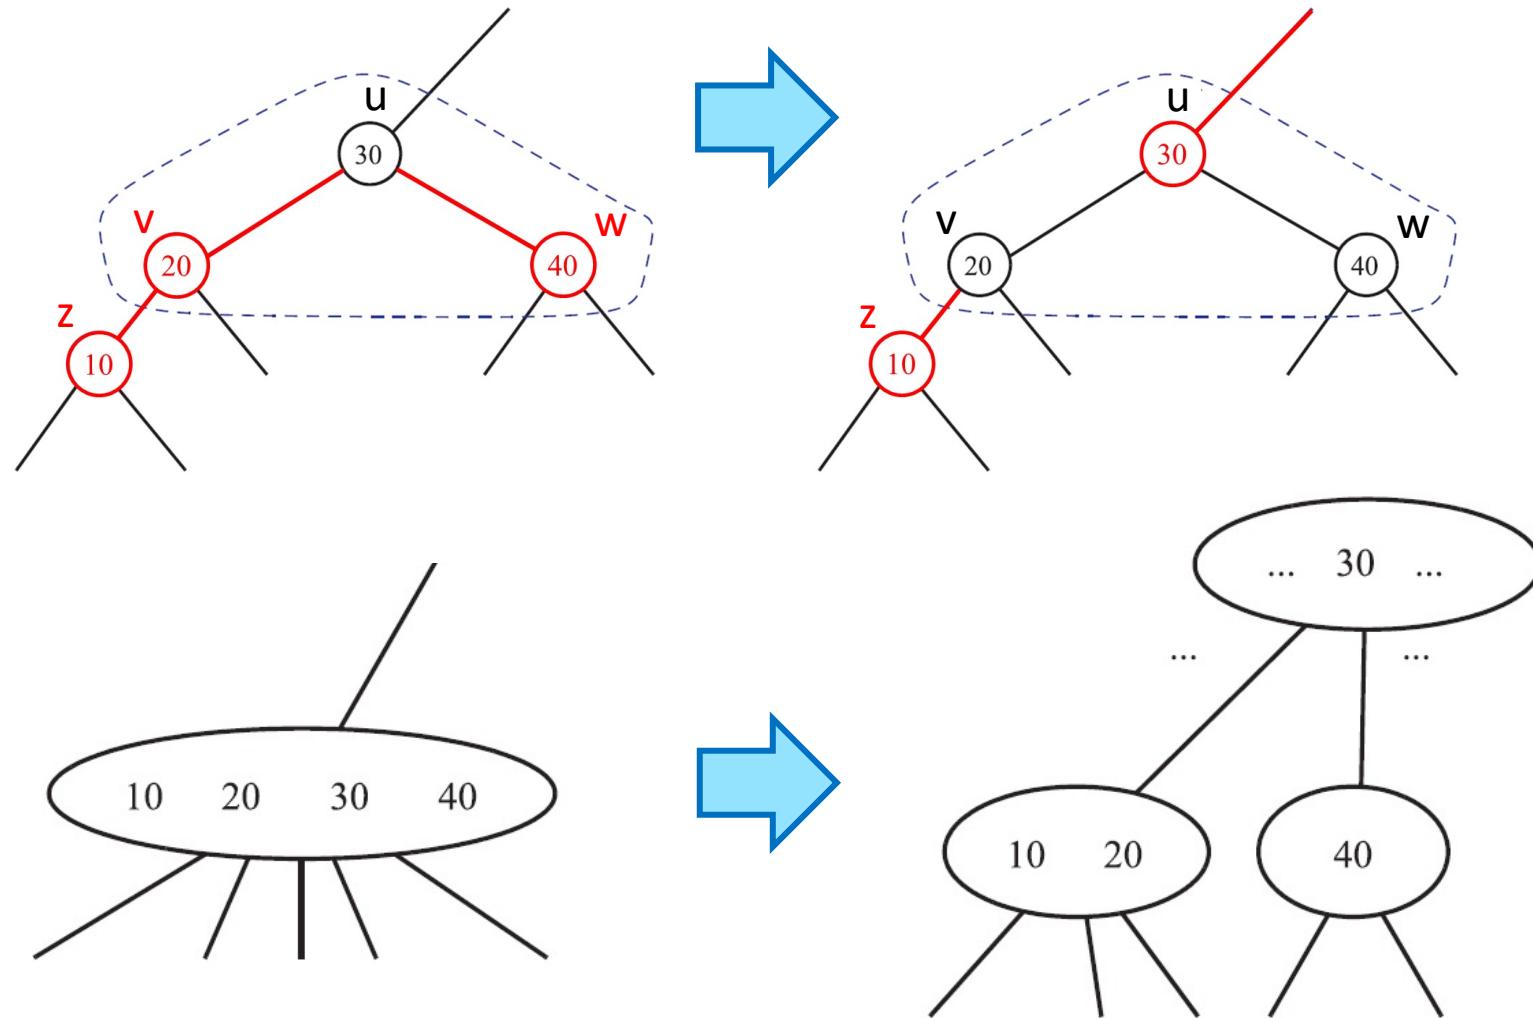
RED-BLACK – RESOLVING DOUBLE RED

- Case 1: if **w** is **black**
 - Rename (u,v,z) as (a,b,c) and use **trinode restructuring**
 - Make a and c **red**, and make b **black**



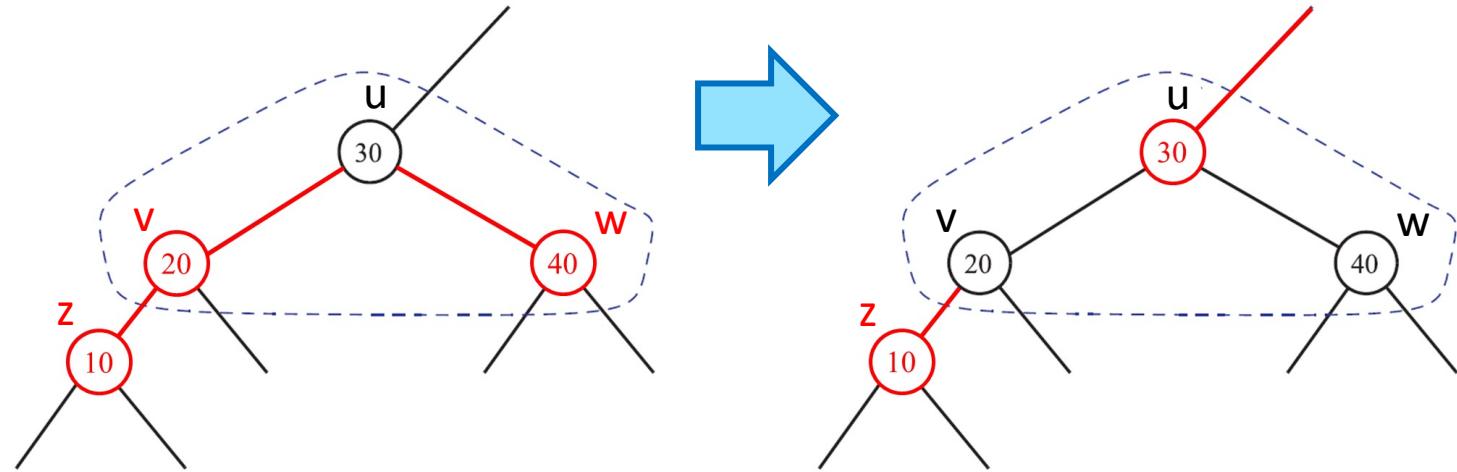
RED-BLACK – RESOLVING DOUBLE RED

- Case 2: if w is red
 - Make v and w **black**
 - Make u **red** (unless it's the root)
 - This operation, called “**recoloring**” is equivalent to a “**split**” operation in the corresponding (2,4) tree!



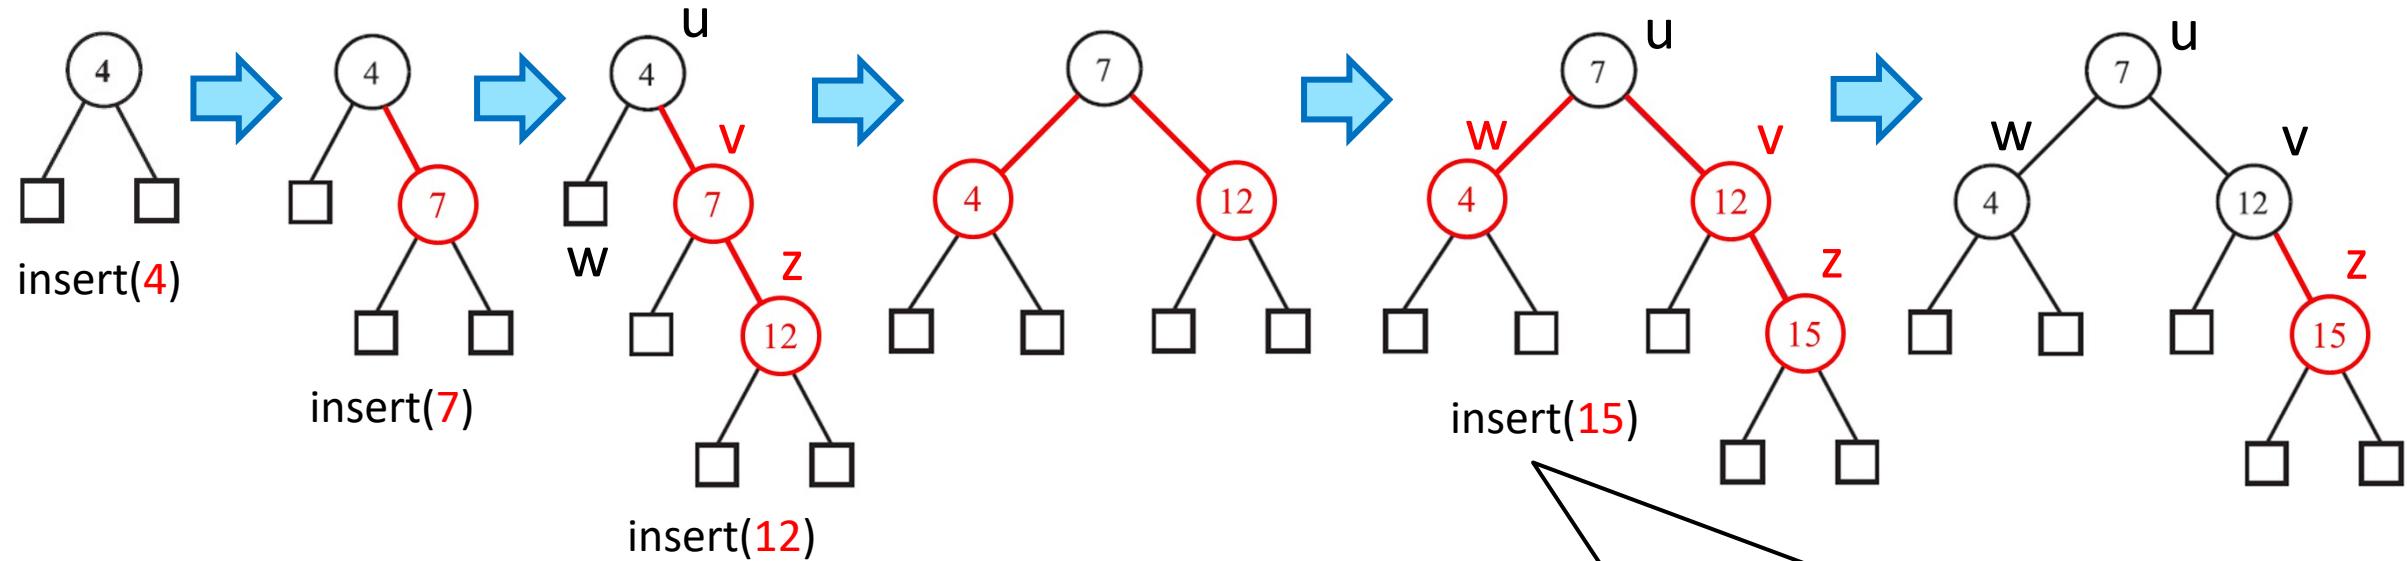
RED-BLACK – RESOLVING DOUBLE RED

- Case 2: if **w** is **red**
 - Make v and w **black**
 - Make u **red** (unless it's the root)



- A **recoloring** operation may introduce a **double red** violation higher up the tree, but such a violation is guaranteed to be solvable via a **recoloring**!
- A **recoloring** operation may introduce yet another **double red** violation, which would need to be resolved by either **trinode restructuring** or **recoloring**, etc.
- Thus, we may need **O(log n)** recoloring operations, and at most one **trinode restructuring** (the latter operation never results in another **double red** violation)

RED-BLACK – INSERTION EXAMPLES



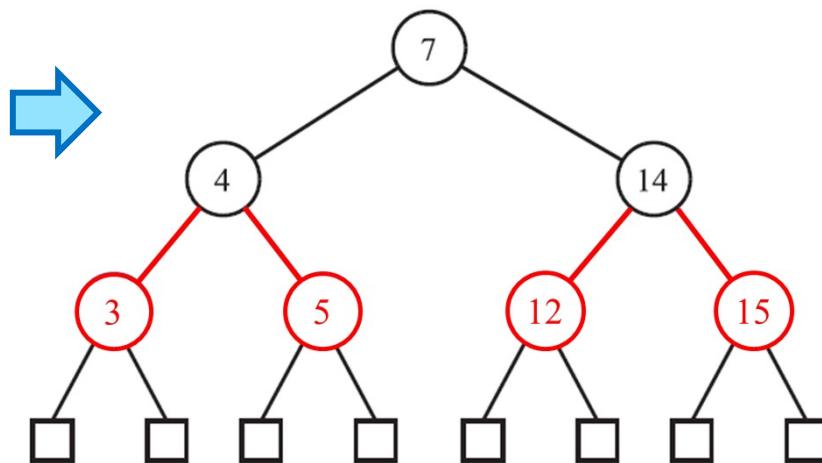
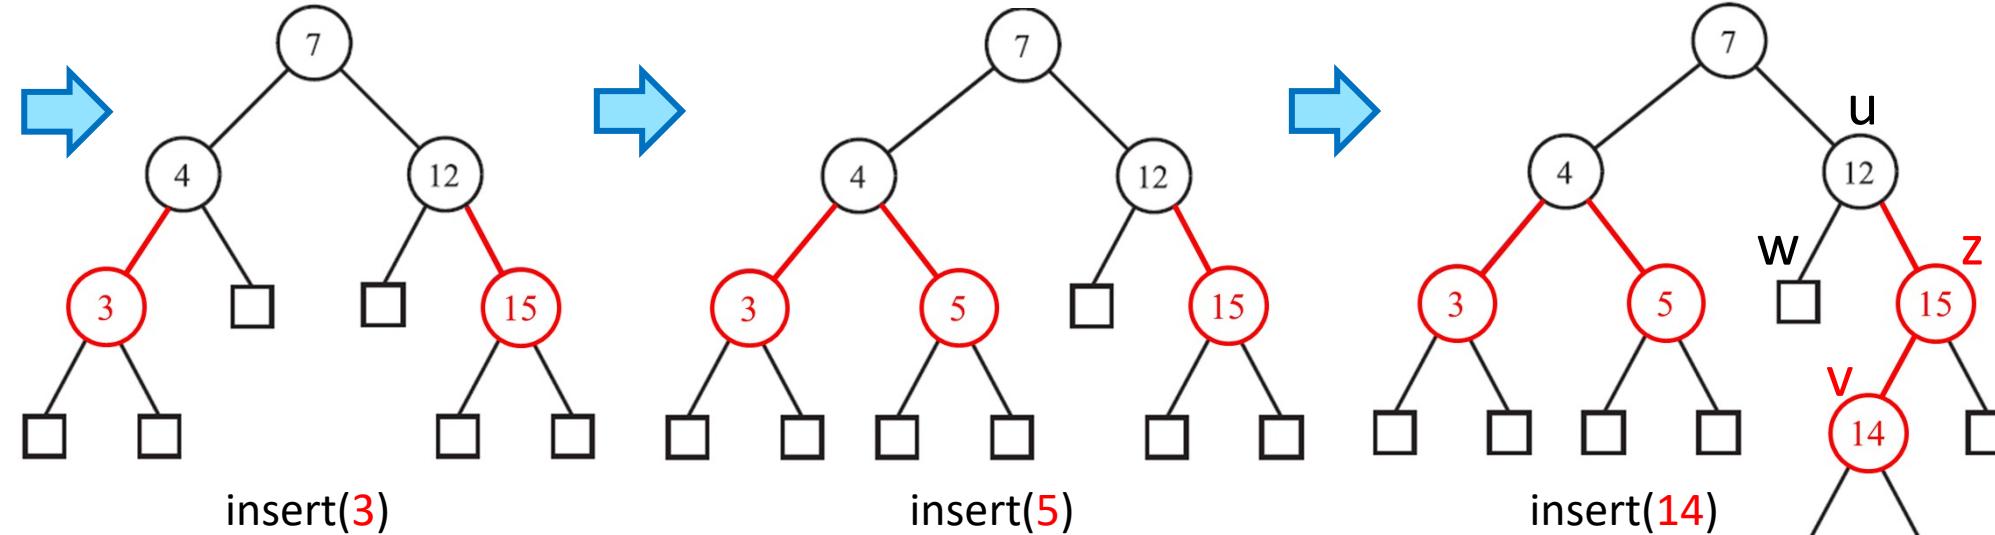
Double red violation!

- Case 1: w is black
 - Rename (u,v,z) as (a,b,c) and use **trinode restructuring**
 - Make a and c **red** and b **black**

Double red violation!

- Case 2: w is red (do recoloring)
 - Make v and w **black**
 - Make u **red** (unless it's the root)

RED-BLACK – INSERTION EXAMPLES



Double red violation!

- Case 1: w is black
 - Rename (u,v,z) as (a,b,c) and use **trinode restructuring**
 - Make a and c red and b black